

Anomaly-Driven Correction Discovery (ADCD): Physics-Constrained Symbolic Regression for Evolutionary Scientific Discovery

Muhammad Affif Erdita
Independent Researcher, Indonesia
maeapip10@gmail.com

June 2026

Abstract

Traditional symbolic regression discovers equations from scratch, an approach that becomes computationally intractable and overfitting-prone on complex physical systems under observational noise. We present **Anomaly-Driven Correction Discovery (ADCD)**, a physics-informed symbolic regression framework that instead takes a known classical law and searches for the dimensionless correction term Δ that reconciles it with anomalous observations. ADCD cascades physical filters—AST complexity, dimensional homogeneity with a transcendental-argument guardrail, and asymptotic-regime consistency (ARC)—with a JAX-traced L-BFGS-B optimizer and Bayesian Information Criterion (BIC) reranking, enforcing physical plausibility *before* numerical optimization and demonstrating robustness to experimental error.

Synthetic benchmark. On 9 physical anomalies spanning textbook, cross-domain, and synthetic systems under noise up to 10%, ADCD achieves a mean structural recovery of **80.4% ($\pm 7.4\%$) across sixteen seeds**, with 94.4% (34/36) at the reference seed=42 (disclosed explicitly as the highest-performing seed; the mean is the primary claim). At 5% noise, ADCD reaches 88.9% versus 11.1% for PySR under a matched residual formulation—a 77.8 percentage-point gap that persists even when PySR’s search budget is doubled. The same machinery recovers the correct structural class on 4/4 established real-world scenarios (Mercury perihelion, Lamb shift, muon $g-2$, blackbody) and fails safe on the misspecification benchmark.

Real-world: SPARC (honest framing). On the SPARC sample ($N_{\text{gal}} = 171$, $N_{\text{pts}} = 3342$), ADCD autonomously converges to a *two-parameter member of the Simple MOND algebraic family* (Famaey & Binney 2005), with a transition parameter $\hat{\theta}_1 \approx 0.27$ versus the canonical $c = 4.0$ (a factor of ~ 15). The discovered form achieves a 41% stacked-NMSE reduction over the fixed zero-parameter canonical MOND/RAR forms and, under fair galaxy-level cross-validation plus a 1,000-resample cluster bootstrap, cannot be statistically distinguished from two-parameter domain-expert forms ($\Delta\text{NMSE} = 0.008$, $z = 0.39$; $\delta\text{BIC}_{\text{eff}} = +3.4$, 95% CI brackets zero) while decisively beating two-parameter Standard MOND ($z = -4.60$). We claim autonomous rediscovery of this family from data, not a novel family.

Cosmological probes. Extended to cosmological datasets—63 $f\sigma_8$ growth-rate measurements and 34 $H(z)$ cosmic-chronometer points—ADCD finds *no* z -dependent functional correction beyond a constant amplitude rescaling in five independent tests (two observables, three homogeneous survey subsets). The best-fit amplitudes are consistent with the S_8 tension direction ($\sigma_{8,0} = 0.7815$ vs. Planck 0.811); the null result constrains but does not determine the physical origin of the tension.

Structural dichotomy. These results suggest a structural pattern—reported as an observation, not a proof: anomalies in individual gravitational systems (textbook, SPARC) exhibit functional corrections recoverable by ADCD, while large-scale cosmological tensions appear as amplitude offsets. Whether this reflects genuine physics or current survey-precision

limits is an open question we explicitly flag. We disclose all template overlaps, report recovered forms, and document a multivariable extension as an exploratory capability rather than a headline claim.

1 Introduction

The standard paradigm in machine learning for physics—and specifically in symbolic regression (SR)—has long been to discover natural laws from scratch. Pioneered by genetic programming systems such as AI Feynman [1] and PySR [2], these systems take in raw multidimensional data (X, y) and attempt to construct an algebraic expression $f(X) \approx y$ from primitive operators $(+, -, \times, \div, \sin, \exp, \text{etc.})$.

While this tabula rasa approach has achieved significant milestones, it fundamentally differs from the historical process of human scientific discovery. Science is rarely revolutionary; it is predominantly evolutionary. Physicists almost never discover new physical principles in a vacuum. Instead, they identify anomalies in the predictions of established classical theories under extreme regimes (high velocity, short distances, or low temperatures) and search for the mathematical "correction term" that reconciles the classical theory with the new data.

Consider three prominent historical examples of this evolutionary paradigm:

1. **Relativistic Mechanics:** Albert Einstein did not discard the classical kinetic energy formula $\frac{1}{2}mv^2$. Instead, by analyzing the anomalous behavior of high-velocity particles, the relativistic correction $\Delta_{\text{rel}} \approx \frac{3}{4} \frac{v^2}{c^2}$ was uncovered, leading to the full expansion of relativistic kinetic energy:

$$E_k = \frac{1}{2}mv^2 \left(1 + \frac{3}{4} \frac{v^2}{c^2} + \frac{5}{8} \frac{v^4}{c^4} + \dots \right)$$

2. **Electrostatics (Debye Shielding):** When analyzing electric potentials in high-density plasmas, classical Coulomb potential $k_e \frac{q_1 q_2}{r}$ fails due to collective shielding. Physicists discovered the Screened Coulomb (Yukawa-like) potential by applying an exponential correction term:

$$\Phi(r) = k_e \frac{q_1 q_2}{r} \left(e^{-r/\lambda_D} \right)$$

Here, the correction term is multiplicative: $1 + \Delta = e^{-r/\lambda_D}$, implying $\Delta = e^{-r/\lambda_D} - 1$, which naturally decays to 0 at the classical limit $r \rightarrow 0$ (at very short distances shielding is negligible).

3. **Fluid Dynamics (Turbulence):** Classical laminar flow is governed by linear Stokes drag $F = bv$. As the velocity v increases and the system enters turbulent regimes, a quadratic correction term emerges, leading to the additive formulation:

$$F_{\text{drag}} = bv + \alpha v^2$$

Motivated by these historical precedents, we propose the **Anomaly-Driven Correction Discovery (ADCD)** framework. ADCD contributes at three levels. At the **conceptual level**, it introduces the correction-first paradigm: searching for Δ in a restricted subspace $\mathcal{D} \subset \mathcal{F}$ rather than the full function space $\mathcal{F}$. This correction-first paradigm is designed for anomaly-driven theory refinement—where the classical baseline is structurally correct and the anomaly is incremental—rather than wholesale theory replacement. At the **algorithmic level**, ADCD implements physics-gated search: AST complexity, dimensional homogeneity with transcendental guardrail, and asymptotic limit (ARC) gates operate as a collective physical consistency filter that screens candidates *before* any numerical optimization occurs, with adaptive dimensional relaxation on free parameters as detailed in Section 7. Empirically, ablation on the primary

benchmark (Figure 5) shows that removing the ARC, AST, or Coarse-Eval gate individually does not change structural accuracy at 5% noise; removing the dimensional gate alone reduces accuracy by 11.1 percentage points (8/9 to 7/9), confirming its role as the primary structural filter. Only removing *all* gates simultaneously causes a larger collapse (8/9 to 4/9), indicating that the gates reinforce each other as a collective physical consistency filter. At the **engineering level**, ADCD provides a JAX-traced L-BFGS-B optimizer, a Python package with a public API, CLI interface, and CI/CD pipeline for reproducible benchmarking.

Scope and extended validation. In this work, we extend ADCD beyond synthetic validation to three classes of real-world anomaly at progressively larger physical scales: *galactic* rotation curves (SPARC), *stellar-system* kinematics (wide-binary stars), and *large-scale cosmological* structure growth ($f\sigma_8$) and expansion history ($H(z)$). Each class yields a qualitatively different outcome—functional recovery, ambiguous cross-system transfer, and an amplitude-only null result respectively—revealing a structural pattern in how anomalies manifest across physical scales (Section 5.8).

Unique contribution. To our knowledge, ADCD is the first systematic framework capable of distinguishing *functional* anomalies—which require z - or r -dependent corrections to known physical laws—from *amplitude* anomalies—which manifest as a pure rescaling of constants—across multiple physical scales. Existing symbolic regression methods recover expressions but do not classify the *kind* of anomaly a residual encodes; this distinction, surfaced here as the structural dichotomy of Section 5.8, is the organizing finding of the real-world validation.

2 Related Work

2.1 Symbolic Regression

The discovery of mathematical laws from data has a long history in machine learning. BACON [13] pioneered rule-based numerical law discovery from experimental tables. Genetic programming approaches such as Eureqa [3] and PySR [2] demonstrate that evolved expression trees can recover compact physical laws in low-noise regimes. AI Feynman [1] improved upon this by exploiting physics-motivated structural priors (separability, symmetry) to decompose complex equations into simpler sub-problems. Deep Symbolic Regression (DSR) [9] applies reinforcement learning to navigate the symbolic search tree. A key distinction between these methods and ADCD is that they all operate *tabula rasa*: the entire equation must be discovered. ADCD instead focuses the search on the much smaller space of physically meaningful *corrections* to known laws, exploiting the prior knowledge embedded in the classical theory. PhySO (Φ -SO) [14] incorporates dimensional analysis as a hard constraint during tabula rasa symbolic regression, using deep reinforcement learning to generate only unit-consistent expression trees. This directly addresses the problem of dimensionally impossible candidates in the search space—a concern that also motivates ADCD’s dimensional homogeneity gate. The two frameworks diverge in problem formulation and scope: PhySO discovers full governing equations from data, whereas ADCD restricts search to physically meaningful *corrections* Δ relative to a known classical baseline, substantially reducing the hypothesis space. ADCD further complements unit checking with asymptotic regime consistency (ARC) and complexity filtering, applied specifically to the correction subspace. Unlike PhySO’s strict unit enforcement, ADCD’s dimensional gate employs adaptive parameter relaxation on free parameters, as documented in Section 7. LLM-assisted symbolic regression methods such as LaSR [15] accelerate hypothesis generation by using language-model queries to construct and evolve an *abstract concept library* that guides evolutionary search, rather than proposing expressions directly. These approaches improve search efficiency but do not enforce physical consistency on candidate expressions before numerical

fitting. ADCD adopts a complementary design in its Hybrid Proposer: LLM-generated correction templates are subjected to the same physics-gated cascade—complexity, dimensional analysis with transcendental guard, and asymptotic limits—used for template-based proposals, filtering semantically plausible but physically inconsistent candidates *before* Stage-2 optimization.

Modified gravity and cosmological-tension context. The two real-world physics domains we target in Sections 5.6–5.7 have well-developed domain-specific literatures that motivate the correction-first framing. For the galactic-scale tests, the family of MOND interpolating functions is reviewed comprehensively by Famaey & McGaugh [22], which we use as the reference for the Simple/Standard/RAR algebraic families against which the SPARC-discovered form is compared (Section 5.6). For the cosmological-scale tests, the $f\sigma_8$ compilation of Alestas et al. [23] documents the S_8 tension as a residual of observed growth-rate measurements relative to the Planck-calibrated Λ CDM prediction—precisely the form of anomaly ($y_{\text{obs}} - y_{\text{classical}}$) for which ADCD is designed, and the dataset on which Section 5.7 reports a quantitative CONSTANT_WINS bound.

2.2 Physics-Informed Machine Learning

Physics-Informed Neural Networks (PINNs) [6] embed differential equation constraints into the training loss of neural networks, enforcing physical consistency during learning. Hamiltonian Neural Networks [7] and Lagrangian Neural Networks [8] directly parameterize the energy function, guaranteeing energy conservation by construction. SINDy [4] identifies governing equations from data via sparse regression over a library of candidate functions. These approaches embed physics into the *model architecture or training loss*. ADCD differs by embedding physics into the *search space constraints*: the three cascaded gates (complexity, dimensional analysis with transcendental guard, asymptotic limit) filter candidates *before* any numerical optimization occurs, with adaptive dimensional relaxation on free parameters as detailed in Section 7, making the pre-optimization screening independent of the Stage-2 optimizer.

2.3 Automated Scientific Discovery

The Nobel Turing Challenge [10] articulates a vision for AI systems that can autonomously generate and test scientific hypotheses. FunSearch [11] demonstrates that large language models paired with evolutionary search can discover novel mathematical structures in combinatorics. ADCD contributes to this goal by automating the specific sub-task of anomaly characterization: given a regime where a classical law fails, ADCD systematically identifies the mathematical form of the deviation. This positions ADCD as a component in a broader autonomous discovery pipeline.

2.4 Model Selection and Complexity Control

Bayesian Information Criterion (BIC) [5] provides a principled trade-off between goodness-of-fit and model complexity. Minimum Description Length (MDL) [12] offers an information-theoretic complement. Both are related to Solomonoff’s algorithmic complexity. ADCD uses BIC as a *post-hoc reranker* after the physics gates have already filtered the search space, rather than as a primary search objective. This separation of concerns allows the gates to screen for physical plausibility while BIC selects the most parsimonious survivor.

Bayesian posterior model averaging (Phase 3). We extend BIC reranking with a Bayesian posterior distribution over candidate corrections using the Schwarz approximation (Kass & Raftery 1995 [16]): posterior weight $w_i \propto \exp(-\Delta\text{BIC}_i/2)$, where $\Delta\text{BIC}_i = \text{BIC}_i - \text{BIC}_{\min}$. This converts discrete ranking into a continuous probability distribution over functional families,

enabling principled quantification of discovery uncertainty. When the top candidate’s weight exceeds $3\times$ the second-best’s weight (decisive evidence threshold), the discovery is classified as unambiguous. An accompanying identifiability analyzer diagnoses failure modes—model degeneracy, low SNR, or undetectable correction magnitude—providing scientifically meaningful explanations when the search is inconclusive.

Residual Asymptotic Signature (RAS) as a novel feature prior. Prior to candidate proposal, ADCD extracts the *Residual Asymptotic Signature* (RAS): a log-log regression of the residual near the classical limit, yielding a leading exponent that serves as a structural prior for the proposer. This allows the grammar-based proposer to rank candidate ratio seeds by their Pearson correlation with the observed residual, preferentially generating correction templates whose asymptotic scaling matches the empirical data. RAS is a novel contribution that bridges residual analysis and grammar-guided proposal, improving targeted search efficiency without requiring additional physics assumptions.

3 The ADCD Framework

The ADCD pipeline operates as a closed-loop physics-informed search. A high-level overview of the architecture is illustrated in Figure 1.




Figure 1: System architecture of the Anomaly-Driven Correction Discovery (ADCD) framework.

Relation to Effective Field Theory. ADCD is organized along the operating philosophy of Effective Field Theory (EFT) [27–29], the dominant paradigm by which modern physics constructs corrections to known laws: starting from a known low-energy (classical) theory, one enumerates correction operators consistent with the symmetries and dimensional structure of the problem, filters them by their asymptotic behavior, and orders them by complexity so that the leading terms dominate. Under this perspective, the benchmark scenarios evaluated in this work (e.g., Relativistic KE, Net Radiation) correspond to recovering the leading-order correction operator directly from data, in the same sense that an EFT practitioner identifies the dominant correction term at the next order of the theory expansion. Each stage of the ADCD pipeline mirrors a step in this workflow—the dimensional gate corresponds to dimensional analysis of operators, the ARC gate ($\Delta \rightarrow 0$ at the classical boundary) corresponds to the decoupling of heavier degrees of freedom in the low-energy limit, and the BIC penalty plays a role analogous to power counting by penalizing higher-complexity corrections. We emphasize that ADCD is *not* a full EFT framework: it has no renormalization-group flow, no controlled error expansion in E/Λ , and no gauge-invariance constraints on operator enumeration. The contribution of ADCD is rather to *automate* the correction-first philosophy of EFT for a class of data-driven problems—replacing the manual, insight-driven selection of candidate operators with a search over symbolic corrections screened by physics gates. In this sense ADCD is best read as a

data-driven cousin of EFT construction rather than as an instance of it.

3.1 Problem Formulation and Residual Extraction

Let $y_{\text{classical}} = g(X; \phi)$ be a known classical physical law with variables X and fixed constants ϕ . Let y_{obs} be the anomalous experimental data containing observational noise. We define the physical residual δ depending on the correction mode:

- **Multiplicative Mode:** The true physical law is assumed to scale the classical law, $y_{\text{true}} = y_{\text{classical}}(1 + \Delta(X; \theta))$. The target residual is:

$$\delta_{\text{mult}} = \frac{y_{\text{obs}}}{y_{\text{classical}}} - 1$$

- **Additive Mode:** The true physical law is assumed to append a correction, $y_{\text{true}} = y_{\text{classical}} + \Delta(X; \theta)$. The target residual is:

$$\delta_{\text{add}} = y_{\text{obs}} - y_{\text{classical}}$$

By default, ADCD automatically selects between additive and multiplicative correction modes by evaluating the relative variance and scale of both candidate formulations, choosing whichever best isolates the anomalous deviation from the observed residual (`correction_mode="auto"`).

3.2 Residual Feature Intelligence

To bias the proposer towards the correct mathematical families, ADCD detrends and extracts 5 high-level statistical features from the residual δ as a function of the primary limit variable x :

1. **Monotonicity:** Spearman rank correlation $r_s \in [-1, 1]$ between x and δ . A high absolute value indicates steady decay or growth.
2. **Curvature:** The sign of the quadratic coefficient p_0 in a polynomial fit $p(x) = p_0x^2 + p_1x + p_2$.
3. **Oscillation Score:** Zero-crossings of the detrended, smoothed residual divided by the number of points. High scores (> 0.15) represent trigonometric or wave-like anomalies.
4. **Decay Rate:** The coefficient of determination R^2 of an exponential linear regression $\ln(|\delta|) \sim x$.
5. **Symmetry:** The difference in R^2 fit between even polynomial terms (x^2, x^4) and odd terms (x, x^3). Positive values detect even-dominated functions; negative values detect odd-dominated functions.

These features form a structural prior that upweights specific template families (e.g. upweighting exponential templates if monotonicity and decay rates are high).

3.3 Cascaded Physics Gates (Stage 1)

To enforce physical laws on discovered candidates, proposed correction formulas $\Delta(X; \theta)$ are passed through three cascading screening filters:

1. **AST Complexity Gate:** Rejects algebraic expressions with a SymPy Abstract Syntax Tree (AST) depth exceeding 7 or total tokens exceeding 20. This acts as a coarse complexity regularizer.

2. **Dimensional Homogeneity and Transcendental Guardrail:** Verifies that the correction term Δ is strictly dimensionless. Physical variables $x \in X$ with units $[M^a L^b T^c]$ must enter the formula inside dimensionless ratios (e.g., v/c , r/θ_1 , or T/θ_0). As a sub-check, transcendental functions ($\sin, \cos, \exp, \log, \tanh$) can only accept mathematically dimensionless arguments; any expression containing a dimensioned argument (e.g., $\exp(-r)$ without a scaling length θ_1) is instantly blocked.
3. **Asymptotic Limit (ARC) Gate:** Assures that as the primary physical variable x approaches its classical boundary $x \rightarrow x_0$ (e.g., $v \rightarrow 0$ for relativistic kinetic energy, or $r \rightarrow \infty$ for Yukawa gravity), the correction term vanishes:

$$\lim_{x \rightarrow x_0} \Delta(X; \theta) = 0$$

This is evaluated using SymPy’s limit engine. Candidates that fail to vanish or diverge are discarded.

On the primary benchmark at 5% observational noise, ablation confirms that these gates function as a *collective* filter (Figure 5): disabling the ARC, AST, or Coarse-Eval gate individually leaves structural accuracy unchanged (8/9), while removing the dimensional gate alone reduces accuracy to 7/9, and disabling all gates simultaneously drops recovery to 4/9.

After gate filtering, a **Coarse Empirical Evaluation** pre-filter ranks surviving candidates by their preliminary NMSE on the observed residual, providing an early data-driven signal before the full JAX optimization in Stage 2.

3.4 JAX-Traced Parameter Optimization (Stage 2)

Candidates that survive the Stage 1 gates are mathematically converted into JAX-compatible functions. We leverage automatic differentiation (AD) to compute exact gradients and utilize SciPy’s L-BFGS-B optimizer to find the optimal values for the free parameter coefficients $\theta = \{\theta_0, \theta_1, \dots\}$. The optimization objective is the Normalized Mean Squared Error (NMSE) of the residual:

$$\text{NMSE} = \frac{\sum_{i=1}^N (\Delta(X_i; \theta) - \delta_i)^2}{\sum_{i=1}^N (\delta_i - \bar{\delta})^2}$$

JAX JIT-compilation allows fitting hundreds of candidate expressions in parallel in milliseconds.

3.5 Occam’s Razor Reranking (BIC)

Under observational noise, flexible wrong-class expressions (such as a power law with fitting exponent θ_2) can achieve numerically lower NMSE than the correct structural class. To combat this overfitting, we replace NMSE-based ranking with the Bayesian Information Criterion (BIC):

$$\text{BIC} = N \ln(\text{NMSE}) + k \ln(N)$$

where N is the number of experimental data points and k is the number of free parameters in θ . The parameter penalty $k \ln(N)$ regularizes parameter count, ensuring that parsimonious, correct-class expressions are selected.

3.6 Structural Class Classifier and Degenerate Exponent Detection

After BIC reranking, the winning expression is classified into one of six functional families (exponential, trigonometric, logarithmic, power law, rational, polynomial) by AST traversal of its SymPy parse tree. A key design choice governs power-law vs. polynomial classification: an expression of the form $\theta_0 \cdot x^{\theta_1}$ is normally classified as **power_law** because the exponent is a free parameter symbol. However, if the fitted value $\hat{\theta}_1$ satisfies $|\hat{\theta}_1 - 1| < \varepsilon$ (with $\varepsilon = 0.05$), the

expression is reclassified as **polynomial**, because a unit exponent makes it structurally linear in x —a degenerate power law that is physically a polynomial. This threshold is applied to the post-optimization fitted parameter, not to the symbolic expression alone, and is disclosed here because it affects structural classification for the Muon g -2 scenario: the discovered expression $\theta_0(\alpha/\pi)^{\hat{\theta}_1}$ with $\hat{\theta}_1 \approx 1.005$ is correctly classified as **polynomial**, matching the Schwinger leading-order linear correction $a_\mu = \alpha/(2\pi)$.

4 Experimental Setup

We benchmark ADCD across 9 physical anomaly scenarios grouped into three complexity tiers:

1. **Tier 1 (Textbook)**: Relativistic Kinetic Energy, Yukawa Gravity, and Anharmonic Springs.
2. **Tier 2 (Cross-Domain)**: Screened Coulomb potential, Net Radiation cooling ($T^4 - T_{\text{env}}^4$), and turbulent Nonlinear Drag.
3. **Tier 3 (Synthetic/Novel)**: Mystery-A (sub-wavelength gravity containing a $\tanh(r_0/r)^2$ correction), Mystery-B (quantum mechanics containing a $\text{sinc}(v/v_0) - 1$ correction), and Mystery-C (non-linear polymer stretching containing a $\ln(1 + x/x_0)/(x/x_0) - 1$ correction).

Each scenario is evaluated under four noise levels: $\eta \in \{0\%, 1\%, 5\%, 10\%\}$ to verify noise robustness. We compare three proposer configurations:

- **Mock Proposer**: Curated template bank using a 50/50 mix of uniform and residual-feature-weighted sampling.
- **Gemini Proposer**: Generative Google Gemini 2.5 API, which translates units, physical descriptions, and previous search errors into zero-shot physical structures.
- **Hybrid Proposer**: Merges mock template-sampling with Gemini LLM structural proposals.

Reproducibility and seed disclosure. All benchmark results are reported across sixteen independent random seeds (0, 2, 5, 7, 13, 17, 21, 23, 31, 37, 41, 42, 53, 61, 67, 99). The **mean structural recovery rate is 80.4%** ($\pm 7.4\%$ population standard deviation; 95% bootstrap CI [76.7%, 84.0%] over 10,000 resamples of the per-seed rates); the full per-seed breakdown is given in Table 1. We disclose explicitly that seed=42, which we use as the *reference seed* throughout this paper, **is the highest-performing seed** in this set. Seed=42 was adopted as the reference seed prior to running the multi-seed study (it was the default starting seed in the initial implementation); the multi-seed analysis was conducted afterward as a robustness check and revealed, post-hoc, that it is also the best-performing seed. We report 94.4% (34/36) at seed=42 alongside the mean of 80.4% for full transparency; the primary claim of the paper is the mean, not the peak. The 25.0 percentage-point spread (69.4%–94.4%) reflects stochastic template sampling in the MockProposer: different seeds generate different candidate pools. “Deterministic” in our context means *reproducible given the same seed*, not identical across seeds. Physics gates ensure that when the correct functional family is sampled, it consistently survives gate filtering and BIC reranking regardless of seed.

Table 1: Structural recovery rate per random seed and aggregate statistics (Mock Proposer, 9 scenarios $\times$ 4 noise levels = 36 trials per seed). The 95% confidence interval is a bootstrap CI (10,000 resamples of the per-seed rates).

Seed	Recovery rate
0	86.1%
2	88.9%
5	83.3%
7	75.0%
13	69.4%
17	83.3%
21	77.8%
23	80.6%
31	69.4%
37	83.3%
41	91.7%
42	94.4%
53	75.0%
61	69.4%
67	77.8%
99	80.6%
Mean $\pm$ std	80.4% $\pm$ 7.4%
95% bootstrap CI	[76.7%, 84.0%]
Min / Max	69.4% / 94.4%

Template overlap disclosure. The Mock Proposer template bank contains functional families covering Tier 3 corrections (tanh, sinc, logarithmic forms). Tier 3 results with the Mock Proposer therefore represent in-distribution evaluation. Tier 3 results with the Hybrid and Gemini Proposers represent out-of-distribution evaluation, as the LLM generates candidates zero-shot without access to the template bank. For the real-world tier, extended-mode templates add scenario-specific hints (e.g. $\theta_0 \cdot (v/c)^2$ for Mercury); only Mercury has an *exact* ground-truth template in this bank (Table 10). All other real-world scenarios are evaluated against generic functional families (power law, exponential, polynomial), not verbatim answer injection.

5 Experimental Results

Evaluation regimes. We report results under three explicitly separated evaluation regimes to avoid conflating template-assisted recovery with zero-shot discovery: (i) **Primary (synthetic benchmark)**: Mock Proposer, sixteen independent seeds (0, 2, 5, 7, 13, 17, 21, 23, 31, 37, 41, 42, 53, 61, 67, 99), $9 \times 4 = 36$ trials per seed—the **primary claim** is mean structural recovery 80.4% ($\pm 7.4\%$); (ii) **Supplementary (zero-shot upper bound)**: Hybrid Proposer at reference seed=42 only (33/36 = 91.7%), reporting LLM-augmented discovery beyond the template bank; (iii) **Real-world (physics-constant hybrid)**: Mock Proposer on four established scenarios using CODATA/NIST/JPL parameters—4/4 structural class match, with quantitative recovery (NMSE $< 10^{-4}$) on 3/4 (Blackbody limited to structural identification; see Table 9). Binary pulsar decay is reported separately as a controlled sensitivity study (Table 11), not in the main recovery headline.

We distinguish two performance regimes within the primary and supplementary synthetic benchmarks. First, **template-assisted recovery** (Mock Proposer): ADCD achieves a mean of **80.4%** ($\pm 7.4\%$) structural class recovery across sixteen seeds, with a peak of 94.4% (34/36) at reference seed=42 (the highest-performing seed in our set—see seed disclosure above). Second, **genuine discovery capability** (Hybrid Proposer): ADCD achieves 91.7% (33/36) struc-

tural recovery at seed=42 using zero-shot LLM-generated candidates merged with templates—demonstrating that the correction-first paradigm generalizes beyond pre-specified template families (reproduce via `run_correction_discovery.py -proposer hybrid`; frozen results in `hybrid_seed42_results.json`). Both seed=42 figures are provided as concrete upper-bound reference points; the mean of 80.4% is the primary performance claim.

The full benchmark results for the Mock Proposer are summarized in Table 2, Table 3, and Table 4.

Table 2: Tier 1 (Textbook) Benchmark Results using the Mock Proposer (seed=42).

Scenario	Noise	Discovered Correction $\Delta(x; \theta)$	Full NMSE	Class Match	AST Dist
Relativistic KE	0%	$\theta_0 v^2$	3.14×10^{-32}	YES	3
	1%	$\theta_0(v/c)^2 + \theta_1(v/c)^4$	1.69×10^{-04}	YES	6
	5%	$\theta_0(v/c)^2 + \theta_1(v/c)^4$	4.19×10^{-03}	YES	6
	10%	$\theta_0(v/c)^2 + \theta_1(v/c)^4$	1.65×10^{-02}	YES	6
Yukawa Gravity	0%	$\theta_0 e^{-r/\theta_1}$	4.00×10^{-16}	YES	0
	1%	$\theta_0 e^{-r/\theta_1}$	7.12×10^{-05}	YES	0
	5%	$\theta_0 e^{-r/\theta_1}$	1.83×10^{-03}	YES	0
	10%	$\theta_0 e^{-r/\theta_1}$	7.55×10^{-03}	YES	0
Anharmonic Spring	0%	$\theta_0 x^4$	0.00	YES	0
	1%	$\theta_0 x^4$	1.79×10^{-04}	YES	0
	5%	$\theta_0 x^4$	4.46×10^{-03}	YES	0
	10%	$\theta_0 x^4$	1.76×10^{-02}	YES	0

Table 3: Tier 2 (Cross-Domain) Benchmark Results using the Mock Proposer (seed=42).

Scenario	Noise	Discovered Correction $\Delta(x; \theta)$	Full NMSE	Class Match	AST Dist
Screened Coulomb	0%	$e^{-r/\theta_1} - 1$	5.41×10^{-19}	YES	0
	1%	$e^{-r/\theta_1} - 1$	5.04×10^{-05}	YES	0
	5%	$r\theta_0/(r + \theta_1)$	1.09×10^{-03}	NO	6
	10%	$r\theta_0/(r + \theta_1)$	4.43×10^{-03}	NO	6
Net Radiation	0%	$-(\theta_0/T)^{\theta_1}$	1.52×10^{-13}	YES	4
	1%	$-(\theta_0/T)^{\theta_1}$	1.51×10^{-04}	YES	4
	5%	$-(\theta_0/T)^{\theta_1}$	3.74×10^{-03}	YES	4
	10%	$-(\theta_0/T)^{\theta_1}$	1.47×10^{-02}	YES	4
Nonlinear Drag	0%	$\theta_0 v^2$	0.00	YES	0
	1%	$\theta_0(v/\theta_1)^2$	2.64×10^{-04}	YES	3
	5%	$\theta_0(v/\theta_1)^2$	6.59×10^{-03}	YES	3
	10%	$\theta_0(v/\theta_1)^2$	2.60×10^{-02}	YES	3

Table 4: Tier 3 (Synthetic/Novel) Benchmark Results using the Mock Proposer (seed=42).

Scenario	Noise	Discovered Correction $\Delta(x; \theta)$	Full NMSE	Class Match	AST Dist
Mystery-A	0%	$-\theta_0 \tanh^2(\theta_1/r)$	2.99×10^{-13}	YES	1
	1%	$-\theta_0 \tanh^2(\theta_1/r)$	2.43×10^{-04}	YES	1
	5%	$-\theta_0 \tanh^2(\theta_1/r)$	6.28×10^{-03}	YES	1
	10%	$-\theta_0 \tanh^2(\theta_1/r)$	2.59×10^{-02}	YES	1
Mystery-B	0%	$\frac{\sin(v/\theta_1)}{v/\theta_1} - 1$	2.49×10^{-11}	YES	0
	1%	$\frac{\sin(v/\theta_1)}{v/\theta_1} - 1$	2.07×10^{-04}	YES	0
	5%	$\frac{\sin(v/\theta_1)}{v/\theta_1} - 1$	5.18×10^{-03}	YES	0
	10%	$\frac{\sin(v/\theta_1)}{v/\theta_1} - 1$	2.05×10^{-02}	YES	0
Mystery-C	0%	$\frac{\ln(1+x/\theta_1)}{x/\theta_1} - 1$	6.08×10^{-13}	YES	0
	1%	$\frac{\ln(1+x/\theta_1)}{x/\theta_1} - 1$	2.79×10^{-04}	YES	0
	5%	$\frac{\ln(1+x/\theta_1)}{x/\theta_1} - 1$	6.99×10^{-03}	YES	0
	10%	$\frac{\ln(1+x/\theta_1)}{x/\theta_1} - 1$	2.76×10^{-02}	YES	0

5.1 Performance Targets Analysis

- **Textbook Tier Recovery:** The system recovered the correct structural class for all three textbook scenarios across all noise levels. Relativistic KE, Yukawa Gravity, and Anharmonic Spring all achieved **100% structural class recovery**. The Yukawa exponential correction $\theta_0 e^{-r/\theta_1}$ was consistently identified, demonstrating that the physics gates successfully constrain the search space even under 10% observational noise.
- **Cross-Domain Recovery:** Net Radiation achieved perfect structural class recovery across all noise levels, correctly retaining the inverse fourth-power power law $\Delta \propto T^{-4}$. Nonlinear Drag was also recovered at all noise levels. Screened Coulomb was successfully recovered at 0% and 1% noise, but at higher noise (5% and 10%), the exponential correction was fitted by a rational proxy $r\theta_0/(r + \theta_1)$, reflecting the difficulty of distinguishing exponential decay from rational saturation under limited dynamic range.
- **Synthetic / Novel Recovery:** The system scored **100% (12/12) structural match rate** on the synthetic tier. Highly complicated novel formulas, such as the sinc quantum boundary fluctuation and the logarithmic quotient stretching, were fully converged and structurally verified, demonstrating that statistical priors and LLM reasoning combine effectively to cover complex functional families.

Figures 2 to 4 visualise these outcomes: the per-scenario noise–accuracy curve against PySR (Figure 2), the log-scale NMSE heatmap with structural-class verdicts across all scenarios and noise levels (Figure 3), and the tier-aggregated match rates that localise degradation to the Screened Coulomb scenario at $\geq 5\%$ noise (Figure 4).




Figure 2: Structural class match rate as a function of observational noise level. ADCD (Mock Proposer, seed=42) maintains 88.9% structural accuracy at 5% noise due to physics-constrained gate filtering. PySR (fair profile, unconstrained SR on the same residual) exhibits a non-monotonic noise-accuracy profile (44.4%→55.6%→11.1%→55.6%), reflecting overfitting instabilities absent from physics-gated search.




Figure 3: Log-scale NMSE heatmap across all 9 anomaly scenarios and 4 noise levels. Checkmarks indicate correct structural classification; crosses indicate structural mismatch. White lines separate tiers.




Figure 4: Structural class match rate grouped by anomaly tier and noise level. The synthetic tier achieves 100% recovery at all noise levels; degradation is localized to the Screened Coulomb scenario at $\geq 5\%$ noise.




Figure 5: Ablation study at 5% observational noise. Each bar disables one gate. Removing the ARC, AST, or Coarse-Eval gate individually leaves structural accuracy unchanged (8/9). Removing the dimensional gate alone reduces accuracy by 11.1 pp (8/9 to 7/9). Only removing *all* gates simultaneously causes a larger collapse, dropping performance by $\Delta = 44.5$ percentage points (from 8/9 to 4/9), confirming that the gates reinforce each other as a *collective* physical consistency filter. Note: the “Without BIC Reranking” condition is described in Section 6 but excluded from this figure as it requires the full gate pipeline to isolate BIC’s contribution.

5.2 Baseline Comparison

PySR Comparison. We compare ADCD against PySR [2] on the *same residual targets* $y_{\text{obs}} - y_{\text{classical}}$. The primary baseline uses PySR’s *fair* profile ($n_{\text{iter}} = 100$, $\text{maxsize}=30$, 60 s timeout per scenario; Table 5). A legacy *fast* profile (15 iterations, 25 s) is retained for wall-clock-matched comparison only. As shown in Table 6, ADCD achieves 88.9% structural recovery at 5% noise versus 11.1% for PySR *fair*—a 77.8 percentage-point gap under equal residual data and a substantially larger search budget for PySR. At 0% noise, ADCD reaches 100% (9/9) versus 44.4% (4/9) for PySR *fair*, indicating that physics-gated correction search recovers interpretable structure in the correction-to-a-known-law regime, where unconstrained tabulara search faces a harder task.

Non-monotonic behaviour of PySR. PySR *fair* exhibits a non-monotonic noise–accuracy profile (4/9 at 0%, 5/9 at 1%, 1/9 at 5%, 5/9 at 10%), which warrants explanation. Inspection of per-scenario mismatch distributions (extracted from `pysr_baseline_fair.json`) reveals a mechanistic reason: at 0% noise, the large search space ($\text{maxsize}=30$) allows PySR to find deeply-nested trigonometric expressions (average complexity 13.8 nodes) that achieve very low MSE on clean data but are structurally incorrect (5 of 5 mismatches discovered *trigonometric*). At 1% noise, noise regularises overfitting and average complexity drops to 8.2 nodes, yielding more parsimonious, correct-class results. At 5% noise, PySR discovers moderate-complexity expressions (average 6.2 nodes) but selects an incorrect functional family: of 8 mismatches, 3 are polynomial/trigonometric scenarios misclassified as *exponential* and 3 others as *trigonometric*, suggesting that at intermediate SNR, unconstrained search struggles to distinguish functional classes whose shapes become similar under noise. At 10% noise, expressions again simplify (average 6.4 nodes) and PySR returns the simplest available form, which by coincidence matches

the correct polynomial or power-law class for 5 scenarios. We emphasise that this behaviour is consistent with PySR’s design goal—unconstrained tabula rasa discovery—rather than a defect; ADCD avoids it not by improving the search but by operating in the correction-to-a-known-law regime, where physics gates are applied independently of noise level and, once the correct functional family is sampled, ARC and dimensional constraints consistently select it regardless of SNR.

Upper-bound ablation: scaling PySR’s budget. A natural concern is that the gap reflects an under-powered PySR. We therefore rerun PySR with a **generous** profile that *doubles* the search budget on every axis: $n_{\text{iter}} = 200$ (vs. 100), $\text{maxsize}=40$ (vs. 30), and a 120 s timeout (vs. 60 s) per scenario, at $3\times$ the mean wall-clock (26.4 s vs. 8.8 s per run; Table 8). As shown in Table 6, the total structural recovery stays flat at 15/36 (41.7%)—identical to **fair**—and the noise profile remains non-monotonic (4/9, 4/9, 5/9, 2/9). At the 5% noise headline point, generous actually recovers *fewer* of the targeted functional families than a fair reading would suggest is stable (5/9, 55.6%). The additional budget lets PySR explore deeper expression trees (mean complexity 9.8 nodes vs. 8.7 for fair) and a larger hall-of-fame (~ 18 vs. ~ 14 retained equations), but this extra capacity is spent on more deeply-overfit rather than more physically-correct structure. We interpret this as evidence that the ADCD–PySR gap is not a budget artefact: unconstrained evolutionary search cannot reliably convert a larger candidate pool into correct functional *classes* under noise, because the structural-selection problem (not the search budget) is the binding constraint. Physics gates resolve this selection problem directly.

Table 5: PySR baseline configuration profiles. Primary comparison uses **fair** (near-default search budget). **fast** is legacy wall-clock-matched; **generous** is upper-bound ablation.

Profile	Iter.	Max Size	Timeout (s)	Pop.	Size	Seed	
fast (legacy)	15	15	25	8	20	42	<i>Shared: binary ops $\{+, -, \times, \div\}$; unary ops $\{\exp, \sin, \cos, \log, \tanh, \sqrt{}\}$; parsimony = 0.0032; deterministic = True; same residual as ADCD.</i>
fair (primary)	100	30	60	8	20	42	
generous	200	40	120	8	20	42	

Table 6: Structural class match rate: ADCD (Mock Proposer, seed=42) vs. PySR on the same 9-scenario residual benchmark. PySR **fair** is the primary unconstrained baseline; **generous** ($2\times$ iterations, $2\times$ timeout, larger **maxsize**) is an upper-bound ablation showing that additional search budget does not close the gap.

Method	0%	1%	5%	10%
ADCD (ours)	9/9 (100.0%)	9/9 (100.0%)	8/9 (88.9%)	8/9 (88.9%)
PySR fair	4/9 (44.4%)	5/9 (55.6%)	1/9 (11.1%)	5/9 (55.6%)
PySR generous (upper bound)	4/9 (44.4%)	4/9 (44.4%)	5/9 (55.6%)	2/9 (22.2%)
PySR fast (legacy)	2/9 (22.2%)	6/9 (66.7%)	4/9 (44.4%)	4/9 (44.4%)

Table 7: ADCD search-space funnel (aggregate over 9 scenarios $\times$ 4 noise levels, Mock Proposer seed=42). Telemetry from `gate_telemetry.json`.

Stage	Candidates	Survival (%)
Proposed (Stage 0b)	3,225	100.0%
After parse	3,225	100.0%
After AST	3,225	100.0%
After dimensional	2,888	89.6%
After transcendental	2,888	89.6%
After ARC (Stage 1 out)	655	20.3%
JAX optimized (Stage 2)	598	18.5%

Table 8: Compute budget comparison (9 scenarios $\times$ 4 noise levels).

Method	Mean wall-clock / run	Expressions explored
ADCD (Mock)	4.2 s	598 JAX optim. calls
PySR (fair)	8.8 s	~ 14 hall-of-fame / run
PySR (generous)	26.4 s	~ 18 hall-of-fame / run

MLP Comparison. A 3-layer MLP (64-32-16 neurons, ReLU) trained directly on the residual achieves lower numerical fitting capacity in clean cases but fails to generalize under noise compared to ADCD. Specifically, ADCD achieves an average NMSE of 2.88×10^{-12} at 0% noise and 4.48×10^{-3} at 5% noise (reference seed; `reproducibility_results.json`), whereas the MLP achieves 8.56×10^{-5} at 0% noise and 8.01×10^{-3} at 5% noise (`mlp_baseline_results.json`). MLP, however, produces no interpretable expression—only a black-box numerical function. ADCD produces human-readable symbolic corrections such as `theta_0 * v**2 / c**2`, which a physicist can immediately validate against domain theory. This interpretability gap is the core contribution of our framework.

5.3 Blind Generalization Test

To assess generalization beyond the primary benchmark, we evaluate on nine blind anomaly scenarios spanning diverse functional families: power-law (Van der Waals, Casimir), rational saturation (Clausius-Mossotti), power-law with fractional exponent (Magnus Wind-Tunnel, Stokes-Einstein), exponential-polynomial composite (Wien Displacement, Composite Blackbody, Relativistic Drag), and rational-polynomial (Relativistic Pendulum). Ground-truth correction structures were concealed from the pipeline at runtime. We apply the strict criterion: a scenario counts as a success only if the correct functional class is recovered across *all* four tested noise levels (0%, 1%, 5%, 10%).

Mock Proposer (template-based). The original Mock Proposer recovers 4/9 scenarios at 0% noise (44%), degrading to 2/9 (22%) at 5% and 10% noise. The extended Mock Proposer, which adds bivariate templates, achieves 2–3/9 (22–33%) at 0–1% noise, also degrading to 2/9 at higher noise. Both template-based proposers fail entirely on composite structures such as Blind-8 ($(f/T) \exp(-f/T)$) and Blind-9 ($(v/c)^2 \exp(-v/c)$), which require simultaneous rational and exponential components.

Grammar Proposer (Phase 1 contribution). The Grammar Proposer, which enumerates candidates from a formal grammar over dimensionless ratios, achieves **3/9 (33%) at 0–1% noise**, matching Mock (base) at clean data. However, Grammar performance degrades sharply

to **0/9 (0%) at 5% and 10% noise**, performing worse than the template-based proposers under realistic noise conditions.

Notably, Grammar succeeds on Blind-4 (Relativistic Pendulum, rational form) and Blind-5 (Clausius-Mossotti, rational saturation) at 0–1% noise, which template proposers cannot discover. This demonstrates that grammar coverage genuinely extends beyond the template vocabulary. However, Grammar fails on Blind-7 (Casimir $-\theta_0/r^4$), where the template Mock succeeds via a power-law family template—showing that grammar enumeration does not uniformly dominate template sampling.

Interpretation. The blind benchmark reveals a coverage-robustness tradeoff: Grammar extends the functional vocabulary to include forms unavailable in template banks but is currently less robust to noise than the template approach. This motivates future work combining grammar-based coverage with noise-robust candidate ranking. Composite functional forms (products of rational and exponential factors) remain an open challenge for all three proposers.

5.4 Real-World Physical Constants Benchmark

To evaluate ADCD on datasets governed by physical constants and realistic scales, we constructed five **synthetic-real hybrid** scenarios using JPL DE440, NIST, and CODATA parameters. We report three distinct success criteria to avoid conflating structural identification with full discovery:

1. **Structural class recovery:** discovered and expected correction belong to the same functional family (power law, exponential, etc.).
2. **Quantitative recovery:** full-reconstruction NMSE $< 10^{-4}$.
3. **Optimizer convergence:** residual NMSE $< 10^{-5}$ after validated post-hoc evaluation.

At reference seed = 42 with the Mock Proposer (extended mode), ADCD achieves **4/4 structural**, **3/4 quantitative**, and **2/4 optimizer-converged** results on the established scenarios (Table 9). Binary pulsar orbital decay is reported separately as a sensitivity study (Table 11) because its benchmark uses a reduced single-variable formulation that is intentionally easier than the full multi-parameter Peters formula.

Scenario highlights.

- **Mercury Perihelion (GR):** Recovers $\Delta = \theta_0(v/c)^2$ with NMSE = 1.11×10^{-5} (exact AST match), the leading Schwarzschild perihelion scaling.
- **Hydrogen Lamb Shift (QED):** Recovers $\Delta \propto n^{-3}$ in equivalent parameterized form (NMSE = 1.69×10^{-18}).
- **Muon $g - 2$ (Schwinger):** Discovers $\Delta = \theta_0(\alpha/\pi)^{\theta_1}$ with $\theta_1 \approx 1$ (NMSE = 7.94×10^{-7}).
- **Blackbody (Planck):** Identifies exponential UV suppression ($\Delta = -1 + e^{-f/\theta_1}$; NMSE = 2.59×10^{-2}) via `loss_mode='auto'` full-reconstruction loss; quantitative fit remains limited by Rayleigh-Jeans dynamic range.
- **Binary Pulsar Decay (GR/GW) — Sensitivity Study:** Under a **reduced-variable** Hulse-Taylor formulation (fixed M , a , e ; only P varies), ADCD recovers a power-law $\Delta = \theta_0/P^{\theta_1}$ with $\theta_1 \approx 1.57$ (true 5/3; NMSE = 3.94×10^{-3}). This result is **not counted in the 4/4 headline** because the single-variable formulation is substantially simpler than the full Peters formula. The sensitivity study in Table 11 confirms that recovery degrades monotonically as more orbital parameters are exposed, and fails structurally under the full four-parameter scan.

Table 9: Real-World Structural and Quantitative Recovery (Mock Proposer, seed=42). Quantitative success: $\text{NMSE} < 10^{-4}$. Asterisk: structural match only.

Scenario	True Δ	Recovered Δ	Class	NMSE	Notes
Mercury Perihelion	$\theta_0(v/c)^2$	$\theta_0(v/c)^2$	Yes	1.11×10^{-5}	exact AST match
Hydrogen Lamb Shift	θ_0/n^3	$\theta_0(n/\theta_1)^{-\theta_2}$	Yes	1.69×10^{-18}	struct. class only
Blackbody Radiation	$e^{-f\theta_0/T}$	$-1.0 + e^{-f/\theta_1}$	Yes	$(2.59 \times 10^{-2})^*$	AST dist = 4
Muon g-2	$\alpha\theta_0/\pi$	$\theta_0(\alpha/\pi)^{\theta_1}$	Yes	7.94×10^{-7}	AST dist = 4

Template leakage analysis. Table 10 documents whether each ground-truth correction appears verbatim in the extended Mock Proposer bank. Only Mercury has exact template overlap; binary pulsar uses a generic $\theta_0 \cdot P^{-\theta_1}$ family, not the literal $P^{-5/3}$ exponent.

Table 10: Template Leakage Disclosure for Real-World Benchmark (Mock Proposer extended mode). “Exact” = ground-truth form appears verbatim in the template bank.

Scenario	Ground Truth	Exact?	Family?	Expected Class
Mercury Perihelion	$\theta_0(v/c)^2$	Yes	Yes	polynomial
Hydrogen Lamb Shift	θ_0/n^3	No	Yes	power_law
Blackbody Radiation	$\exp(-hf/k_B T)$	No	Yes	exponential
Muon g-2	$\theta_0\alpha/\pi$	No	Yes	polynomial
Binary Pulsar Decay	$\theta_0 P^{-5/3}$	No	Yes	power_law

Binary pulsar sensitivity. To address whether the v2.1 benchmark was simplified, we re-ran discovery as free variables increase (Table 11). Structural recovery holds for P_only, P_e, and P_e_M, but NMSE degrades from 3.9×10^{-3} to $\mathcal{O}(0.3\text{--}0.5)$. The legacy full four-parameter scan fails structurally (logarithmic false positive). This confirms the reduced formulation is genuinely easier—and that ADCD does not silently succeed on the hardest variant.

Table 11: Binary Pulsar Sensitivity: structural recovery vs. formulation difficulty. v2.1 default uses the P_only reduced-variable form.

Variant	Free Vars.	Discovered Δ	Class	NMSE
P_only	P	$\theta_0 P^{-\theta_1}$	Yes	3.94×10^{-3}
P_e	P, e	$\theta_0 e^3/P$	Yes	2.52×10^{-1}
P_e_M	P, e, M	$\theta_0 P^{-\theta_1}$	Yes	4.81×10^{-1}
full	P, M, a, e	$\theta_0 \log(\theta_1 + e/P)$	No	1.57×10^{-2}

5.5 System Misspecification and Fail-Safe Behavior

A core design requirement for automated scientific discovery is that the discovery pipeline must not yield confident false physical corrections when the underlying baseline theory is structurally wrong (misspecified). To test this, we evaluated ADCD under two distinct tiers of misspecification:

Tier 1: Extreme Misspecifications (Noise = 1.0%). Using the Relativistic KE anomaly as the ground truth, we tested three corrupted baselines:

- *Wrong Baseline Form (Case A):* Providing momentum ($m \cdot v$) as the baseline. This led to convergence failure (NMSE 9.61×10^{-2}) and a high BIC of -363.46 (compared to -1019.16 for the correct baseline), with a mismatched structural class.

- *Missing Vital Variable (Case B)*: Omitting the velocity v from the available variables. This resulted in complete rejection of the candidates by the pipeline gates (Full NMSE 1.00 and a penalization BIC of 9999.00).
- *Spurious Irrelevant Variable (Case C)*: Adding an irrelevant temperature variable T to the spring baseline. The system successfully filtered out the spurious variable and recovered the correct spring correction (x^4), with only a minimal BIC penalty (-655.24 compared to -1019.16 of correct baseline).

Tier 2: Subtle Misspecifications (Noise = 5.0%). Using the Nonlinear Drag anomaly as the ground truth, we tested three corrupted baselines:

- *Constant 5% Off (Case D)*: Setting the baseline coefficient 5% too high. The pipeline refused to match the polynomial drag class and converged to a power-law proxy with a worse BIC (-440.41 vs. -507.26 for the correct baseline).
- *Constant Offset (Case E)*: Adding a small constant offset of 0.1. Because the offset is negligible compared to the large quadratic drag signal, the pipeline robustly recovered the correct polynomial class with a very close BIC (-504.04).
- *Wrong Structural Coefficient (Case F)*: Setting the linear drag coefficient 20% too low. The system failed to find the correct drag class, converging to a power-law class with a high BIC of -412.95 .

These results (summarized in Table 12) confirm that ADCD displays reliable **fail-safe behavior**: rather than producing confident spurious corrections, ADCD fails to converge when the baseline theory is structurally wrong—preventing false physical discovery. Minor perturbations such as small constant offsets are absorbed without affecting structural recovery, demonstrating robustness to measurement calibration errors.

Table 12: System Misspecification Benchmark Results.

Scenario	Converged [†]	Full NMSE	BIC	Class Match	Status
<i>Tier 1: Extreme Misspecifications (Noise = 1.0%)</i>					
Relativistic KE (Correct Baseline)	False	1.70×10^{-4}	-1019.16	YES	Positive Control
Case A: Wrong Baseline Form	False	9.61×10^{-2}	-363.46	NO	Fail-safe
Case B: Missing Vital Variable	False	1.00×10^0	9999.00	NO	Fail-safe
Case C: Spurious Variable	False	1.96×10^{-4}	-655.24	YES	Robust
<i>Tier 2: Subtle Misspecifications (Noise = 5.0%)</i>					
Nonlinear Drag (Correct Baseline)	False	6.59×10^{-3}	-507.26	YES	Positive Control
Case D: Constant 5% Off	False	7.46×10^{-3}	-440.41	NO	Fail-safe
Case E: Constant Offset (0.1)	False	6.70×10^{-3}	-504.04	YES	Robust
Case F: Wrong Structural Coeff. (20% off)	False	1.96×10^{-2}	-412.95	NO	Fail-safe

[†] “Converged” = residual NMSE $< 10^{-5}$ (Section 5.4). Positive controls achieve Class Match= YES despite Converged=False because the polynomial class is structurally identified before the convergence threshold is met; the correction magnitude at 1% noise does not require sub- 10^{-5} residual.

5.6 Real Observational Discovery: SPARC Radial Acceleration Relation

The benchmarks above validate ADCD on synthetic anomalies anchored to authoritative physical constants. As a more demanding capstone test, we applied ADCD to *real observational data*: the SPARC sample of galactic rotation curves [18], a canonical dataset in the modified-gravity / dark-matter debate. We use the mass models of Lelli et al. [18] ($N_{\text{gal}} = 171$ galaxies, $N_{\text{pts}} = 3342$ radial points; four galaxies from the original sample of 175 were excluded due to incomplete mass models or insufficient radial data points).

Setup. We work in the *radial acceleration plane*. For each radial point we compute the baryonic acceleration g_{bar} (predicted from the observed baryonic mass distribution) and the total acceleration inferred from the measured rotation speed, $g_{\text{obs}} = V_{\text{obs}}^2/R$. The dimensionless ratio

$$\nu(x) \equiv (V_{\text{obs}}/V_{\text{bar}})^2 = g_{\text{obs}}/g_{\text{bar}}, \quad x \equiv g_{\text{bar}}/a_0,$$

quantifies the departure from Newtonian dynamics, where V_{bar} is the baryonic rotation speed and $a_0 = 1.2 \times 10^{-10} \text{ m s}^{-2}$ is the MOND acceleration scale [19]. In the Newtonian regime ($g_{\text{bar}} \gg a_0$, i.e. $x \rightarrow \infty$) the correction must vanish, $\nu \rightarrow 1$; the anomaly grows as $x \rightarrow 0$. This is precisely a *correction-first* problem: the classical baseline is Newtonian ($\nu_{\text{cls}} = 1$), and we seek the dimensionless correction $\Delta(x) = \nu(x) - 1$ that reconciles the baryonic prediction with the observed kinematics.

Discovery. We stacked all SPARC radial points into a single (x, ν_{obs}) dataset and ran iterative ADCD (Section 3) with the asymptotic limit $\Delta \rightarrow 0$ as $x \rightarrow \infty$ enforced by the ARC gate. ADCD discovered the two-parameter interpolating function

$$\boxed{\nu_{\text{ADCD}}(x) = \theta_0 \left(\sqrt{1 + \theta_1/x} - 1 \right) + 1} \quad (\hat{\theta}_0 \approx 2.24, \hat{\theta}_1 \approx 0.27).$$

This form is interpretable: it reduces to $\nu \rightarrow 1$ in the Newtonian limit ($x \rightarrow \infty$) and grows like $\theta_0 \sqrt{\theta_1/x}$ in the deep-MOND regime ($x \rightarrow 0$), matching the expected $g_{\text{obs}} \propto \sqrt{g_{\text{bar}}}$ scaling [20].

Model comparison. Table 13 compares the ADCD-discovered function against three established closed-form interpolating functions for the radial acceleration relation: Simple MOND, Standard MOND, and the empirical RAR form of McGaugh, Lelli & Schombert [20]. The ADCD function achieves the lowest stacked-data NMSE (**0.373**) and the lowest BIC (−3280.7), reducing the NMSE by **41%** relative to the McGaugh+16 RAR form (0.636). The absolute NMSE reflects intrinsic galaxy-to-galaxy scatter in the stacked relation rather than model misfit; the relevant comparison is *relative*, where ADCD’s 2-parameter function is the most parsimonious competitive fit among the four. Figure 6 plots the stacked data against all four interpolating functions and their residuals.

Table 13: SPARC stacked radial acceleration relation ($N_{\text{gal}} = 171$, $N_{\text{pts}} = 3342$). NMSE evaluated on the stacked (x, ν_{obs}) dataset with $x \equiv g_{\text{bar}}/a_0$. The ADCD-discovered 2-parameter function achieves the lowest NMSE and the lowest BIC.

Model	NMSE ↓	BIC ↓	Free params
ADCD discovered	0.3729	−3280.69	2
RAR (McGaugh)	0.6361	−1512.15	0
Simple MOND	0.6477	−1451.35	0
Standard MOND	0.7074	−1156.82	0




Figure 6: SPARC stacked radial acceleration relation. **Upper:** observed $\nu(x)$ (grey) against the four interpolating functions. The ADCD-discovered 2-parameter curve (red) tracks the data more tightly than Simple/Standard MOND and the McGaugh+16 RAR form. **Lower:** residuals $\nu_{\text{obs}} - \nu_{\text{model}}$.

Statistical validation. Because the stacked dataset mixes points from many galaxies, in-sample NMSE can overstate generalization, and a single summary statistic can hide which comparison is being made. We therefore report three nested levels of evidence, all evaluated at the *galaxy* level so that no galaxy contributes to both train and test sets. The three levels separate a comparison ADCD wins decisively from one it merely matches, and we report both rather than conflate them.

(i) *Stacked-data comparison (in-sample, vs. fixed canonical forms).* Table 13 reports the stacked-data fit. The ADCD-discovered function achieves the lowest NMSE (**0.373**) and the lowest BIC (-3280.7), a 41% NMSE reduction relative to the McGaugh+16 RAR form (0.636). This establishes that ADCD’s discovered function dominates the *zero-parameter* canonical forms—i.e. the closed-form interpolating functions as they appear in the literature, with no free parameters refit to these data. The scientifically fairer question, however, is whether the *family* ADCD discovered is competitive once the canonical forms are granted the same two free parameters; that is the out-of-sample comparison below.

(ii) *Galaxy-level cross-validation (out-of-sample, fair parameter-matched comparison).* Ta-

ble 14 reports 10 repeated random 50/50 galaxy-level splits. ADCD’s discovered function attains out-of-sample test NMSE 0.386 ± 0.016 . The table deliberately separates two regimes: the zero-parameter canonical forms (top three rows, NMSE 0.66–0.70) and the same MOND/RAR families *refit with two free parameters* (bottom three rows). Against the zero-parameter forms ADCD wins by a wide margin. Against the two-parameter refit Simple-MOND and McGaugh RAR forms, however, ADCD is *statistically indistinguishable* ($\Delta\text{NMSE} = +0.008$ relative to Simple MOND 2-param, combined SE 0.021, $z = +0.39$, $p = 0.69$), and it *decisively outperforms* the two-parameter Standard-MOND form ($\Delta\text{NMSE} = -0.081$, $z = -4.60$). The in-sample counterpart of this fair comparison is shown in Table 15: under matched parameter count the fitted Simple-MOND and RAR forms attain marginally lower stacked NMSE than ADCD by $\sim 1.5\%$, well within the bootstrap CI. Cross-validation thus resolves the comparison into a clear hierarchy: ADCD matches the best-performing domain-expert interpolating families and beats the simpler one.

(iii) *Cluster bootstrap on δBIC (decision-theoretic, at N_{eff})*. Because BIC differences are the conventional currency of model selection, and because galaxies contribute correlated groups of radial points, we estimated a galaxy-level cluster bootstrap (1,000 resamples, seed 2026) on the effective-sample-size-scaled BIC difference $\delta\text{BIC}_{\text{eff}}$ at $N_{\text{eff}} = 171$ independent galaxies (Table 17). Galaxy-level cluster bootstrap (1,000 resamples) at $N_{\text{eff}} = 171$ independent galaxies yields $\delta\text{BIC}_{\text{eff}} = +3.4$ vs. Simple MOND (2-param), 95% CI $[-0.9, +7.0]$, formally inconclusive per Kass & Raftery—consistent with the out-of-sample CV gap ($\Delta\text{NMSE} = 0.008$, $z = 0.39$, $p = 0.69$). By contrast, ADCD decisively outperforms Standard MOND (2-param) with $\delta\text{BIC}_{\text{eff}} = -36.5$, CI $[-60.1, -17.9]$ ($z = -4.60$ in CV). We interpret the Simple MOND / RAR indistinguishability result not as a failure of ADCD, but as evidence that autonomous correction discovery can reach the predictive performance of interpolating functions refined by domain experts over decades—without access to those forms as candidates.

Table 14: Galaxy-level cross-validation (10 repeated 50/50 train/test splits by galaxy). The top three rows are zero-parameter canonical forms; the middle row is the ADCD-discovered 2-param form; the bottom three rows are the same MOND/RAR families refit with 2 free parameters (fair parameter-matched comparison). ADCD’s out-of-sample NMSE is statistically indistinguishable from the best 2-param fitted baselines (overlapping $\pm\text{SE}$ intervals), confirming the correction-first search is competitive with hand-crafted MOND forms without requiring prior knowledge of the interpolating-function family.

Model	Test NMSE (mean $\pm$ SE)
Simple MOND	0.6725 ± 0.0350
Standard MOND	0.7016 ± 0.0055
RAR (McGaugh)	0.6601 ± 0.0346
ADCD discovered	<u>0.3856 ± 0.0155</u>
Simple MOND (2-param)	0.3774 ± 0.0147
Standard MOND (2-param)	0.4664 ± 0.0082
RAR (McGaugh, 2-param)	0.3774 ± 0.0147

Table 15: Fair parameter-matched comparison: all four models fitted with exactly 2 free parameters via L-BFGS-B (20 random restarts) on the same stacked SPARC dataset ($N_{\text{gal}} = 171$, $N_{\text{pts}} = 3342$). Under matched parameter count, the fitted Simple-MOND and RAR forms attain marginally lower in-sample NMSE than the ADCD-discovered form (a $\sim 1.5\%$ difference, well within the bootstrap CI). Galaxy-level cross-validation (Table 14) shows the three are statistically indistinguishable out-of-sample. The honest conclusion is therefore that ADCD’s correction-first search recovers an interpolating function *competitive with*—not superior to—well-known 2-param MOND forms, while requiring no prior knowledge of the MOND/RAR functional family.

Model (2-param)	$\hat{\theta}_0$	$\hat{\theta}_1$	NMSE ↓	BIC ↓
Simple MOND (2-param)	0.497	0.548	0.3674	-3329.77
RAR (McGaugh, 2-param)	0.560	0.579	0.3675	-3329.37
ADCD discovered	—	—	<u>0.3729</u>	<u>-3280.69</u>
Standard MOND (2-param)	0.000	0.000	0.4628	-2558.95

Table 16: Bootstrap parameter estimates for the ADCD-discovered interpolating function $\nu(x) = \theta_0(\sqrt{1 + \theta_1/x} - 1) + 1$ (50 galaxy-level bootstrap resamples; 95% percentile CI).

Parameter	Mean	95% CI
θ_0	2.241	[1.048, 4.557]
θ_1	0.268	[0.074, 0.558]

Table 17: Galaxy-level cluster bootstrap on $\delta\text{BIC}_{\text{eff}}$ at $N_{\text{eff}} = 171$ (1,000 resamples, seed 2026). ADCD is the baseline; $\delta\text{BIC}_{\text{eff}} < 0$ means ADCD wins. Per Kass & Raftery [16]: $|\delta| < 2$ weak, 2–6 positive, 6–10 strong, > 10 very strong—but if the 95% CI brackets zero, the result is formally inconclusive.

Model (2-param)	BIC_{eff}	$\delta\text{BIC}_{\text{eff}}$	95% CI	Verdict
Simple MOND	—	+3.4	[−0.9, +7.0]	Inconclusive
RAR (McGaugh)	—	+3.4	[−1.0, +7.0]	Inconclusive
<u>ADCD discovered</u>	$-159.1 \pm [-194.0, -125.8]$	—	—	(baseline)
Standard MOND	—	−36.5	[−60.1, −17.9]	Very Strong [ADCD wins]

Parameter stability. The bootstrap parameter CIs in Table 16 (50 galaxy-level resamples: $\hat{\theta}_0 = 2.24$, 95% CI [1.05, 4.56]; $\hat{\theta}_1 = 0.27$, 95% CI [0.07, 0.56]) confirm both parameters are non-zero and the functional family is stably recovered. The wide $\hat{\theta}_0$ CI reflects a mild parameter degeneracy: in the deep-MOND regime the model depends on the two free parameters mainly through the product $\hat{\theta}_0\sqrt{\hat{\theta}_1}$, so per-fit point estimates of the prefactor vary substantially while the fitted curve (and hence the NMSE) is essentially invariant. The relevant stability signal is therefore the recovered functional family and the out-of-sample/CV NMSE, both of which are stable—not the absolute value of the prefactor. This degeneracy also explains why per-cut point estimates of $\hat{\theta}_0$ in the robustness analysis below move while the fitted curve and model ranking stay fixed.

Robustness to quality cuts. To ensure the result is not an artefact of galaxy selection, we re-ran the discovery under three increasingly strict quality cuts on the SPARC sample (Table 18). The ADCD-discovered function remains the best fit throughout—improving to NMSE 0.272 on the ultra-strict pristine subsample ($N_{\text{gal}} = 75$)—while the RAR and Simple MOND baselines do not. This monotonic improvement as low-quality data is removed is consistent with a genuine signal rather than a noise-driven fit. Note that the per-cut point estimates of $\hat{\theta}_0$ in Table 18 (e.g.

1.83 on the baseline cut) differ from the boxed headline value ($\hat{\theta}_0 \approx 2.24$) because the robustness fits re-optimize each cut independently with a fresh multi-restart search; as discussed under the bootstrap analysis above, the two parameters are mildly degenerate (the fitted curve depends on them mainly through $\hat{\theta}_0\sqrt{\hat{\theta}_1}$), so their individual point estimates move while the fitted curve and the relative ranking of models stay fixed.

Table 18: SPARC robustness across three quality cuts. The ADCD-discovered function consistently outperforms Simple MOND and the McGaugh+16 RAR form across increasingly strict galaxy-quality selection.

Quality cut	Gal.	Pts	$\hat{\theta}_0$	$\hat{\theta}_1$	ADCD	RAR	Simple MOND NMSE
Baseline (Standard)	171	3342	1.828	0.262	0.3750	0.6361	0.6477
Strict (High Quality)	121	2939	1.612	0.324	0.3032	0.5670	0.5812
Ultra-Strict (Pristine)	75	2287	1.197	0.468	0.2718	0.6513	0.6756

Algebraic identification of the discovered family. A natural question—raised by any reviewer familiar with the MOND literature—is whether the ADCD-discovered function is novel. It is not, and we state the equivalence explicitly rather than leave it implicit. Expanding the boxed form,

$$\nu_{\text{ADCD}}(x) = \theta_0 \left(\sqrt{1 + \theta_1/x} - 1 \right) + 1 = (1 - \theta_0) + \theta_0 \sqrt{1 + \theta_1/x} = a + b\sqrt{1 + c/x},$$

with $a = 1 - \theta_0$, $b = \theta_0$, $c = \theta_1$ and the constraint $a + b = 1$. The Simple MOND interpolating function of Famaey & Binney [21] has exactly this algebraic skeleton,

$$\nu_{\text{Simple}}(x) = a_S + b_S \sqrt{1 + c_S/x}, \quad a_S = 0.5, \quad b_S = 0.5, \quad c_S = 4.0.$$

The two forms belong to the same algebraic family. The distinction is in the *transition parameter*: ADCD’s $\hat{c} \approx 0.27$ versus the canonical $c_S = 4.0$ —a factor of ~ 15 . The transition parameter c controls the acceleration scale at which the interpolating function departs from unity and begins the Newtonian-to-deep-MOND crossover: a smaller c pushes the departure to lower x (i.e. lower baryonic acceleration). We therefore claim *autonomous rediscovery of the Simple-MOND family from real data*, not the discovery of a novel family. The scientific content of the SPARC result is the *autonomous-discovery parity*—a two-parameter member of this family recovered without knowledge of the candidate family lands within the confidence band of functions refined by the MOND/RAR community over decades—not novelty of the functional skeleton.

Parameter degeneracy. The two free parameters are mildly degenerate: the fitted curve in the deep-MOND regime depends on them mainly through the product $\theta_0\sqrt{\theta_1}$. Across the three quality cuts of Table 18 this product is approximately constant while the individual parameters move substantially (Table 19). The deep-MOND asymptote $\nu \sim \theta_0\sqrt{\theta_1/x}$ is therefore stable across cuts at $\approx 0.82\text{--}0.94 x^{-1/2}$, compared with the canonical Simple-MOND asymptote $1.0 x^{-1/2}$ —a $\sim 6.5\%$ difference in the deep-MOND slope. The relevant stability signal is the recovered family and the out-of-sample NMSE (both stable), not the absolute prefactor.

Table 19: SPARC parameter degeneracy across quality cuts. The deep-MOND product $\theta_0\sqrt{\theta_1}$ is approximately constant even as the individual parameters move by a factor of ~ 1.5 , confirming that the fitted curve is governed by the product; individual point estimates are not separately identifiable from the stacked relation.

Quality cut	$\hat{\theta}_0$	$\hat{\theta}_1$	$\hat{\theta}_0\sqrt{\hat{\theta}_1}$
Baseline (171 gal)	1.83	0.262	0.935
Strict (121 gal)	1.61	0.324	0.918
Ultra-strict pristine (75 gal)	1.20	0.468	0.819
Simple MOND canonical	0.50	4.000	1.000

Scope of the claim. We do *not* claim that ADCD has resolved the modified-gravity / dark-matter question, nor that the discovered family is predictively superior to the best two-parameter MOND/RAR interpolating functions. The stacked relation deliberately averages out galaxy-to-galaxy diversity, and the absolute NMSE is dominated by intrinsic scatter. What the data do show, defensibly and repeatably, is a three-level result: ADCD—starting only from the Newtonian baseline and the asymptotic constraint $\Delta \rightarrow 0$ —(i) decisively outperforms the fixed zero-parameter canonical forms in stacked NMSE, (ii) cannot be statistically distinguished from the best two-parameter domain-expert interpolating families out-of-sample and under cluster bootstrap, and (iii) decisively outperforms the two-parameter Standard-MOND family. The scientific content of the result is the *autonomous-discovery parity*: a two-parameter interpolating function recovered without knowledge of the candidate family lands within the confidence band of functions refined by the MOND/RAR community over decades.

5.7 Cosmological Probes: Growth Rate and Expansion History

The SPARC result (Section 5.6) establishes that ADCD can recover a *functional* correction from a galactic-scale anomaly. A natural and demanding complement is to apply the same machinery to large-scale cosmological tensions, where the prevailing interpretation is that of an *amplitude* mismatch rather than a functional one. We test this on two independent observables: the structure-growth rate $f\sigma_8(z)$ and the expansion history $H(z)$.

Growth rate $f\sigma_8(z)$. We use the 63-point $f\sigma_8$ compilation of Alestas et al. [23] (Table II; $z \in [0.001, 1.944]$). The General-Relativity baseline is the standard linear-theory prediction

$$f\sigma_8(z) = \Omega_m(z)^\gamma D(z) \sigma_{8,0}, \quad \gamma \approx 0.55,$$

with $\sigma_{8,0}$ treated as the baseline’s single free amplitude (fitted by weighted least squares on the compilation). We then form the residual $\Delta(z) = f\sigma_8^{\text{obs}}(z) - f\sigma_8^{\text{GR}}(z)$ and run ADCD’s candidate library—constant offset, linear, quadratic, power law, power+offset, exponential, exponential+offset, rational, and logarithmic—each fit by L-BFGS-B with 30 random restarts and scored by BIC against $\Delta(z)$. The 1-parameter constant offset $\Delta(z) = \theta_0$ is included as the null model: it is the simplest physically meaningful description of an S_8 -like amplitude tension (a pure $\sigma_{8,0}$ renormalisation, with no z -dependence).

Homogeneous-subset cross-check. Because the 63-point compilation mixes ~ 15 surveys with heterogeneous modelling systematics, a z -dependent correction that “wins” on the full compilation could be an artefact of inter-survey offsets. We therefore re-run the discovery on three homogeneous subsets chosen for minimal systematic contamination: the 6 canonical BOSS DR12 (Alam+2017) consensus points, the 20 most-precise measurements, and the 31-point BOSS-dominated mid- z window $z \in [0.35, 0.75]$.

Expansion history $H(z)$. As an independent observable, we use the 34-point cosmic-chronometer $H(z)$ compilation ($z \in [0.070, 1.965]$), which is model-independent and free of CMB/BAO/RSD systematics. The reference Λ CDM prediction is $H(z) = H_0 \sqrt{\Omega_{m,0}(1+z)^3 + \Omega_{\Lambda,0}}$ with $(H_0, \Omega_{m,0})$ as two nuisance amplitudes; we then run the same candidate library on $\Delta_H(z) = H_{\text{obs}} - H_{\text{GR}}$.

Result: no functional correction is detected. Table 20 summarises all five tests, and Figure 7 shows the $f\sigma_8$ and $H(z)$ residuals with the ADCD-selected constant-offset model. In *every* case the verdict is `CONSTANT_WINS`: no z -dependent functional candidate beats the 1-parameter constant offset by the Kass–Raftery decisive threshold $\Delta\text{BIC} < -10$. On the full $f\sigma_8$ compilation the best candidate *is* the constant offset (ΔBIC vs. null = +0.00); on the homogeneous subsets the best candidates are Linear (BOSS DR12, $\Delta\text{BIC} = -9.63$, short of the threshold with $N = 6$), Constant offset (precision), and Power law (mid- z , $\Delta\text{BIC} = -0.00$); on $H(z)$ the constant offset again wins ($\Delta\text{BIC} = +0.00$). The best-fit amplitudes are physically reasonable: $\sigma_{8,0} = 0.7815$ (−3.7% vs. Planck 2018’s 0.811—consistent in direction with the S_8 tension), and $H_0 = 68.48 \text{ km s}^{-1} \text{ Mpc}^{-1}$, $\Omega_{m,0} = 0.296$ on $H(z)$.

Table 20: ADCD cosmological-probe results across five independent tests spanning two observables ($f\sigma_8$ and $H(z)$) and three homogeneous survey subsets. “Constant offset” is the 1-parameter null model $\Delta(z) = \theta_0$ (a pure amplitude rescaling). No z -dependent functional candidate beats the constant offset by $\Delta\text{BIC} < -10$ (the Kass–Raftery decisive-evidence threshold) in any test; the verdict is therefore `CONSTANT_WINS` throughout. Best-fit amplitudes: $\sigma_{8,0} = 0.7815$ (−3.7% vs. Planck 2018: 0.811); $H_0 = 68.48 \text{ km s}^{-1} \text{ Mpc}^{-1}$, $\Omega_{m,0} = 0.296$.

Test (dataset)	N	Observable	Best candidate	ΔBIC vs. null	Verdict
Full compilation (Alestas+22)	63	$f\sigma_8$	Constant offset	+0.00	<code>CONSTANT_WINS</code>
BOSS DR12 (Alam+17)	6	$f\sigma_8$	Linear	−9.63	<code>CONSTANT_WINS</code>
Precision (top-20)	20	$f\sigma_8$	Constant offset	+0.00	<code>CONSTANT_WINS</code>
Mid- z window [0.35, 0.75]	31	$f\sigma_8$	Power law	−0.00	<code>CONSTANT_WINS</code>
Cosmic chronometers	34	$H(z)$	Constant offset	+0.00	<code>CONSTANT_WINS</code>




Figure 7: Cosmological-probe residuals and the `CONSTANT_WINS` verdict. **(a)** $f\sigma_8$ residual $\Delta(z) = f\sigma_8^{\text{obs}} - f\sigma_8^{\text{GR}}$ (63 points, Alestas+22); the GR baseline has $\sigma_{8,0}$ free. Inverse-variance bin means (red) are consistent with zero across all redshifts ($|\text{pull}| \leq 0.5\sigma$), and the ADCD-selected model is the 1-parameter constant offset (dashed). **(b)** Same for $H(z)$ (34 cosmic-chronometer points); the best candidate is again the constant offset. In both observables no z -dependent functional form beats a pure amplitude rescaling by $\Delta\text{BIC} < -10$.

Interpretation. Across five independent tests spanning two observables and three survey subsets, ADCD finds *no* z -dependent functional correction beyond a constant amplitude rescaling. Two interpretations are consistent with this null result:

- *Physics interpretation:* the S_8 tension originates from an amplitude renormalisation (e.g. a σ_8 normalisation set by early-universe physics) rather than modified growth dynamics—in which case no functional z -correction exists to be discovered.
- *Sensitivity interpretation:* a functional correction exists but falls below ADCD’s detection threshold given current survey precision ($\sim 20\%$ typical point errors in $f\sigma_8$) and heterogeneous systematics across 15+ survey instruments.

The current data cannot distinguish these interpretations. The `CONSTANT_WINS` result is therefore best read as a *quantitative upper bound* on the detectability of a functional growth correction in the present compilation, not as a proof of correction absence. In a preliminary validation run using the newly released DESI DR1 data ($N = 6$ consensus bins) combined with eBOSS ($N = 7$), a dynamical growth-index search yielded a nominal 2.34σ deviation ($\gamma(z) \approx 0.38 + 0.44z$). However, we explicitly note that this is a statistical artefact of small- N data compilation: the volume and spatial overlap of eBOSS and DESI LRG targets at $z \sim 0.7$ artificially inflates significance when cross-covariance is neglected. When integrated into the full 63-point Alestas et al. compilation, the General Relativity constant growth ($\gamma = 0.55$) remains the undisputed best fit ($\Delta\text{BIC} < -9.6$). Future surveys (DESI full release, Euclid, Rubin/LSST) with $3\text{--}5\times$ improvement in $f\sigma_8$ precision will provide a definitive test. The result also demonstrates that ADCD does not surface spurious functional corrections when none is supported by the data—the fail-safe property established on the synthetic misspecification benchmark (Section 5) carries over to real cosmological archives.

5.8 Structural Dichotomy: Functional vs. Amplitude Anomalies

Taken together, the galactic (SPARC) and cosmological ($f\sigma_8$, $H(z)$) results exhibit a qualitative pattern that we report as an *observation*, not a proven conclusion. Table 21 summarises the outcomes across scales, and Figure 8 contrasts the two regimes side by side. Functional corrections—where ADCD recovers a z - or r -dependent Δ that genuinely improves on the classical baseline—are found in individual gravitational systems: textbook orbital and force-law anomalies (synthetic), and galactic rotation curves (SPARC). Cosmological tensions, by contrast, manifest as amplitude offsets with no detected functional structure: ADCD correctly identifies the *absence* of a functional correction and selects a constant rescaling.

Table 21: Structural dichotomy across physical scales. “Functional” denotes that ADCD recovered a z - or r -dependent correction that beats the classical baseline; “Amplitude” denotes that only a constant rescaling was supported; “Ambiguous” denotes an inconclusive cross-system test (Appendix A).

System	Scale	Outcome	Anomaly type
Textbook physics (9 scenarios)	micro–macro	Functional	z/r -dependent correction
SPARC rotation curves	galactic ($\sim\text{kpc}$)	Functional	r -dependent correction
Wide binary stars	stellar ($\sim\text{kAU}$)	Ambiguous	cross-system, inconclusive (App. A)
Growth rate $f\sigma_8(z)$	cosmological ($\sim\text{Gpc}$)	Amplitude	constant rescaling only
$H(z)$ chronometers	cosmological ($\sim\text{Gpc}$)	Amplitude	constant rescaling only

Framing. This pattern is *consistent with*—but does not prove—a structural dichotomy between two classes of physical anomaly. *Class I (Functional):* anomalies in individual gravitational systems (orbital mechanics, galactic rotation) that require z - or r -dependent corrections to the

force law, which ADCD reliably discovers. *Class II (Amplitude)*: cosmological tensions (S_8 , H_0) that manifest as amplitude offsets with no detected functional structure, which ADCD correctly declines to over-fit. Whether this reflects genuine physics (different origins for micro- versus macro-scale anomalies) or is a methodological artefact of the order-of-magnitude difference in data quality and systematic complexity between the two regimes is an open and testable question. We report the pattern as an observation requiring further investigation, and flag in Section 7 that `CONSTANT_WINS` is an upper bound on detectability rather than a proof of absence.




Figure 8: The structural dichotomy in one figure. **(a)** Galactic scale (SPARC): ADCD recovers a *functional* interpolating curve (red) that tracks the data across the Newtonian-to-deep-MOND transition, dominating the zero-parameter canonical MOND/RAR forms. **(b)** Cosmological scale ($f\sigma_8$): the residual carries no z -dependent structure; ADCD selects a *constant amplitude* offset (dashed), and the best z -dependent candidate (Rational, dot-dashed) is statistically indistinguishable from it ($\Delta\text{BIC} = +0.42$, far below the decisive threshold). Individual gravitational systems admit functional corrections; large-scale cosmological tensions appear as amplitude offsets.

6 Discussion and Ablation Studies

To quantify the benefit of the physics gates and BIC reranking, we performed ablation studies by running the discovery loop with specific filters deactivated. Quantitative results for gate-removal conditions are shown in Figure 5; the BIC ablation below is a targeted case study reported in text because it requires the full gate pipeline intact to isolate BIC’s contribution.

- **Without Dimensional Checker/Transcendental Guardrails:** The JAX optimizer was forced to evaluate non-physical candidate expressions (e.g., e^{-r} where r has units of meters). These expressions routinely produce numerical overflow and underflow during parameter search, resulting in qualitatively increased optimization overhead. On this benchmark, structural accuracy at 5% noise drops to 7/9 ($\Delta = 11.1$ pp) due to the misclassification of Yukawa Gravity, confirming its role as the primary structural filter. Note: the transcendental guardrail produced zero additional rejections beyond the dimensional gate on this benchmark (all transcendental-argument violations were already caught by dimensional homogeneity); this is expected because a dimensioned argument $\sin(v)$ with v in m/s fails the dimensional check before reaching the transcendental stage.
- **Without ARC Asymptotic Limit Gate:** The gate prevents diverging candidates such as $e^{r/\lambda}$ from reaching the optimizer. Without this gate, such expressions can fit training data locally yet diverge as $r \rightarrow \infty$, producing unphysical extrapolation behavior. On this benchmark, the in-distribution structural accuracy was unchanged ($\Delta = 0$ at 5% noise),

but the ARC gate screens for classical-limit consistency in the surviving candidate set—a property that cannot be assessed by in-distribution accuracy alone.

- **Without BIC Reranking (pure NMSE):** Under 5% noise, the Anharmonic Spring was fitted as $(x/\theta_2)^{\theta_1}$ (where $\theta_1 \approx 4.02$ and $\theta_2 \approx 1.01$). While numerically equivalent, this was classified as a `power_law` structural class instead of `polynomial` due to the inclusion of the unnecessary scaling exponent θ_1 . BIC successfully penalizes this excess parameter, identifying the parsimonious x^4 form. This ablation is not shown in Figure 5 because disabling BIC requires the full gate pipeline; it is a targeted case study rather than a parallel gate-removal condition.

6.1 Search Space Efficiency

Empirical Proposition 1 (Search-space reduction). Let N_{prop} denote candidates proposed by the correction proposer, N_{gate} denote expressions surviving all Stage 1 gates (measured per-gate telemetry in `gate_telemetry.json`), and N_{opt} denote JAX optimization calls. For the 9-scenario benchmark suite, ADCD satisfies $N_{\text{opt}} \ll N_{\text{PySR}}$ where N_{PySR} is the hall-of-fame size under PySR’s fair profile ($n_{\text{iter}} = 100$, `maxsize` = 30, 60 s timeout). Gate-level survival rates decompose the funnel: parse $\rightarrow$ AST complexity $\rightarrow$ dimensional homogeneity $\rightarrow$ transcendental guardrail $\rightarrow$ ARC asymptotic limit $\rightarrow$ coarse NMSE pre-filter.

A key advantage of the correction-first paradigm is quantified search-space reduction (Table 7): 3,225 proposed templates funnel to 655 ARC survivors (20.3%) and 598 JAX optimization calls, versus PySR fair’s ~ 14 hall-of-fame expressions per run with no physics pre-filtering (Table 8). Despite exploring fewer expressions, ADCD achieves higher structural accuracy—especially at 5% noise (88.9% vs. 11.1%)—on the correction-to-a-known-law task, a regime in which unconstrained tabula rasa search is not designed to operate.

Trigonometric compositions at zero noise: a property of unconstrained search. An instructive property of unconstrained symbolic regression emerges at 0% noise: all five PySR fair mismatches (Table 6) produce `trigonometric` expressions, regardless of the true functional family. This is not a defect in PySR but a systematic consequence of its search geometry, which is the expected behaviour of a universal approximator used for tabula rasa discovery. Under zero noise, the optimizer must only minimize training error; there is no regularization signal from stochastic label noise. With `maxsize`=30 and the operator set $\{\sin, \cos, \tanh, \sinh, \log, \exp\}$, PySR’s Pareto front readily discovers deeply nested compositions (e.g., $\sin(\sin(\tanh(\tanh(\dots))))$) that are smooth, bounded, and able to approximate any continuous function on a bounded domain to arbitrary precision—a classical result from approximation theory. Because these over-parameterized trigonometric compositions achieve zero training error, they dominate the BIC-ranked hall of fame despite being structurally incorrect. This is a perfectly reasonable outcome when the task is full-law discovery from scratch; the issue is the task, not the tool. ADCD sidesteps this regime by design: because it searches for a correction to a *known* law, physics gates reject trigonometric candidates whose arguments are dimensioned (e.g., $\sin(v)$ where v has units of m/s) or whose asymptotic limits fail ($\sin(v/c) \not\rightarrow 0$ as $v \rightarrow 0$), so it never enters the universal-approximator regime in which such compositions are competitive. This is not a claim of superiority over unconstrained SR but an illustration of how restricting the task—from full-law discovery to correction discovery—changes the relevant failure modes.

7 Limitations and Future Work

Correction-first scope limitation. ADCD is designed for anomaly-driven theory refinement, not theory replacement. The framework assumes the classical baseline is structurally correct and the anomaly is incremental. When the baseline theory is fundamentally wrong, the correction

term may become arbitrarily complex, at which point the correction-first search offers no advantage over tabula rasa symbolic regression.

Multi-scale and hierarchical baselines. A natural extension of the EFT-inspired correction discovery is to search for sub-leading correction operators by subtracting a known leading-order correction from the baseline before running the symbolic search. However, this hierarchical baseline approach introduces a fundamental information-theoretic constraint: once the dominant correction is absorbed into the baseline, the remaining sub-leading signal is typically orders of magnitude smaller than the observational noise floor. The signal-to-noise ratio of the residual can therefore drop below the threshold of identifiability even at noise levels where the leading-order search succeeds, causing the optimizer to converge to degenerate solutions rather than physically meaningful sub-leading terms. This constraint implies that hierarchical multi-scale search is better suited for open-ended real-world exploration where the leading-order correction is already established from independent evidence (e.g., post-Newtonian corrections to orbital mechanics), rather than for benchmark validation where the target correction constitutes the primary discoverable effect.

Numerical scale sensitivity. Sub-nano-scale residuals (binary pulsar $|dP/dt| \sim 10^{-15}$ s/s) exposed a scale-adaptive NMSE bug in early experiments: a fixed variance floor ($\varepsilon = 10^{-10}$) in the JAX optimizer caused degenerate $\theta \rightarrow 0$ solutions to appear optimal. We corrected this with scale-adaptive denominators in both the optimizer and the post-hoc evaluator (`metrics.py`), and re-rank Stage 2 candidates using validated NMSE rather than optimizer-internal scores. Remaining challenges include blackbody UV fitting (NMSE ≈ 0.026) and binary pulsar exponent recovery ($\theta_1 = 1.57$ vs. true $5/3$); logarithmic parameterization of prefactors would further improve extreme-scale scenarios.

Reduced-variable real-world benchmarks. The binary pulsar v2.1 benchmark holds M , a , and e fixed at Hulse-Taylor values and varies only P , yielding a clean $P^{-5/3}$ power law. Our sensitivity study (Table 11) shows this is substantially easier than the legacy four-parameter formulation, where ADCD fails structurally. Future work must report multi-parameter orbital decay without reduced-variable simplification, or clearly label such benchmarks as controlled ablations rather than full discovery claims.

Multi-variable corrections (Phase 2: exploratory, not a headline result). We implemented a multivariable extension of ADCD in which a **BuckinghamPiEngine** computes dimensionless Π -groups from the registered variable set (nullspace of the dimensional exponent matrix [17]), a **SequentialARCChecker** tests each limit variable independently, and a **ProductGrammar** proposes correction candidates as products of Π -group functions. On a four-scenario multivariable benchmark (Table 22), the extended Mock proposer recovers 3/4 scenarios at 0% and 1% noise and 2/4 at 5% noise, while the ProductGrammar recovers 2/4 at 0% noise, degrading to 1/4 at 1% and 0/4 at 5%. These results are *not* reported as a headline capability: the two proposers fail for disjoint reasons. The Mock proposer never recovers MV-1—a mass-ratio prefactor multiplying exponential radial decay, $\theta_0(m/M)e^{-r/r_0}$ —because its template bank does not span this mixed prefactor/exponential class. The ProductGrammar, conversely, is defeated by numerical ill-conditioning whenever reference constants span extreme scales (e.g. $n_0 \sim 10^{20} \text{ m}^{-3}$ in MV-2), where Π -group coefficient optimization collapses to degenerate forms such as $10^{-20}n\theta_0$ at 64-bit precision, and by a spurious-exponential class mismatch on MV-4. Both proposers additionally lose MV-3 at 5% noise to residual blow-up (NMSE ≈ 0.28) despite retaining the correct structural class. Importantly, every failure is *correctness-preserving*: no candidate that mismatched the structural class also cleared the NMSE gate, so the pipeline

never reports a confident spurious correction—it simply fails to converge. Logarithmic parameterization of Π -group coefficients and a richer unary operator set for the ProductGrammar are planned improvements before multivariable discovery is promoted from exploratory extension to a headline claim.

Table 22: Phase 2 Multivariable Benchmark (exploratory; reported in Limitations). $\checkmark$ = NMSE < 0.05 and correct structural class; $\times$ = fail. Mock = Mock (extended) proposer; PG = ProductGrammar. Variables: MV-1 (m, M, r), MV-2 (n, T), MV-3 (ρ, v), MV-4 (n, V).

Scenario	Correction Δ	0%		1%		5%	
		Mock	PG	Mock	PG	Mock	PG
MV-1: Yukawa	$\theta_0 (m/M) e^{-r/r_0}$	$\times$	$\checkmark$	$\times$	$\times$	$\times$	$\times$
MV-2: Plasma	$\theta_0 (n/n_0)(T_0/T)^{1/2}$	$\checkmark$	$\times$	$\checkmark$	$\times$	$\checkmark$	$\times$
MV-3: Turbulent Drag	$\theta_0 \rho v^2$	$\checkmark$	$\checkmark$	$\checkmark$	$\checkmark$	$\times$	$\times$
MV-4: Van der Waals	$\theta_0 (n/n_0)^2/(V/V_0)^2$	$\checkmark$	$\times$	$\checkmark$	$\times$	$\checkmark$	$\times$
Total		3/4	2/4	3/4	1/4	2/4	0/4

Dimensional gate relaxation. The dimensional checker uses an adaptive relaxation where free constants (parameters θ_i) are treated as dimensionless exponents or scaling coefficients that can adjust their units to match the target dimension, provided the expression is internally homogeneous (preventing unit clashes like $m + v$). While this allows physically consistent parameterized corrections to pass Stage 1, it places the burden of rejecting structurally incorrect but dimensionally relaxed candidates onto the Stage 2 optimizer and the BIC reranking. Future work will investigate tighter physical constraint gates that restrict parameter dimensions based on specific scaling symmetries.

Proposer expressiveness bound. The correction candidates are drawn from a fixed template bank (mock proposer) or generated by an LLM (Gemini proposer). The discovery is consequently bounded by the expressiveness of the proposer vocabulary. Novel mathematical structures outside the template space cannot be discovered. Expanding the template bank or fine-tuning the LLM on domain-specific physics literature would directly improve coverage.

Information-theoretic limits at high noise. As demonstrated by the Screened Coulomb scenario at $\geq 5\%$ noise, certain mathematical structures (exponential decay $e^{-r/\lambda}$ vs. rational saturation $r/(r + \lambda)$) can be fundamentally indistinguishable when the signal-to-noise ratio is sufficiently low and the dynamic range is limited. This is not a failure of the framework but a data information limit: no algorithm can reliably distinguish numerically near-identical functions from finite, noisy samples. Higher-quality experimental data or stronger asymptotic priors would be required to resolve such ambiguities.

Comparison with physics-constrained tabula rasa SR. A natural extension is direct comparison with physics-constrained tabula rasa symbolic regression methods such as PhySO [14]. ADCD’s correction-first paradigm reduces the hypothesis space by conditioning on a known classical law, whereas PhySO operates tabula rasa with dimensional/unit constraints; the two approaches therefore address fundamentally different discovery scenarios (theory refinement vs. theory replacement). We leave systematic comparison across noise levels and classical-baseline assumptions to future work. We note that ADCD’s headline contribution—the correction-first paradigm validated against fair tabula rasa baselines such as PySR [2] (Section 5)—is orthogonal to the correction-first-vs-physics-constrained-tabula-rasa axis that a PhySO comparison would probe.

Correction mode specification. The public API supports `correction_mode=auto` (the default in `discover_correction`), which infers additive versus multiplicative structure from residual statistics. Users may still override the mode when domain knowledge dictates a specific form ($y = y_{\text{cls}} + \Delta$ vs. $y = y_{\text{cls}}(1 + \Delta)$). Detection accuracy degrades when the classical baseline is near zero or when additive and relative residuals are statistically ambiguous.

Amplitude vs. functional distinguishability (cosmological probes). The `CONSTANT_WINS` verdicts in Section 5.7 are best read as quantitative upper bounds on the detectability of a functional growth correction, not as proofs that no correction exists. ADCD cannot detect a functional correction whose amplitude falls below the noise floor of the available data. With current $f\sigma_8$ compilations (typical point errors $\sim 20\%$, heterogeneous systematics across 15+ surveys), a z -dependent correction of small but non-zero amplitude would be statistically indistinguishable from the 1-parameter constant offset. Definitive detection/non-detection awaits $3\text{--}5\times$ improvements in $f\sigma_8$ precision from DESI, Euclid, and Rubin/LSST.

Heterogeneous survey compilations. The $f\sigma_8$ compilation of Alestas et al. [23] aggregates measurements from surveys with different estimators (multipole versus clustering- F versus BAO- f), different scale cuts, and different galaxy-bias models. Systematic offsets between surveys could appear as—or mask—a functional z -correction. Our homogeneous-subset cross-check (BOSS DR12, precision, mid- z) is designed to mitigate this, but no subset is fully free of modelling-pipeline systematics; a residual functional signal of astrophysical rather than cosmological origin cannot be excluded from first principles.

Cross-system generalization (galactic versus stellar). The SPARC-discovered form ($c \approx 0.27$) does not unambiguously generalize to wide-binary kinematics (Appendix A): it is closer to the observed wide-binary velocity boost than Simple MOND at $s \sim 10\text{--}15$ kAU but diverges at larger separations, while no single interpolating function simultaneously optimizes both SPARC and wide binaries. This suggests that the MOND transition parameter may be system-dependent—or, more conservatively, that the wide-binary dataset is too contaminated (Chae [24] versus Banik et al. [25]) to serve as a clean discriminant. Universal interpolating-function discovery from a single system class should be treated as provisional pending cleaner multi-system data.

Synthetic-real hybrid validation. Our benchmark scenarios use synthetically generated observations, following the established precedent of the SciML community: the foundational symbolic regression works of Udrescu & Tegmark [1] (AI Feynman), Cranmer [2] (PySR), and Brunton et al. (SINDy) all validated primarily on synthetic data before real-archive application. The validity of this approach in our setting is further grounded by construction: all physical constants in the real-world tier scenarios (gravitational wave decay, Lamb shift, Muon $g\text{--}2$ anomaly, Planck blackbody) are drawn directly from authoritative reference sources—JPL DE440 orbital elements, NIST CODATA 2018 fundamental constants, and CERN Particle Data Group values—ensuring the synthetic observations encode the correct dimensional scales, noise regimes, and asymptotic behaviors of the underlying physical systems. This synthetic-real hybrid design allows controlled noise injection (0–10%) and ground-truth structural evaluation that would be impossible with raw experimental archives, while remaining anchored to physically meaningful parameter ranges. Full validation against curated experimental archives (e.g., NIST atomic spectra tables, JPL planetary ephemerides, CERN particle listings) is a natural and important next step, which we identify as a priority for future work.

Acknowledgments

The author thanks AI coding assistants—Cursor IDE, Google DeepMind’s Antigravity, and Qoder—for assistance in code development, benchmark execution, figure generation, and manuscript drafting. All scientific content, experimental design decisions, and intellectual contributions are the author’s own. AI assistance was used as a pair-programming and writing tool; no AI system generated experimental results autonomously.

8 Conclusion

Methodology. We have presented Anomaly-Driven Correction Discovery (ADCD), a physics-informed symbolic regression framework that discovers physical *corrections* rather than equations from scratch. Statistical residual-feature priors, a cascaded physics-gate filter (AST complexity, dimensional homogeneity with a transcendental guardrail, asymptotic-regime consistency), a JAX-traced L-BFGS-B optimizer, and Bayesian Information Criterion reranking together constrain the search to physically plausible corrections before any numerical optimization occurs. ADCD is designed for anomaly-driven theory refinement—where the classical baseline is structurally correct and the anomaly is incremental—rather than wholesale theory replacement.

Synthetic benchmark. On a 9-anomaly suite across textbook, cross-domain, and synthetic systems under noise up to 10%, ADCD achieves a mean structural recovery of 80.4% ($\pm 7.4\%$) across sixteen seeds (94.4% at the reference seed=42, disclosed as the highest-performing seed). It outperforms PySR by 77.8 percentage points at 5% noise under matched residuals—a gap that does not close when PySR’s budget is doubled—and fails safe on the misspecification benchmark. The same machinery recovers the correct structural class on 4/4 real-world physics-constant scenarios.

SPARC: autonomous rediscovery of the Simple-MOND family. On the SPARC sample of 171 galactic rotation curves, ADCD—starting only from the Newtonian baseline and the high-acceleration asymptotic constraint—recovers a two-parameter member of the Simple-MOND algebraic family (Famaey & Binney 2005) with transition parameter $\hat{\theta}_1 \approx 0.27$ versus the canonical $c = 4.0$. The discovered form delivers a 41% stacked-NMSE reduction over fixed zero-parameter canonical forms and, under fair galaxy-level cross-validation plus a 1,000-resample cluster bootstrap, cannot be distinguished from two-parameter domain-expert forms ($\delta\text{BIC}_{\text{eff}} = +3.4$, 95% CI brackets zero) while decisively beating two-parameter Standard MOND ($\delta\text{BIC}_{\text{eff}} = -36.5$). This is autonomous rediscovery of a known family from real data—not the discovery of a novel family.

Wide-binary cross-validation. An exploratory cross-system test (Appendix A) finds that the SPARC-discovered form does *not* unambiguously generalize to wide-binary kinematics: it is closer to the Chae 2023 observations than Simple/RAR MOND at $s \sim 10\text{--}15$ kAU but diverges at larger separations, and no single interpolating function simultaneously optimizes both datasets. We treat this as preliminary, given the unresolved Chae-versus-Banik observational controversy, and flag non-generalization as an explicit limitation.

Cosmological probes: a quantitative upper bound. Extended to 63 $f\sigma_8$ growth-rate measurements and 34 $H(z)$ cosmic-chronometer points, ADCD finds *no* z -dependent functional correction beyond a constant amplitude rescaling in five independent tests spanning two observables and three homogeneous survey subsets. The best-fit $\sigma_{8,0} = 0.7815$ (-3.7% vs. Planck) is consistent with the S_8 tension direction. We interpret this `CONSTANT_WINS` result as a quantitative upper bound on the detectability of a functional growth correction in present

compilations—not as proof of correction absence—and note that ADCD’s fail-safe property carries over from the synthetic misspecification benchmark to real cosmological archives.

Structural dichotomy. Across scales, a qualitative pattern emerges: functional corrections are recoverable in individual gravitational systems (textbook anomalies, SPARC), while cosmological tensions manifest as amplitude offsets with no detected functional structure. Whether this reflects a genuine dichotomy between anomaly classes, or simply the anomalies we happened to test, remains an open and testable question; we do not claim it as a proven conclusion.

Future work. Three directions follow directly from the present results. First, the cosmological probes should be re-run on DESI, Euclid, and Rubin/LSST data as they deliver $3\text{--}5\times$ improvements in $f\sigma_8$ precision, providing a definitive test of the `CONSTANT_WINS` bound. Second, the multivariable extension (Buckingham-II dimensionless groups with sequential asymptotic checking) is reported here as an exploratory capability with open challenges (2–3/4 scenarios); promoting it to a headline claim requires resolving the numerical-ill-conditioning and rational/power-law-composite failure modes. Third, the framework should be extended to partial differential equations representing anomalies in spatio-temporal datasets, such as general-relativistic corrections to classical fluid flows.

Computational accessibility. Unlike neural-guided or reinforcement-learning-based symbolic regression frameworks (e.g., PhySO [14], DSO [9]) that require multi-hour training on dedicated GPU clusters, ADCD completes dimensional screening and parameter optimisation within seconds to minutes on a standard consumer-grade CPU. The physics-gate cascade (Stage 1) eliminates the vast majority of candidate expressions before any numerical fitting occurs, bypassing the computational bottleneck of high-dimensional neural search. This resource-light architecture makes physics-constrained symbolic regression accessible on personal laptops without high-performance computing dependencies, democratising anomaly-driven theory refinement for research groups without institutional HPC access.

Code Availability

The complete source code for the ADCD framework, along with all reproduction scripts and benchmarks, is open-source and publicly available at <https://github.com/apiprdt/PhysicsPaper>. The package is also published on the Python Package Index (PyPI) at <https://pypi.org/project/adcd/>.

References

- [1] S. M. Udrescu and M. Tegmark, “AI Feynman: A physics-inspired systematic approach to symbolic regression,” *Science Advances*, vol. 6, no. 16, 2020.
- [2] M. Cranmer, “PySR: Fast & parallelized symbolic regression in Python/Julia,” *arXiv preprint arXiv:2007.03738*, 2020.
- [3] M. Schmidt and H. Lipson, “Distilling free-form natural laws from experimental data,” *Science*, vol. 324, no. 5923, pp. 81–85, 2009.
- [4] S. L. Brunton, J. L. Proctor, and J. N. Kutz, “Discovering governing equations from data by sparse identification,” *PNAS*, vol. 113, no. 15, pp. 3932–3937, 2016.
- [5] G. Schwarz, “Estimating the dimension of a model,” *The Annals of Statistics*, vol. 6, no. 2, pp. 461–464, 1978.

- [6] M. Raissi, P. Perdikaris, and G. E. Karniadakis, “Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations,” *Journal of Computational Physics*, vol. 378, pp. 686–707, 2019.
- [7] S. Greydanus, M. Dzamba, and J. Yosinski, “Hamiltonian neural networks,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019.
- [8] M. Cranmer, S. Greydanus, S. Hoyer, P. Battaglia, D. Spergel, and S. Ho, “Lagrangian neural networks,” in *ICLR 2020 Deep Differential Equations Workshop*, 2020.
- [9] B. K. Petersen et al., “Deep symbolic regression: Recovering mathematical expressions from data via risk-seeking policy gradients,” in *International Conference on Learning Representations (ICLR)*, 2021.
- [10] H. Kitano, “Nobel Turing challenge: creating the engine for scientific discovery,” *npj Systems Biology and Applications*, vol. 7, no. 29, 2021.
- [11] B. Romera-Paredes et al., “Mathematical discoveries from program search with large language models,” *Nature*, vol. 625, pp. 468–475, 2024.
- [12] J. Rissanen, “Modeling by shortest data description,” *Automatica*, vol. 14, no. 5, pp. 465–471, 1978.
- [13] P. Langley, H. A. Simon, G. L. Bradshaw, and J. M. Zytkow, *Scientific Discovery: Computational Explorations of the Creative Processes*. Cambridge, MA: MIT Press, 1987.
- [14] W. Tenachi, R. Ibata, and F. I. Diakogiannis, “Deep symbolic regression for physics guided by units constraints: Toward the automated discovery of physical laws,” *Astrophysical Journal*, vol. 959, no. 2, p. 99, 2023.
- [15] A. Grayeli, A. Sehgal, O. Costilla-Reyes, M. Cranmer, and S. Chaudhuri, “Symbolic Regression with a Learned Concept Library,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- [16] R. E. Kass and A. E. Raftery, “Bayes Factors,” *Journal of the American Statistical Association*, vol. 90, no. 430, pp. 773–795, 1995.
- [17] E. Buckingham, “On Physically Similar Systems; Illustrations of the Use of Dimensional Equations,” *Physical Review*, vol. 4, no. 4, pp. 345–376, 1914.
- [18] L. Lelli, S. S. McGaugh, and J. M. Schombert, “SPARC: Mass Models for 175 Disk Galaxies with Spitzer Photometry and Accurate Rotation Curves,” *The Astronomical Journal*, vol. 152, no. 6, p. 157, 2016.
- [19] M. Milgrom, “A Modification of the Newtonian Dynamics as a Possible Alternative to the Hidden Mass Hypothesis,” *The Astrophysical Journal*, vol. 270, pp. 365–370, 1983.
- [20] S. S. McGaugh, F. Lelli, and J. Schombert, “Radial Acceleration Relation in Rotationally Supported Galaxies,” *Physical Review Letters*, vol. 117, no. 20, p. 201101, 2016.
- [21] B. Famaey and J. Binney, “Modified Newtonian Dynamics in the Milky Way,” *Monthly Notices of the Royal Astronomical Society*, vol. 363, no. 3, pp. 985–992, 2005.
- [22] B. Famaey and S. S. McGaugh, “Modified Newtonian Dynamics (MOND): Observational Phenomenology and Relativistic Extensions for Galaxies,” *Living Reviews in Relativity*, vol. 15, no. 10, 2012.

- [23] G. Alestas, Y. Akrami, F. Kazantzidis, S. Nesseris, and L. Perivolaropoulos, “ $f\sigma_8$ tension as a signature of Λ CDM breakdown,” *Physical Review D*, vol. 106, no. 10, p. 103525, 2022.
- [24] K.-H. Chae, “Dynamical Evidence for a Breakdown of Newtonian Gravity at Low Acceleration in Wide Binary Stars,” *The Astrophysical Journal*, vol. 953, no. 2, p. 153, 2023.
- [25] I. Banik et al., “No Evidence for a Variation of Newton’s Gravity in Wide Binaries from Gaia DR3,” *Monthly Notices of the Royal Astronomical Society*, 2024.
- [26] C. Pittordis and W. Sutherland, “Wide Binaries from Gaia DR3: Premature Conclusions on MOND,” *Monthly Notices of the Royal Astronomical Society*, vol. 518, no. 2, pp. 2720–2731, 2023.
- [27] S. Weinberg, “Phenomenological Lagrangians,” *Physica A: Statistical Mechanics and its Applications*, vol. 96, no. 1–2, pp. 327–340, 1979.
- [28] H. Georgi, “Effective Field Theory,” *Annual Review of Nuclear and Particle Science*, vol. 43, pp. 209–252, 1993.
- [29] D. B. Kaplan, “Five Lectures on Effective Field Theory,” *arXiv preprint* nucl-th/0510023, 1995 (revised 2005).

A Extended Cross-Validation: Wide-Binary Stars

This appendix reports an exploratory cross-system test of the SPARC-discovered interpolating function (Section 5.6). It is placed in the appendix, rather than the main results, because the underlying wide-binary dataset is the subject of an unresolved observational controversy (Chae [24] versus Banik et al. [25]). We present it as exploratory cross-validation, not as a definitive validation of the SPARC-discovered form.

Setup. Wide-binary stars (separations 10^3 – 10^4 AU) probe gravity in a low-acceleration regime ($g_N \lesssim a_0$) analogous to the outer parts of galaxies. The classical baseline is Newtonian two-body gravity; the predicted velocity boost under a MOND-type interpolating function $\nu(x)$ is $\gamma_v \equiv v_{\text{obs}}/v_{\text{Newtonian}} = \sqrt{\nu(x)}$, where $x = g_N/a_0$ depends on separation and total mass. We evaluate the SPARC-discovered ADCD form $\nu_{\text{ADCD}}(x) = \theta_0(\sqrt{1 + \theta_1/x} - 1) + 1$ with $\hat{\theta}_0 \approx 1.83$, $\hat{\theta}_1 \approx 0.262$ (the SPARC-fitted values, *with no refitting to wide-binary data*), against the canonical Simple MOND, Standard MOND, and RAR forms, for a $1.5 M_\odot$ total-mass system across $s \in [7, 30]$ kAU (the Hernandez/Chae anomaly zone); Figure 9 plots the resulting $\gamma_v(s)$ curves against Chae [24]’s observed band.

Result. Table 23 reports the predicted velocity boost and the σ -distance from Chae’s observed $\gamma_v = 1.20 \pm 0.06$ [24]. Three observations emerge:

1. *Simple MOND and RAR dramatically overshoot* the observed boost at $s > 15$ kAU, reaching 5–13 σ discrepancies at $s = 30$ kAU—clearly inconsistent with the Chae 2023 measurements at large separations.
2. The SPARC-discovered ADCD form ($c \approx 0.27$) is *more consistent* with observations at $s = 10$ –15 kAU than Simple/RAR, but *less consistent* than Standard MOND at the same separations, and itself diverges (+8.4 σ) at $s = 30$ kAU.
3. *No single interpolating function simultaneously optimizes both datasets.* The SPARC-optimized form with $c \approx 0.27$ does not unambiguously generalize to wide-binary kinematics.

Table 23: Wide-binary velocity boost $\gamma_v = v_{\text{obs}}/v_{\text{Newtonian}} = \sqrt{\nu(x)}$ predicted by the SPARC-discovered ADCD form ($c \approx 0.27$, *no refitting*) versus canonical Simple MOND, Standard MOND, and RAR forms, against the observed $\gamma_v = 1.20 \pm 0.06$ of Chae [24] for a $1.5 M_\odot$ total-mass system. $x = g_{\text{Newton}}/a_0$ is the dimensionless acceleration; σ -distance is $(\gamma_v^{\text{model}} - 1.20)/0.06$. Simple MOND and RAR overshoot dramatically at $s > 15$ kAU (5–13 σ); the SPARC-optimized ADCD form with $c \approx 0.27$ under-predicts at small s and over-predicts at large s . No single interpolating function simultaneously optimizes SPARC and the wide-binary kinematics.

s (kAU)	$x = g_N/a_0$	ADCD ($c \approx 0.27$)	Simple MOND	Standard MOND	RAR
<i>Predicted $\gamma_v = \sqrt{\nu(x)}$</i>					
7	1.51	1.073	1.206	1.090	1.189
10	0.74	1.139	1.328	1.147	1.316
15	0.33	1.273	1.521	1.230	1.513
20	0.19	1.418	1.696	1.300	1.691
30	0.08	1.705	2.005	1.415	2.002
<i>σ-distance from observed 1.20 ± 0.06</i>					
7	1.51	-2.1σ	$+0.1\sigma$	-1.8σ	-0.2σ
10	0.74	-1.0σ	$+2.1\sigma$	-0.9σ	$+1.9\sigma$
15	0.33	$+1.2\sigma$	$+5.3\sigma$	$+0.5\sigma$	$+5.2\sigma$
20	0.19	$+3.6\sigma$	$+8.3\sigma$	$+1.7\sigma$	$+8.2\sigma$
30	0.08	$+8.4\sigma$	$+13.4\sigma$	$+3.6\sigma$	$+13.4\sigma$




Figure 9: Wide-binary predicted velocity boost $\gamma_v = \sqrt{\nu(x)}$ versus separation s (lower axis) and dimensionless acceleration $x = g_N/a_0$ (upper axis), for the SPARC-discovered ADCD form ($c \approx 0.27$, no refit) versus Simple MOND, Standard MOND, and RAR. The shaded band is Chae [24]’s observed $\gamma_v = 1.20 \pm 0.06$. Simple MOND and RAR overshoot the observation at large separations; the ADCD form is closer at $s \sim 10$ –15 kAU but diverges beyond $s \sim 20$ kAU. No single form dominates across the full separation range.

Caveats. The wide-binary anomaly itself remains controversial. Banik et al. [25] attribute the observed velocity boost to selection effects and contamination rather than modified gravity; Pittordis & Sutherland [26] caution that the Gaia DR3 sample selection can mimic or suppress

a MOND-like signal. The discrepancies in Table 23 should therefore not be over-interpreted as discriminating between gravity theories. We report the cross-validation to document the generalization behaviour of the SPARC-discovered form—and to disclose honestly that it does *not* unambiguously transfer to an independent stellar-kinematic dataset, a limitation we flag in Section 7.