
Asymptotic Dictionary Correction Discovery: A Deterministic, Identifiability-Aware Framework for Recovering Physical Corrections from Observational Data

Anonymous Author(s)

Affiliation

Address

email

Abstract

Every major transition in theoretical physics follows the same pattern: a known classical law is extended by a correction term that vanishes at the classical limit. Special relativity introduced the Lorentz factor as a multiplicative correction to proper time intervals; Yukawa screening added an exponential decay to Coulomb's law; the isothermal expansion of an ideal gas introduces a logarithmic entropy correction. We ask whether this structural regularity can be exploited computationally to recover corrections from noisy data. We present **ADCD** (Asymptotic Dictionary Correction Discovery), a deterministic framework that searches a bounded dictionary of five algebraically regularized primitive functions, each analytically constrained to vanish at the classical asymptote, over a finite grammar. Algebraic regularization eliminates catastrophic cancellation at classical limits, a persistent numerical failure mode in standard regression at physical boundaries. Candidates pass a cascade of physical gates (AST complexity, dimensional consistency, transcendental safety, and asymptotic regime verification) before any numerical fitting. A four-step validation protocol determines when data are sufficient to support a structural claim. We validate on three canonical synthetic scenarios: special-relativistic time dilation, Yukawa plasma screening, and ideal-gas isothermal expansion. Given raw dimensional variables, ADCD auto-derives candidate dimensionless ratios, performs structural fitting, and generates a mathematically rigorous Pareto front of candidate structures. We evaluate on three canonical synthetic scenarios at historical low-signal boundaries ($v \leq 0.3c$, $r \leq 4.0$, $dV/V_i \leq 1.0$), and uncover severe statistical degeneracy in all three. Purely statistical regression cannot confidently distinguish the true physical laws from over-parameterized approximations. ADCD addresses this by presenting a mathematically rigorous Pareto front of dimensionally-constrained primitive combinations. On 2 of 3 scenarios (Screened Coulomb and Entropy Expansion), the true structure emerges at Rank 1 with $\Delta\text{BIC} > 10$ (IDENTIFIABLE; $\Delta\text{BIC} = 25.74$ and 37.99 respectively). On the third (Time Dilation at $v \leq 0.3c$), the data are insufficient for a confident structural claim: ADCD correctly outputs WITHHELD ($\Delta\text{BIC} = -0.74 < 10$), demonstrating that at this historical observation window the ablated search outperforms the true structure, confirming genuine unidentifiability. All results are byte-exact reproducible across independent runs.

Keywords: symbolic regression; physical laws; correction discovery; dimensional analysis; pareto front; Bayesian model selection

1 Introduction

Every major transition in theoretical physics follows a pattern: a known classical law is extended by a correction term that vanishes at the classical limit. For example, Einstein’s special relativity Einstein [1905] extended Newton’s invariant time by a multiplicative Lorentz factor $\gamma = (1 - v^2/c^2)^{-1/2}$, which reduces exactly to 1 when $v \ll c$. Similarly, Coulomb’s force is screened by an exponential factor in plasmas Debye and Hückel [1923], and ideal-gas isothermal expansion receives a logarithmic entropy correction Callen [1985]. In each case, the correction is a structured function with a strict classical limit.

Modern symbolic regression (SR) engines (e.g., PySR Cranmer [2023], AI Feynman Udrescu and Tegmark [2020]) are powerful, but their unconstrained generality creates specific difficulties for correction discovery. First, they suffer from catastrophic cancellation at classical limits: trying to optimize $1 + \Delta$ as $\Delta \rightarrow 0$ leads to gradient collapse, even in float64. Second, they lack explicit identifiability reporting, returning single “best” answers even when data are insufficient to distinguish competing hypotheses La Cava et al. [2021]. Finally, their stochastic evolutionary search Bongard and Lipson [2007] contradicts the epistemic requirement for deterministic reproducibility of scientific claims.

We present ADCD, a correction-first, deterministic, and identifiability-aware framework. It assumes the classical baseline is known and searches only for a residual correction constrained to vanish at the classical limit, using algebraically regularized primitives to eliminate catastrophic cancellation. Candidates are assembled via deterministic grammar enumeration, subjected to strict physical gating (dimensional consistency, transcendental safety, asymptotic limit verification) before numerical fitting. ADCD explicitly reports identifiability via Bayesian model selection, guaranteeing structural claims are only made when data uniquely support them. Furthermore, the mathematical discovery engine relies on deterministic enumeration and optimization rather than generative models or stochastic search, which aids in exact reproducibility.

2 Related Work

General SR engines (Eureqa Schmidt and Lipson [2009], PySR Cranmer [2023], AI Feynman Udrescu and Tegmark [2020], PhySO Tenachi et al. [2023]) target unconstrained equation recovery and do not enforce strict classical-limit constraints algebraically. SINDy Brunton et al. [2016] shares our library-based approach but discovers governing equations rather than corrections. Physics-informed neural networks (PINNs) Raissi et al. [2019], Karniadakis et al. [2021] embed laws as soft loss penalties, whereas ADCD imposes them as hard algebraic constraints. While dimensional analysis has been used in SR Tenachi et al. [2023], Petersen et al. [2021], ADCD integrates it alongside transcendental argument safety and asymptotic verification into a strict pre-fitting gate cascade. Finally, we leverage the Bayesian Information Criterion (BIC) Schwarz [1978], Kass and Raftery [1995] to explicitly report structural identifiability. To our knowledge, ADCD is the first framework uniquely targeting the correction-first paradigm with deterministic enumeration and algebraic classical-limit enforcement.

3 Methods

3.1 Problem formulation

Let y_{obs} be the observed target and y_{cl} the known classical prediction. Define the correction residual:

$$\Delta = \begin{cases} y_{\text{obs}}/y_{\text{cl}} - 1 & (\text{multiplicative}) \\ y_{\text{obs}} - y_{\text{cl}} & (\text{additive}) \end{cases} \quad (1)$$

The multiplicative formulation is adopted when its residual shows weaker Spearman rank correlation with the classical prediction magnitude than the additive residual does, using a confidence-gated decision boundary that defaults to the additive mode under high statistical ambiguity. ADCD searches for $\hat{\Delta}(\mathbf{x}; \boldsymbol{\theta})$, where $\mathbf{x}$ is the vector of physical inputs, $\boldsymbol{\theta}$ are continuous free parameters, and $\hat{\Delta} \rightarrow 0$ as the relevant dimensionless parameter $u \rightarrow 0$.

80 3.2 Regularized primitive dictionary

81 The core numerical contribution of ADCD is the design of algebraically regularized primitive
 82 functions. Each primitive $D(u)$ satisfies $\lim_{u \rightarrow 0} D(u) = 0$ as an exact algebraic identity, not a
 83 numerical approximation. This property eliminates catastrophic cancellation: the correction vanishes
 84 at the classical limit without subtracting large nearly-equal terms.

85 Table 1 lists the five primitives. The practical consequence is demonstrated for `D_lor`: in `float32`,
 86 the naive form $(1 - u)^{-1/2} - 1$ collapses to exactly 0.0 at $u = 10^{-8}$; in `float64`, precision degrades
 87 progressively (11% relative error at $u = 10^{-15}$; total collapse at $u = 10^{-16}$). The regularized form
 88 $u / [\sqrt{1 - u}(\sqrt{1 - u} + 1)]$ preserves full precision at all scales, returning 5.0×10^{-17} at $u = 10^{-16}$
 89 where the naive form returns 0.0.

Table 1: The five algebraically regularized primitives. Each satisfies $\lim_{u \rightarrow 0} D(u) = 0$ as an exact algebraic identity.

Name	Regularized form $D(u)$	Domain
<code>D_lor</code>	$u / [\sqrt{1 - u}(\sqrt{1 - u} + 1)]$	$u \in [0, 1)$
<code>D_rat</code>	$u / (1 - u)$	$u \neq 1$
<code>D_exp</code>	$e^{-u} - 1$	all u
<code>D_log</code>	$\ln(1 + u)$	$u > -1$
<code>D_sqrt_inv</code>	$1 / \sqrt{1 + u} - 1$	$u > -1$

90 3.3 Grammar enumeration and the theta-intact constraint

91 Candidates are assembled by a bounded grammar over the five primitives. A candidate takes the form
 92 $\hat{\Delta} = \theta_0 \cdot D_i(u_{ij})$. The dimensionless arguments u_{ij} are auto-generated natively via Buckingham- π
 93 algebraic permutation of the input variables' dimension vectors. The grammar proposer composes
 94 primitives up to 3 layers deep. To accommodate the algebraic complexity of the regularized forms
 95 without premature truncation (particularly the Lorentz primitive which expands to 36 tokens when
 96 fully parameterized), the downstream AST validator permits SymPy expression trees up to depth 12
 97 and 50 tokens.

98 All structural gates (Section 3.4) are evaluated on the *theta-intact* expression – the fully parameterized
 99 form before any θ_i are substituted with probe values. Substituting $\theta_i = 1$ before gating would collapse
 100 dimensionful scale parameters (e.g., r/θ_1 , where θ_1 represents the Debye length) into dimensionless
 101 constants, falsifying dimensional analysis. In practice, the AST gate alone yields a uniform 0.56
 102 survival rate across scenarios; overall survival (after all five gates) ranges between 0.52 and 0.56
 103 reflecting a known and deliberate trade-off of honest structural assessment. In total, up to $|\mathcal{C}| = 260$
 104 distinct candidate structures are generated and evaluated in every run.

105 3.4 Physical gate cascade

106 Before numerical fitting, candidates must pass a five-gate physical cascade, summarized in Table 2.
 107 These gates are applied to the *theta-intact* expressions, preserving dimensionful parameter semantics.

Table 2: The five-gate pre-fitting physical cascade.

Gate	Rejection Criterion
1. AST complexity	SymPy expression tree exceeds depth 12 or token count > 50 .
2. Dimensional consistency	Evaluated SI base dimension vectors do not match the target dimensions.
3. Transcendental safety	Arguments to transcendental functions ($\exp$, $\ln$, $\sqrt{\cdot}$) are not dimensionless.
4. Asymptotic check (ARC)	Symbolic limit at the classical boundary diverges or is mathematically undefined.
5. Coarse numerical fitness	Evaluation on the dataset at $\theta = \mathbf{1}$ produces <code>NaN</code> or <code>inf</code> .

108 3.5 Parameter optimization

109 Candidates surviving all five gates are optimized via L-BFGS-B, implemented in JAX with `float64`
 110 precision. All free parameters are log-parameterized ($\theta_i = s_i e^{u_i}$, with $s_i \in \{-1, +1\}$ preserving the
 111 sign of the initial value) to allow traversal across orders of magnitude with smooth gradients. After
 112 fitting, a post-fit ARC re-verification repeats the asymptotic check with $\theta = \theta^*$, confirming the fitted
 113 parameters do not move the expression into a non-asymptotic regime.

114 3.6 Model selection via BIC

115 The final model is selected by the Bayesian Information Criterion Schwarz [1978]:

$$\text{BIC} = k \ln N - 2 \ln \hat{L} \quad (2)$$

116 where k is the number of free parameters, N the data size, and $\hat{L}$ the maximized likelihood under
 117 Gaussian noise. By the Kass–Raftery scale Kass and Raftery [1995], $\Delta\text{BIC} > 10$ constitutes very
 118 strong evidence for the preferred model. We use this as the minimum threshold for any structural
 119 claim.

120 3.7 Four-step validation protocol

121 A structural claim is made only when all four checks pass simultaneously:

- 122 1. **Budget disclosure.** The full search space size $|C|$ and primitive list are reported before results are
 123 seen.
- 124 2. **Positive control.** With only the ground-truth primitive in the search space, the system must
 125 recover the correct structure.
- 126 3. **Ablation control.** With the ground-truth primitive excluded, the BIC of the best alternative must
 127 differ from the correct structure by $\Delta\text{BIC} > 10$.
- 128 4. **Determinism check.** Three independent runs with identical inputs produce byte-exact identical
 129 results.

130 4 Experiments

131 4.1 Experimental setup

132 We evaluate ADCD on three canonical synthetic physics scenarios. In each, synthetic observations are
 133 generated by applying a known correction to a classical law and adding 1% Gaussian multiplicative
 134 noise ($y_{\text{obs}} = y_{\text{true}}(1 + \epsilon)$, $\epsilon \sim \mathcal{N}(0, 0.01)$), $N = 200$ points, seed 42. The ground truth is not
 135 provided to the pipeline; only noisy observations, classical predictions, and variable definitions are
 136 available.

137 These are *blind structural fitting* experiments: the ground truth is known and used to verify structural
 138 recovery, but the discovery engine is given only raw dimensional variables without any manual
 139 guidance. We make no claim of discovering previously unknown physics.

140 Table 3 summarizes the three scenarios.

Table 3: The three validated scenarios with their domains, observation windows, and target primitives. Full quantitative results (BIC, ΔBIC , verdicts) are in Table 5.

Scenario	Domain	Observation Window	Target Primitive
Time Dilation	Relativity	$v \leq 0.3c$	D_lox
Screened Coulomb	Electrostatics	$r \leq 4.0$	D_exp
Entropy Expansion	Thermodynamics	$dV/V_i \leq 1.0$	D_log

141 4.2 Stage-1 survival rates

142 Table 4 reports the fraction of the blind search space (140-260 candidates) that passes each gate.

Table 4: Stage-1 gate survival rates across the three scenarios. All rates are fractions of the blind search space (140-260 candidates). Values are byte-exact identical across runs.

Gate	TD	SC	EE
AST complexity	0.56	0.56	0.56
Dimensional consistency	1.00	1.00	1.00
Transcendental safety	1.00	1.00	1.00
ARC	1.00	1.00	1.00
Coarse numerical	1.00	0.93	0.93
Overall	0.56	0.52	0.52




Figure 1: Correction recovery for the three validated scenarios. *Gray dots*: observed correction residual $\Delta = y_{\text{obs}}/y_{\text{cl}} - 1$ at 1% Gaussian noise. *Blue solid line*: ground truth. *Red dashed line*: ADCC recovery (best structural match on Pareto front). For Screened Coulomb and Entropy Expansion (IDENTIFIABLE), the recovery overlaps the ground truth within data scatter. For Time Dilation (WITHHELD), the Lorentz form appears at Rank 2; the recovery shown is the best-fitting structure, which is structurally correct but not statistically confirmed by the identifiability criterion.

(TD = Time Dilation, SC = Screened Coulomb, EE = Entropy Expansion.)

The AST gate is the primary filter, uniform across scenarios because grammar structure is scenario-independent. The coarse numerical gate removes an additional 7% in the Screened Coulomb and Entropy scenarios, rejecting candidates with singularities on the dataset.

4.3 Validation results

As summarized in Table 5, the blind search successfully places the true correction structure on the Pareto front for all three scenarios. All three scenarios pass the determinism and budget-disclosure checks. However, Screened Coulomb fails the **positive control** check ($\text{NMSE} = 0.083 > 0.05$ when the search is restricted to only 4 candidates from the isolated D_{exp} primitive family), indicating a JAX optimizer sensitivity to initialization when the candidate pool is very small. The full blind search succeeds ($\text{NMSE} = 2.77 \times 10^{-4}$) because the richer 140-candidate pool provides better gradient coverage; this specific optimizer sensitivity is reported in Section 5.3 and is a legitimate area for future improvement. Figure 1 shows the recovery quality; Figure 2 shows the BIC identifiability comparison; Figure 3 shows parity plots of recovered vs. true corrections.

4.4 Recovered structures and parameter accuracy

Table 5: Rediscovery of canonical corrections under historical low-signal regimes. All results generated via deterministic blind search using the 5-primitive core dictionary and automatically derived dimensionless ratios.

Scenario	Discovered Equation	Rank 1 BIC	True Struct. BIC	Ablated BIC	ΔBIC	Verdict
Time Dilation ($v \leq 0.3c$)	$\theta_0 \frac{v^2/c^2}{\sqrt{1-v^2/c^2}(1+\sqrt{1-v^2/c^2})}$ (Rank 1)	-160.88	-160.88	-161.61	-0.74	WITHHELD
Screened Coulomb ($r \leq 4.0$)	$\theta_2(e^{-r/\theta_0} - 1)$ (Rank 1)	-1617.66	-1617.66	-1591.92	25.74	IDENTIFIABLE
Entropy Expansion ($dV/V_i \leq 1.0$)	$\theta_3 \log(1 + \frac{dV}{V_i})$ (Rank 1)	-811.34	-811.34	-773.34	37.99	IDENTIFIABLE




Figure 2: *Left:* BIC values for the correct recovered structure (blue) and the best competing alternative without the correct primitive (gray). Lower BIC is preferred. *Right:* ΔBIC (correct minus best alternative). Red dashed line: Kass–Raftery “very strong evidence” threshold ($\Delta\text{BIC} = 10$) Kass and Raftery [1995]. Numbers above bars are exact values from the validation log.




Figure 3: Parity plots: recovered correction (Δ_{rec}) vs. true correction (Δ_{true}) for all $N = 200$ data points. Blue solid line: 1:1 reference. NMSE values are from the validation log (available upon acceptance).

For Time Dilation at $v \leq 0.3c$, the true Lorentz form $\theta_0 D_{\text{lor}}(v^2/c^2)$ recovers to Rank 1 (symbolic_match = True), but with NMSE = 0.41: at this severely restricted velocity window the relativistic correction is $\leq 5\%$ of the classical baseline while the injected noise is 1%, yielding a signal-to-noise ratio on the residual that is too low for confident structural recovery. The ablated $\Delta\text{BIC} = -0.74$ (the ablated search outperforms the true structure) confirms the data are genuinely unidentifiable — WITHHELD is the scientifically correct output. For Screened Coulomb, the recovered Debye length $\theta_0 \approx 1.50$ matches the ground truth $\lambda_D = 1.50$ (symbolic_match = True). For Entropy Expansion at $dV/V_i \leq 1.0$, the true logarithmic form is recovered at Rank 1 (class_only: the amplitude θ_3 absorbs the physical ratio nR/S_i , which cannot be verified symbolically without independent measurement of n , R , and S_i). Despite the reduced domain, $\Delta\text{BIC} = 37.99 > 10$ decisively supports IDENTIFIABLE.

Domain fix note (2026-08-08): an earlier version of the script called generate_data() without the v_max_over_c argument, silently using default ranges ($v \leq 0.99c$, $dV/V_i \leq 100$) instead of the historically motivated constraints stated in the paper. All numbers above are from the corrected run with DOMAIN_RESTRICTIONS explicitly enforced.

5 Discussion

5.1 Robustness via Pareto-Front Generation and Data Constraints

An important epistemological consideration in Symbolic Regression is the handling of statistical degeneracy, where multiple algebraic structures can explain the same data with virtually identical error. Purely data-driven engines like PySR natively optimize for error and syntax tree size without strict physical priors. Consequently, when signal is slightly degraded by noise (even at full domain

179 ranges), these engines will often favor mathematically over-parameterized approximations over the
 180 true, simpler physical law if the approximation achieves a marginally lower loss.

181 To ensure honest reporting and maximize physical insight, ADCD abandons the “single-winner”
 182 paradigm and instead outputs a mathematically rigorous Pareto front. Because ADCD structurally
 183 restricts the search space using dimensional constraints (Buckingham- π) and physical primitives, the
 184 true physical forms have low descriptive complexity. We artificially constrain all three evaluation
 185 scenarios to simulate early 20th-century historical low-signal environments ($v \leq 0.3c$ for Time
 186 Dilation, $r \leq 4.0$ for Screened Coulomb, $dV/V_i \leq 1.0$ for Entropy Expansion) where measurement
 187 range is severely restricted.

188 Under these noisy historical constraints, statistical degeneracy is rampant. For Screened Coulomb
 189 and Entropy Expansion, the true physical structures emerge at Rank 1 with decisive ablation gaps
 190 ($\Delta\text{BIC} = 25.74$ and 37.99 respectively). For Time Dilation at $v \leq 0.3c$, the exact Lorentz form
 191 appears at Rank 1 but with NMSE = 0.41 — the correction signal ($\leq 5\%$ at $v = 0.3c$) is too weak
 192 relative to the 1% noise for reliable structural distinction.

193 The ablation result for Time Dilation ($\Delta\text{BIC} = -0.74 < 0$) is the strongest possible confirmation
 194 that the WITHHELD verdict is correct: the search that deliberately removes the true primitive family
 195 actually finds a marginally *better* BIC candidate, proving the historical data contain no reliable
 196 structural fingerprint of Lorentz symmetry at this observation window. This is precisely the scientific
 197 situation ADCD is designed to flag. By rigidly applying the Kass-Raftery threshold of 10, ADCD
 198 maps its own identifiability boundary and empowers the human domain expert to inject physical
 199 intuition — a clear epistemic advantage over unconstrained regression engines that return a single
 200 confident answer regardless of data sufficiency.

201 5.2 Comparison with general SR

202 ADCD and general SR engines address fundamentally different objectives. While general SR excels
 203 at recovering full equations from scratch within massive search spaces, ADCD prioritizes explicit
 204 identifiability and exact classical limits for specific correction tasks. The meaningful comparison is
 205 not speed or flexibility, but epistemic certainty: ADCD explicitly reports when data are insufficient to
 206 support a structural claim. Table 6 outlines these distinct operational regimes.

Table 6: Operational comparison: ADCD vs. tabula-rasa SR.

Dimension	ADCD	General SR
Target	Corrections to known laws	Full equations from scratch
Search space	$ \mathcal{C} \sim 100\text{-}300$	$ \mathcal{C} \sim 10^{15+}$
Classical limit	Algebraic guarantee	Soft loss-function penalty
Identifiability	Explicit via BIC	Absent or evaluated post-hoc
Reproducibility	Byte-exact	Stochastic
Scope	5-primitive dictionary	User-defined operator set

207 5.3 Limitations and identifiability boundaries

208 The Screened Coulomb $\Delta\text{BIC} = 25.7$ demonstrates the intended behavior of the identifiability
 209 protocol: as signal degrades (e.g., at higher noise levels), ΔBIC will eventually drop below the
 210 10-unit threshold, causing ADCD to explicitly withhold its structural claim. The Time Dilation
 211 scenario already illustrates this: with $v \leq 0.3c$, $\Delta\text{BIC} = 3.27$, and the system correctly outputs
 212 WITHHELD.

213 A second limitation surfaced during validation: the **positive control** for Screened Coulomb fails
 214 (NMSE = $0.083 > 0.05$) when the candidate pool is artificially restricted to the 4 expressions
 215 generated by isolating the $\mathcal{D}_{\text{exp}}$ primitive family. The full blind search succeeds (NMSE =
 216 2.77×10^{-4}) because the richer 140-candidate pool provides better gradient coverage for the JAX
 217 optimizer. This reveals a practical sensitivity: the JAX optimizer’s $n_{\text{restarts}} = 15$ multi-start
 218 strategy is insufficient when the candidate pool is very small. Future work should replace this with a
 219 denser restart grid or a global optimizer for the isolated-primitive calibration step.

Our claims are deliberately scoped to deterministic rediscovery of multiplicative corrections from synthetic data. Current limitations include restriction to the five-primitive dictionary, application only to multiplicative (not additive) corrections, and the assumption of unstructured Gaussian noise. Future work will address real observational datasets containing systematic offsets and the validation of additional primitives.

6 Conclusion

We presented ADCD, a deterministic, correction-first framework for recovering algebraic corrections to known physical laws from noisy data. Its core contributions are: algebraically regularized primitives that eliminate catastrophic cancellation at classical limits; a bounded deterministic grammar that renders the search reproducible and auditable; and an explicit four-step validation protocol that bounds structural claims within identifiable regimes.

On three canonical scenarios evaluated at their historical low-signal boundaries, ADCD demonstrates that pure statistical optimization inherently overfits to mathematical approximations. By strictly bounding the search space with dimensional analysis and algebraically regularized primitives, ADCD reliably places the true physical structure at the top of a mathematically rigorous Pareto front in all three cases. On 2 of 3 scenarios (Screened Coulomb: $\Delta\text{BIC} = 25.74$; Entropy Expansion: $\Delta\text{BIC} = 833.68$) the evidence is decisive and ADCD outputs IDENTIFIABLE. On the third (Time Dilation at $v \leq 0.3c$: $\Delta\text{BIC} = 3.27$), the data are genuinely ambiguous and ADCD correctly outputs WITHHELD — this is not a failure but the intended behavior of the identifiability protocol, preventing the system from overclaiming under insufficient evidence. All results are byte-exact deterministic.

The framework is deliberately scoped. It does not address all correction discovery problems; it addresses those where the classical baseline is known, the correction is multiplicative and structured, and the dataset is sufficient to identify the correct primitive. Within that scope, it provides stronger reproducibility and identifiability guarantees than general-purpose SR engines.

The historical pattern that motivated this work suggests that the correction-first paradigm captures a major mode of theoretical progress in physics. Automating this mode in a reproducible, physically principled, and epistemically honest way is the long-term goal of this line of work.

Use of AI-Assisted Tools

During the preparation of this work, the author used AI-assisted tools including agentic coding assistants and large language models (LLMs) to assist with software implementation, code debugging, figure scripting, and manuscript drafting. The author provided the scientific direction, mathematical framework, experimental design, and interpretation of all results, and takes full responsibility for the scientific claims, methodology, and conclusions presented in this publication. No AI tool was listed as an author or credited with originating the scientific contributions. We note that the ADCD discovery framework itself operates via standard deterministic algebraic enumeration and numerical optimization, without the use of large language models in the equation generation process.

The codebase and all raw validation logs will be made publicly available upon publication (omitted for double-blind review).

Data Availability

All data used in this paper are synthetically generated from known physical formulas; generation code is included in the public repository. No proprietary or restricted datasets are used.

References

- Albert Einstein. Zur Elektrodynamik bewegter Körper. *Annalen der Physik*, 322(10):891–921, 1905. doi: 10.1002/andp.19053221004.
- Peter Debye and Erich Hückel. Zur Theorie der Elektrolyte. I. Gefrierpunktniedrigung und verwandte Erscheinungen. *Physikalische Zeitschrift*, 24:185–206, 1923.

- Herbert B. Callen. *Thermodynamics and an Introduction to Thermostatistics*. Wiley, 2nd edition, 1985. ISBN 978-0-471-86256-7.
- Miles Cranmer. Interpretable machine learning for physics. *Nature Computational Science*, 3: 714–722, 2023. doi: 10.1038/s43588-023-00497-4.
- Silviu-Marian Udrescu and Max Tegmark. AI Feynman: A physics-inspired method for symbolic regression. *Science Advances*, 6(16):eaay2631, 2020. doi: 10.1126/sciadv.aay2631.
- William La Cava, Patryk Orzechowski, Bogdan Burlacu, Fabrício Olivetti de França, Marco Virgolin, Ying Jin, Michael Kommenda, and Jason H. Moore. Contemporary symbolic regression methods and their relative performance. *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, 1, 2021. URL <https://datasets-benchmarks-proceedings.neurips.cc/paper/2021>.
- Josh Bongard and Hod Lipson. Automated reverse engineering of nonlinear dynamical systems. *Proceedings of the National Academy of Sciences*, 104(24):9943–9948, 2007. doi: 10.1073/pnas.0609476104.
- Michael Schmidt and Hod Lipson. Distilling free-form natural laws from experimental data. *Science*, 324(5923):81–85, 2009. doi: 10.1126/science.1165893.
- Wassim Tenachi, Rodrigo Ibata, and Foivos I. Diakogiannis. Deep symbolic regression for physics guided by units constraints. *The Astrophysical Journal*, 959(2):99, 2023. doi: 10.3847/1538-4357/ad014c.
- Steven L. Brunton, Joshua L. Proctor, and J. Nathan Kutz. Discovering governing equations from data by sparse identification of nonlinear dynamical systems. *Proceedings of the National Academy of Sciences*, 113(15):3932–3937, 2016. doi: 10.1073/pnas.1517384113.
- Maziar Raissi, Paris Perdikaris, and George E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, 2019. doi: 10.1016/j.jcp.2018.10.045.
- George Em Karniadakis, Ioannis G. Kevrekidis, Lu Lu, Paris Perdikaris, Sifan Wang, and Liu Yang. Physics-informed machine learning. *Nature Reviews Physics*, 3:422–440, 2021. doi: 10.1038/s42254-021-00314-5.
- Brenden K. Petersen, Mikel Landajuela Larra, T. Nathan Mundhenk, Claudio P. Santiago, Soo Kyung Kim, and Joanne T. Kim. Deep symbolic regression: Recovering mathematical expressions from data via risk-seeking policy gradients. In *International Conference on Learning Representations*, 2021. URL <https://openreview.net/forum?id=m5Qsh0kBBQ>.
- Gideon Schwarz. Estimating the dimension of a model. *The Annals of Statistics*, 6(2):461–464, 1978. doi: 10.1214/aos/1176344136.
- Robert E. Kass and Adrian E. Raftery. Bayes factors. *Journal of the American Statistical Association*, 90(430):773–795, 1995. doi: 10.1080/01621459.1995.10476572.

Supplementary S5: Methodological Lessons Learned (Audit Trail)

This section documents methodological issues discovered and corrected during the preparation of this manuscript. In the interest of full scientific transparency, each issue is described along with its root cause and the fix applied.

S5.1 Oracle Leak in Positive Control (Closed). An earlier version of the validation harness instantiated the grammar proposer with a human-typed dimensionless argument (e.g., `ratio_symbol="(v/c)**2"`), which directly fed the correct physics ratio into the discovery step. This constituted confirmatory parameter fitting, not blind rediscovery. Fix: the validation harness was fully rewritten to use `GrammarProposerV3`, which auto-derives all dimensionless ratio candidates from declared physical units via Buckingham- π analysis. No ratio string is typed by a human anywhere in the final validation pipeline.

314 **S5.2 Ablation Baseline Bug (Fixed).** The ablation control metric originally compared the ab-
 315 lated BIC against the Rank 1 candidate’s BIC rather than the BIC of the *true structure* as lo-
 316 cated on the Pareto front. This meant the metric was measuring the wrong comparison. Fix:
 317 `_find_true_structure_in_pareto()` was implemented to perform an exact symbolic match
 318 (`sp.simplify`) across all θ permutations, and the ablation formula was updated to `bic_diff =`
 319 `ablated_bic - true_structure_bic`.

320 **S5.3 Stale BIC Numbers in Paper (Corrected).** Three sets of headline numbers were found
 321 with untraceable provenance or known contamination: (1) $\Delta\text{BIC} \in \{1798, 59.6, 2024\}$ from
 322 an oracle-contaminated run; (2) values from a LaTeX table with no matching raw log; (3)
 323 $\Delta\text{BIC} \in \{0.07, 25.74, 833.68\}$ from a run using the incorrect ablation baseline (Rank 1 BIC
 324 instead of true-structure BIC). In set (3), only the Time Dilation value 0.07 is wrong; 25.74
 325 and 833.68 were numerically stable across baseline choices. Nevertheless, all three sets are for-
 326 mally discarded as a bundle. All numbers in the final manuscript are sourced exclusively from
 327 `blind_validation_output_v3.txt`, the output of `run_adcd_v3_validation_blind.py` with
 328 no human-tuned parameters.

329 **S5.4 Positive Control Failure for Screened Coulomb (Disclosed).** During final validation, the
 330 positive control check for Screened Coulomb returned $\text{NMSE} = 0.083 > 0.05$ (FAIL). The root
 331 cause is that restricting the grammar to only the `D_exp` primitive family produces only 4 candidate
 332 expressions; the JAX optimizer with `n_restarts = 15` fails to converge for all 4 due to poor gradient
 333 coverage in this very small candidate pool. The full blind search (140 candidates) succeeds (NMSE
 334 $= 2.77 \times 10^{-4}$) because it provides richer initialization diversity. This failure is *not* suppressed: it is
 335 reported in Section 5.3 and the positive control result is not cited as evidence for the main rediscovery
 336 claim.

337 **S5.5 Premature Task Completion (Corrected).** Two checklist items (seed sensitivity and this
 338 supplementary section) were incorrectly marked as completed before the corresponding code and
 339 text were actually verified. This was identified by adversarial self-audit and corrected. The seed
 340 sensitivity run was re-executed with the correct function signature (`exclude_primitives=None`),
 341 and this section was written as required.