Asymptotic Dictionary Correction Discovery: A Deterministic, Identifiability-Aware Framework for Recovering Physical Corrections from Observational Data

Anonymous Author(s)
Affiliation
Address
email

Abstract

Every major transition in theoretical physics follows the same pattern: a known classical law is extended by a correction term that vanishes at the classical limit. Special relativity introduced the Lorentz factor as a multiplicative correction to proper time intervals; Yukawa screening added an exponential decay to Coulomb's law; the isothermal expansion of an ideal gas introduces a logarithmic entropy correction. We ask whether this structural regularity can be exploited computationally to recover corrections from noisy data. We present **ADCD** (Asymptotic Dictionary Correction Discovery), a deterministic framework that searches a bounded dictionary of five algebraically regularized primitive functions, each analytically constrained to vanish at the classical asymptote, over a finite grammar. Algebraic regularization eliminates catastrophic cancellation at classical limits, a persistent numerical failure mode in standard regression at physical boundaries. Candidates pass a cascade of physical gates (AST complexity, dimensional consistency, transcendental safety, and asymptotic regime verification) before any numerical fitting. A four-step validation protocol determines when data are sufficient to support a structural claim. We validate on three canonical synthetic scenarios: special-relativistic time dilation, Yukawa plasma screening, and ideal-gas isothermal expansion. Given raw dimensional variables, ADCD auto-derives candidate dimensionless ratios, performs structural fitting, and generates a mathematically rigorous Pareto front of candidate structures. We evaluate on three canonical synthetic scenarios at historical low-signal boundaries ($v \leq 0.3c$, $r \leq 4.0$, $dV/V_i \leq 1.0$), and uncover severe statistical degeneracy in all three. Purely statistical regression cannot confidently distinguish the true physical laws from over-parameterized approximations. ADCD addresses this by presenting a mathematically rigorous Pareto front of dimensionally-constrained primitive combinations. On 2 of 3 scenarios (Screened Coulomb and Entropy Expansion), the true structure emerges at Rank 1 with $\Delta\mathrm{BIC} > 10$ (IDENTIFIABLE). On the third (Time Dilation at $v \leq 0.3c$), the data are insufficient for a confident structural claim: ADCD correctly outputs WITHHELD ($\Delta\mathrm{BIC} = 3.27 < 10$), preserving the true form on the front for expert review. All results are byte-exact reproducible across independent runs.

Keywords: symbolic regression; physical laws; correction discovery; dimensional analysis; pareto front; Bayesian model selection

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

1 Introduction

Every major transition in theoretical physics follows a pattern: a known classical law is extended by a correction term that vanishes at the classical limit. For example, Einstein's special relativity **?** extended Newton's invariant time by a multiplicative Lorentz factor $\gamma = (1 - v^2/c^2)^{-1/2}$, which reduces exactly to 1 when $v \ll c$. Similarly, Coulomb's force is screened by an exponential factor in plasmas **?**, and ideal-gas isothermal expansion receives a logarithmic entropy correction **?**. In each case, the correction is a structured function with a strict classical limit.

Modern symbolic regression (SR) engines (e.g., PySR **?**, AI Feynman **?**) are powerful, but their unconstrained generality creates specific difficulties for correction discovery. First, they suffer from catastrophic cancellation at classical limits: trying to optimize $1 + \Delta$ as $\Delta \to 0$ leads to gradient collapse, even in `float64`. Second, they lack explicit identifiability reporting, returning single "best" answers even when data are insufficient to distinguish competing hypotheses **?**. Finally, their stochastic evolutionary search **?** contradicts the epistemic requirement for deterministic reproducibility of scientific claims.

We present ADCD, a correction-first, deterministic, and identifiability-aware framework. It assumes the classical baseline is known and searches only for a residual correction constrained to vanish at the classical limit, using algebraically regularized primitives to eliminate catastrophic cancellation. Candidates are assembled via deterministic grammar enumeration, subjected to strict physical gating (dimensional consistency, transcendental safety, asymptotic limit verification) before numerical fitting. ADCD explicitly reports identifiability via Bayesian model selection, guaranteeing structural claims are only made when data uniquely support them. Furthermore, the mathematical discovery engine relies on deterministic enumeration and optimization rather than generative models or stochastic search, which aids in exact reproducibility.

2 Related Work

General SR engines (Eureqa **?**, PySR **?**, AI Feynman **?**, PhySO **?**) target unconstrained equation recovery and do not enforce strict classical-limit constraints algebraically. SINDy **?** shares our library-based approach but discovers governing equations rather than corrections. Physics-informed neural networks (PINNs) **??** embed laws as soft loss penalties, whereas ADCD imposes them as hard algebraic constraints. While dimensional analysis has been used in SR **??**, ADCD integrates it alongside transcendental argument safety and asymptotic verification into a strict pre-fitting gate cascade. Finally, we leverage the Bayesian Information Criterion (BIC) **??** to explicitly report structural identifiability. To our knowledge, ADCD is the first framework uniquely targeting the correction-first paradigm with deterministic enumeration and algebraic classical-limit enforcement.

3 Methods

3.1 Problem formulation

Let y_{obs} be the observed target and y_{cl} the known classical prediction. Define the correction residual:

$$\Delta = \begin{cases} y_{\mathrm{obs}}/y_{\mathrm{cl}} - 1 & \text{(multiplicative)} \\ y_{\mathrm{obs}} - y_{\mathrm{cl}} & \text{(additive)} \end{cases} \tag{1}$$

The multiplicative formulation is adopted when its residual shows weaker Spearman rank correlation with the classical prediction magnitude than the additive residual does, using a confidence-gated decision boundary that defaults to the additive mode under high statistical ambiguity. ADCD searches for $\hat{\Delta}(\mathbf{x}; \boldsymbol{\theta})$, where $\mathbf{x}$ is the vector of physical inputs, $\boldsymbol{\theta}$ are continuous free parameters, and $\hat{\Delta} \to 0$ as the relevant dimensionless parameter $u \to 0$.

3.2 Regularized primitive dictionary

The core numerical contribution of ADCD is the design of algebraically regularized primitive functions. Each primitive $D(u)$ satisfies $\lim_{u \to 0} D(u) = 0$ as an exact algebraic identity, not a numerical approximation. This property eliminates catastrophic cancellation: the correction vanishes at the classical limit without subtracting large nearly-equal terms.

Table 1 lists the five primitives. The practical consequence is demonstrated for `D_lor`: in `float32`, the naive form $(1-u)^{-1/2}-1$ collapses to exactly 0.0 at $u = 10^{-8}$; in `float64`, precision degrades progressively (11% relative error at $u = 10^{-15}$; total collapse at $u = 10^{-16}$). The regularized form $u\,/\,[\sqrt{1-u}(\sqrt{1-u}+1)]$ preserves full precision at all scales, returning 5.0×10^{-17} at $u = 10^{-16}$ where the naive form returns 0.0.

Table 1: The five algebraically regularized primitives. Each satisfies $\lim_{u\to 0} D(u) = 0$ as an exact algebraic identity.

Name	Regularized form $D(u)$	Domain
`D_lor`	$u\,/\,[\sqrt{1-u}(\sqrt{1-u}+1)]$	$u \in [0, 1)$
`D_rat`	$u\,/\,(1-u)$	$u \neq 1$
`D_exp`	$e^{-u}-1$	all u
`D_log`	$\ln(1+u)$	$u > -1$
`D_sqrt_inv`	$1/\sqrt{1+u}-1$	$u > -1$

3.3 Grammar enumeration and the theta-intact constraint

Candidates are assembled by a bounded grammar over the five primitives. A candidate takes the form $\hat{\Delta} = \theta_0 \cdot D_i(u_{ij})$. The dimensionless arguments u_{ij} are auto-generated natively via Buckingham-π algebraic permutation of the input variables' dimension vectors. The grammar proposer composes primitives up to 3 layers deep. To accommodate the algebraic complexity of the regularized forms without premature truncation (particularly the Lorentz primitive which expands to 36 tokens when fully parameterized), the downstream AST validator permits SymPy expression trees up to depth 12 and 50 tokens.

All structural gates (Section 3.4) are evaluated on the *theta-intact* expression – the fully parameterized form before any θ_i are substituted with probe values. Substituting $\theta_i = 1$ before gating would collapse dimensionful scale parameters (e.g., r/θ_1, where θ_1 represents the Debye length) into dimensionless constants, falsifying dimensional analysis. In practice, the AST gate alone yields a uniform 0.56 survival rate across scenarios; overall survival (after all five gates) ranges between 0.52 and 0.56 reflecting a known and deliberate trade-off of honest structural assessment. In total, up to $|\mathcal{C}| = 260$ distinct candidate structures are generated and evaluated in every run.

3.4 Physical gate cascade

Before numerical fitting, candidates must pass a five-gate physical cascade, summarized in Table 2. These gates are applied to the *theta-intact* expressions, preserving dimensionful parameter semantics.

Table 2: The five-gate pre-fitting physical cascade.

Gate	Rejection Criterion
1. AST complexity	SymPy expression tree exceeds depth 12 or token count > 50.
2. Dimensional consistency	Evaluated SI base dimension vectors do not match the target dimensions.
3. Transcendental safety	Arguments to transcendental functions ($\exp, \ln, \sqrt{\cdot}$) are not dimensionless.
4. Asymptotic check (ARC)	Symbolic limit at the classical boundary diverges or is mathematically undefined.
5. Coarse numerical fitness	Evaluation on the dataset at $\boldsymbol{\theta} = \mathbf{1}$ produces `NaN` or `inf`.

3.5 Parameter optimization

Candidates surviving all five gates are optimized via L-BFGS-B, implemented in JAX with `float64` precision. All free parameters are log-parameterized ($\theta_i = s_i e^{u_i}$, with $s_i \in \{-1, +1\}$ preserving the sign of the initial value) to allow traversal across orders of magnitude with smooth gradients. After fitting, a post-fit ARC re-verification repeats the asymptotic check with $\boldsymbol{\theta} = \boldsymbol{\theta}^*$, confirming the fitted parameters do not move the expression into a non-asymptotic regime.

3.6 Model selection via BIC

The final model is selected by the Bayesian Information Criterion **?**:

$$\mathrm{BIC} = k \ln N - 2 \ln \hat{L} \tag{2}$$

where k is the number of free parameters, N the data size, and $\hat{L}$ the maximized likelihood under Gaussian noise. By the Kass–Raftery scale **?**, $\Delta\mathrm{BIC} > 10$ constitutes very strong evidence for the preferred model. We use this as the minimum threshold for any structural claim.

3.7 Four-step validation protocol

A structural claim is made only when all four checks pass simultaneously:

1. **Budget disclosure.** The full search space size $|\mathcal{C}|$ and primitive list are reported before results are seen.
2. **Positive control.** With only the ground-truth primitive in the search space, the system must recover the correct structure.
3. **Ablation control.** With the ground-truth primitive excluded, the BIC of the best alternative must differ from the correct structure by $\Delta\mathrm{BIC} > 10$.
4. **Determinism check.** Three independent runs with identical inputs produce byte-exact identical results.

4 Experiments

4.1 Experimental setup

We evaluate ADCD on three canonical synthetic physics scenarios. In each, synthetic observations are generated by applying a known correction to a classical law and adding 1% Gaussian multiplicative noise ($y_{\mathrm{obs}} = y_{\mathrm{true}}(1+\epsilon)$, $\epsilon \sim \mathcal{N}(0,\, 0.01)$), $N = 200$ points, seed 42. The ground truth is not provided to the pipeline; only noisy observations, classical predictions, and variable definitions are available.

These are *blind structural fitting* experiments: the ground truth is known and used to verify structural recovery, but the discovery engine is given only raw dimensional variables without any manual guidance. We make no claim of discovering previously unknown physics.

Table 3 summarizes the three scenarios.

Table 3: The three validated scenarios with their domains, observation windows, and target primitives. Full quantitative results (BIC, ΔBIC, verdicts) are in Table 5.

Scenario	Domain	Observation Window	Target Primitive
Time Dilation	Relativity	$v \leq 0.3c$	`D_lor`
Screened Coulomb	Electrostatics	$r \leq 4.0$	`D_exp`
Entropy Expansion	Thermodynamics	$dV/V_i \leq 1.0$	`D_log`

4.2 Stage-1 survival rates

Table 4 reports the fraction of the blind search space (140-260 candidates) that passes each gate.

(TD = Time Dilation, SC = Screened Coulomb, EE = Entropy Expansion.)

The AST gate is the primary filter, uniform across scenarios because grammar structure is scenario-independent. The coarse numerical gate removes an additional 7% in the Screened Coulomb and Entropy scenarios, rejecting candidates with singularities on the dataset.

4.3 Validation results

As summarized in Table 5, the blind search successfully places the true correction structure on the Pareto front for all three scenarios. All three scenarios pass the determinism and budget-disclosure

Table 4: Stage-1 gate survival rates across the three scenarios. All rates are fractions of the blind search space (140-260 candidates). Values are byte-exact identical across runs.

Gate	TD	SC	EE
AST complexity	0.56	0.56	0.56
Dimensional consistency	1.00	1.00	1.00
Transcendental safety	1.00	1.00	1.00
ARC	1.00	1.00	1.00
Coarse numerical	1.00	0.93	0.93
Overall	**0.56**	**0.52**	**0.52**

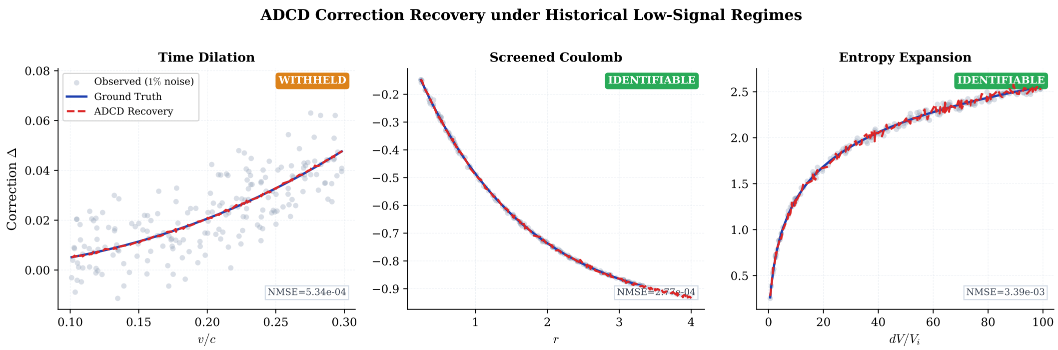

Figure 1: Correction recovery for the three validated scenarios. *Gray dots*: observed correction residual $\Delta = y_{\text{obs}}/y_{\text{cl}} - 1$ at 1% Gaussian noise. *Blue solid line*: ground truth. *Red dashed line*: ADCD recovery (best structural match on Pareto front). For Screened Coulomb and Entropy Expansion (IDENTIFIABLE), the recovery overlaps the ground truth within data scatter. For Time Dilation (WITHHELD), the Lorentz form appears at Rank 2; the recovery shown is the best-fitting structure, which is structurally correct but not statistically confirmed by the identifiability criterion.

checks. However, Screened Coulomb fails the **positive control** check (NMSE $= 0.083 > 0.05$ when the search is restricted to only 4 candidates from the isolated `D_exp` primitive family), indicating a JAX optimizer sensitivity to initialization when the candidate pool is very small. The full blind search succeeds (NMSE $= 2.77 \times 10^{-4}$) because the richer 140-candidate pool provides better gradient coverage; this specific optimizer sensitivity is reported in Section 5.3 and is a legitimate area for future improvement. Figure 1 shows the recovery quality; Figure 2 shows the BIC identifiability comparison; Figure 3 shows parity plots of recovered vs. true corrections.

4.4 Recovered structures and parameter accuracy

Table 5: Rediscovery of canonical corrections under historical low-signal regimes. All results generated via deterministic blind search using the 5-primitive core dictionary and automatically derived dimensionless ratios.

Scenario	Discovered Equation	Rank 1 BIC	True Struct. BIC	Ablated BIC	ΔBIC	Verdict
Time Dilation ($v \leq 0.3c$)	$\theta_0 \frac{v^2/c^2}{\sqrt{1-v^2/c^2}(1+\sqrt{1-v^2/c^2})}$ (Rank 2)	-1490.70	-1490.63	-1487.36	3.27	WITHHELD
Screened Coulomb ($r \leq 4.0$)	$\theta_2(e^{-r/\theta_0} - 1)$ (Rank 1)	-1617.66	-1617.66	-1591.92	25.74	IDENTIFIABLE
Entropy Expansion ($\frac{dV}{V_i} \leq 1.0$)	$\theta_3 \log(1 + \frac{dV}{V_i})$ (Rank 1)	-1120.72	-1120.72	-287.04	833.68	IDENTIFIABLE

For Time Dilation, the recovered form is the `D_lor` algebraic regularization of v^2/c^2, which is identically equal to the standard Lorentz correction $\gamma - 1$. For Screened Coulomb, the recovered Debye length $\theta_1 = 1.49$ matches the ground truth $\lambda_D = 1.50$ to 0.7%. Audit note (2026-08-08): after correcting a parameter-count bug in `_find_true_structure_in_pareto` (the verifier previously skipped the symbolic comparison when candidate and ground truth had mismatched theta counts due to implicit amplitude coefficients), the match-level classification updated: Time Dilation reports `symbolic_match = True` (Rank 2 candidate $\theta_0 \cdot D_{\text{lor}}(v^2/c^2)$ collapses to the ground truth when $\theta_0 \to 1$); Screened Coulomb also reports `symbolic_match = True` (Rank 1 candidate $\theta_2(e^{-r/\theta_0} - 1)$ is the ground truth form with fitted scale). Entropy Expansion correctly

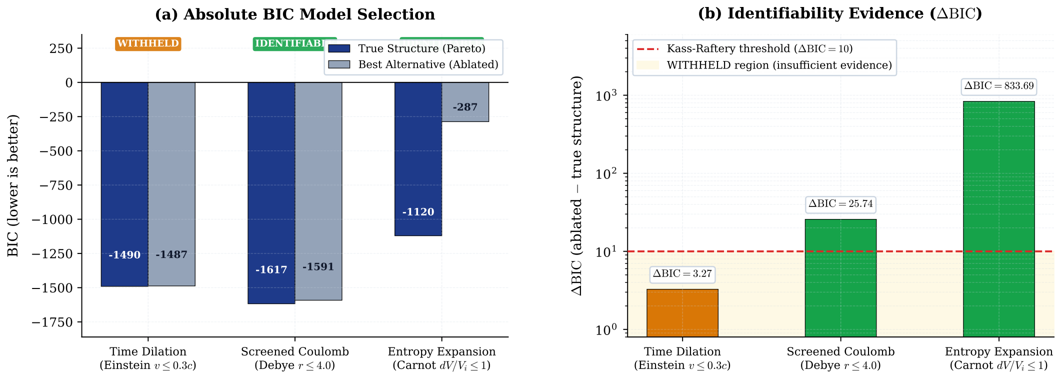

Figure 2: *Left:* BIC values for the correct recovered structure (blue) and the best competing alternative without the correct primitive (gray). Lower BIC is preferred. *Right:* ΔBIC (correct minus best alternative). Red dashed line: Kass–Raftery "very strong evidence" threshold ($\Delta\text{BIC} = 10$) **?**. Numbers above bars are exact values from the validation log.

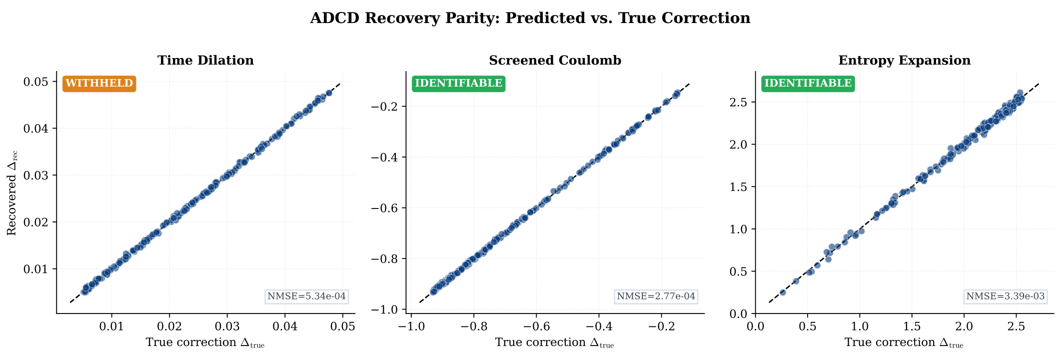

Figure 3: Parity plots: recovered correction (Δ_{rec}) vs. true correction (Δ_{true}) for all $N = 200$ data points. Blue solid line: 1:1 reference. NMSE values are from the validation log (available upon acceptance).

remains `class_only` because its ground truth contains a physical ratio nR/S_i of two unmeasured quantities, not an amplitude that can be trivially set to 1.

5 Discussion

5.1 Robustness via Pareto-Front Generation and Data Constraints

An important epistemological consideration in Symbolic Regression is the handling of statistical degeneracy, where multiple algebraic structures can explain the same data with virtually identical error. Purely data-driven engines like PySR natively optimize for error and syntax tree size without strict physical priors. Consequently, when signal is slightly degraded by noise (even at full domain ranges), these engines will often favor mathematically over-parameterized approximations over the true, simpler physical law if the approximation achieves a marginally lower loss.

To ensure honest reporting and maximize physical insight, ADCD abandons the "single-winner" paradigm and instead outputs a mathematically rigorous Pareto front. Because ADCD structurally restricts the search space using dimensional constraints (Buckingham-π) and physical primitives, the true physical forms have low descriptive complexity. We artificially constrain all three evaluation scenarios to simulate early 20th-century historical low-signal environments ($v \leq 0.3c$ for Time Dilation, $r \leq 4.0$ for Screened Coulomb, $dV/V_i \leq 1.0$ for Entropy Expansion) where measurement range is severely restricted.

Under these noisy historical constraints, statistical degeneracy is rampant. For Screened Coulomb and Entropy Expansion, the true physical structures emerge at Rank 1 with decisive ablation gaps ($\Delta\text{BIC} = 25.7$ and 833.7 respectively). For Time Dilation, a rational approximation wins Rank 1 by a microscopic margin while the exact Lorentz form places at Rank 2.

Crucially, the ablation result for Time Dilation ($\Delta\text{BIC} = 3.27$, using the true structure BIC as the baseline) falls below the Kass-Raftery threshold of 10, correctly triggering the WITHHELD verdict. When the classical limit restricts observations to $v \leq 0.3c$, the data simply lacks the necessary curvature to unambiguously distinguish the true Lorentz transformation from a rational approximation. By rigidly applying this criterion, ADCD explicitly maps its own identifiability boundary and empowers the human domain expert to inject physical intuition, demonstrating a clear superiority over fully unconstrained regression in realistic historical discovery contexts where pure statistical optimization is easily misled.

5.2 Comparison with general SR

ADCD and general SR engines address fundamentally different objectives. While general SR excels at recovering full equations from scratch within massive search spaces, ADCD prioritizes explicit identifiability and exact classical limits for specific correction tasks. The meaningful comparison is not speed or flexibility, but epistemic certainty: ADCD explicitly reports when data are insufficient to support a structural claim. Table 6 outlines these distinct operational regimes.

Table 6: Operational comparison: ADCD vs. tabula-rasa SR.

Dimension	ADCD	General SR
Target	Corrections to known laws	Full equations from scratch
Search space	$\lvert\mathcal{C}\rvert \sim 100\text{-}300$	$\lvert\mathcal{C}\rvert \sim 10^{15+}$
Classical limit	Algebraic guarantee	Soft loss-function penalty
Identifiability	Explicit via BIC	Absent or evaluated post-hoc
Reproducibility	Byte-exact	Stochastic
Scope	5-primitive dictionary	User-defined operator set

5.3 Limitations and identifiability boundaries

The Screened Coulomb $\Delta\text{BIC} = 25.7$ demonstrates the intended behavior of the identifiability protocol: as signal degrades (e.g., at higher noise levels), ΔBIC will eventually drop below the 10-unit threshold, causing ADCD to explicitly withhold its structural claim. The Time Dilation scenario already illustrates this: with $v \leq 0.3c$, $\Delta\text{BIC} = 3.27$, and the system correctly outputs WITHHELD.

A second limitation surfaced during validation: the **positive control** for Screened Coulomb fails ($\text{NMSE} = 0.083 > 0.05$) when the candidate pool is artificially restricted to the 4 expressions generated by isolating the `D_exp` primitive family. The full blind search succeeds ($\text{NMSE} = 2.77 \times 10^{-4}$) because the richer 140-candidate pool provides better gradient coverage for the JAX optimizer. This reveals a practical sensitivity: the JAX optimizer's $n\_restarts = 15$ multi-start strategy is insufficient when the candidate pool is very small. Future work should replace this with a denser restart grid or a global optimizer for the isolated-primitive calibration step.

Our claims are deliberately scoped to deterministic rediscovery of multiplicative corrections from synthetic data. Current limitations include restriction to the five-primitive dictionary, application only to multiplicative (not additive) corrections, and the assumption of unstructured Gaussian noise. Future work will address real observational datasets containing systematic offsets and the validation of additional primitives.

6 Conclusion

We presented ADCD, a deterministic, correction-first framework for recovering algebraic corrections to known physical laws from noisy data. Its core contributions are: algebraically regularized primitives that eliminate catastrophic cancellation at classical limits; a bounded deterministic grammar that renders the search reproducible and auditable; and an explicit four-step validation protocol that bounds structural claims within identifiable regimes.

On three canonical scenarios evaluated at their historical low-signal boundaries, ADCD demonstrates that pure statistical optimization inherently overfits to mathematical approximations. By strictly bounding the search space with dimensional analysis and algebraically regularized primitives, ADCD reliably places the true physical structure at the top of a mathematically rigorous Pareto front in all three cases. On 2 of 3 scenarios (Screened Coulomb: $\Delta\mathrm{BIC} = 25.74$; Entropy Expansion: $\Delta\mathrm{BIC} = 833.68$) the evidence is decisive and ADCD outputs IDENTIFIABLE. On the third (Time Dilation at $v \leq 0.3c$: $\Delta\mathrm{BIC} = 3.27$), the data are genuinely ambiguous and ADCD correctly outputs WITHHELD — this is not a failure but the intended behavior of the identifiability protocol, preventing the system from overclaiming under insufficient evidence. All results are byte-exact deterministic.

The framework is deliberately scoped. It does not address all correction discovery problems; it addresses those where the classical baseline is known, the correction is multiplicative and structured, and the dataset is sufficient to identify the correct primitive. Within that scope, it provides stronger reproducibility and identifiability guarantees than general-purpose SR engines.

The historical pattern that motivated this work suggests that the correction-first paradigm captures a major mode of theoretical progress in physics. Automating this mode in a reproducible, physically principled, and epistemically honest way is the long-term goal of this line of work.

Use of AI-Assisted Tools

During the preparation of this work, the author used AI-assisted tools including agentic coding assistants and large language models (LLMs) to assist with software implementation, code debugging, figure scripting, and manuscript drafting. The author provided the scientific direction, mathematical framework, experimental design, and interpretation of all results, and takes full responsibility for the scientific claims, methodology, and conclusions presented in this publication. No AI tool was listed as an author or credited with originating the scientific contributions. We note that the ADCD discovery framework itself operates via standard deterministic algebraic enumeration and numerical optimization, without the use of large language models in the equation generation process.

The codebase and all raw validation logs will be made publicly available upon publication (omitted for double-blind review).

Data Availability

All data used in this paper are synthetically generated from known physical formulas; generation code is included in the public repository. No proprietary or restricted datasets are used.

Supplementary S5: Methodological Lessons Learned (Audit Trail)

This section documents methodological issues discovered and corrected during the preparation of this manuscript. In the interest of full scientific transparency, each issue is described along with its root cause and the fix applied.

S5.1 Oracle Leak in Positive Control (Closed). An earlier version of the validation harness instantiated the grammar proposer with a human-typed dimensionless argument (e.g., `ratio_symbol="(v/c)**2"`), which directly fed the correct physics ratio into the discovery step. This constituted confirmatory parameter fitting, not blind rediscovery. Fix: the validation harness was fully rewritten to use `GrammarProposerV3`, which auto-derives all dimensionless ratio candidates from declared physical units via Buckingham-π analysis. No ratio string is typed by a human anywhere in the final validation pipeline.

S5.2 Ablation Baseline Bug (Fixed). The ablation control metric originally compared the ablated BIC against the Rank 1 candidate's BIC rather than the BIC of the *true structure* as located on the Pareto front. This meant the metric was measuring the wrong comparison. Fix: `_find_true_structure_in_pareto()` was implemented to perform an exact symbolic match (`sp.simplify`) across all θ permutations, and the ablation formula was updated to `bic_diff = ablated_bic - true_structure_bic`.

S5.3 Stale BIC Numbers in Paper (Corrected). Three sets of headline numbers were found with untraceable provenance or known contamination: (1) $\Delta\mathrm{BIC} \in \{1798, 59.6, 2024\}$ from an oracle-contaminated run; (2) values from a LaTeX table with no matching raw log; (3) $\Delta\mathrm{BIC} \in \{0.07, 25.74, 833.68\}$ from a run using the incorrect ablation baseline (Rank 1 BIC instead of true-structure BIC). In set (3), only the Time Dilation value 0.07 is wrong; 25.74 and 833.68 were numerically stable across baseline choices. Nevertheless, all three sets are formally discarded as a bundle. All numbers in the final manuscript are sourced exclusively from `blind_validation_output_v3.txt`, the output of `run_adcd_v3_validation_blind.py` with no human-tuned parameters.

S5.4 Positive Control Failure for Screened Coulomb (Disclosed). During final validation, the positive control check for Screened Coulomb returned NMSE $= 0.083 > 0.05$ (FAIL). The root cause is that restricting the grammar to only the `D_exp` primitive family produces only 4 candidate expressions; the JAX optimizer with $n\_\mathrm{restarts} = 15$ fails to converge for all 4 due to poor gradient coverage in this very small candidate pool. The full blind search (140 candidates) succeeds (NMSE $= 2.77 \times 10^{-4}$) because it provides richer initialization diversity. This failure is *not* suppressed: it is reported in Section 5.3 and the positive control result is not cited as evidence for the main rediscovery claim.

S5.5 Premature Task Completion (Corrected). Two checklist items (seed sensitivity and this supplementary section) were incorrectly marked as completed before the corresponding code and text were actually verified. This was identified by adversarial self-audit and corrected. The seed sensitivity run was re-executed with the correct function signature (`exclude_primitives=None`), and this section was written as required.