
Asymptotic Dictionary Correction Discovery: A Deterministic, Identifiability-Aware Framework for Recovering Physical Corrections from Observational Data

Anonymous Author(s)

Affiliation

Address

email

Abstract

We present **ADCD** (Asymptotic Dictionary Correction Discovery), a deterministic framework for recovering structured corrections to known physical laws from noisy data. Given a classical baseline, ADCD searches for the residual correction that vanishes exactly at the classical limit, using three interlocking mechanisms: algebraically regularized primitives that enforce asymptotic vanishing as an exact algebraic identity; a five-dimensional (MLT Θ Q) Buckingham- π engine that auto-derives dimensionless ratio arguments without manual specification; and a Phenomenon-Specific Taxonomy encoding peer-reviewed functional priors for each physics domain. A four-step validation protocol—budget disclosure, positive control, BIC-based ablation control, and determinism check—issues one of two verdicts: IDENTIFIABLE or WITHHELD. On three canonical scenarios at historical low-signal windows, Screened Coulomb ($\Delta\text{BIC} = 30.65$, exact symbolic match) and Entropy Expansion ($\Delta\text{BIC} = 14.58$, class-level match) are IDENTIFIABLE; Time Dilation at $v \leq 0.3c$ is WITHHELD because the 4.8% correction is indistinguishable from 1% noise—the intended, epistemically honest behavior. All results are byte-exact deterministic.

Keywords: symbolic regression; correction discovery; dimensional analysis; Bayesian prior; model selection; identifiability

1 Introduction

Physical laws frequently take a correction-first form: a known classical baseline extended by a term that vanishes at the classical limit Einstein [1905], Debye and Hückel [1923], Callen [1985]. Recovering that correction from observational data is distinct from full equation discovery: the baseline is fixed, only the residual must be identified, and that residual must vanish exactly where classical physics holds.

General-purpose symbolic regression (SR) engines such as PySR Cranmer [2023], AI Feynman Udrescu and Tegmark [2020], and PhySO Tenachi et al. [2023] excel at equation discovery from scratch but face three difficulties in the correction-first setting: (i) *catastrophic cancellation*: floating-point precision collapses as $u \rightarrow 0$, corrupting optimizer feedback IEEE Task P754 [2019]; (ii) *stochastic irreproducibility*: evolutionary search cannot guarantee byte-identical results across runs Bongard and Lipson [2007], Schmidt and Lipson [2009]; and (iii) *silent overclaiming*: a single best-fit equation is returned regardless of whether the data carry sufficient power to distinguish competing structures La Cava et al. [2021]. ADCD addresses each failure mode in turn: alge-

braically regularized primitives eliminate cancellation; deterministic grammar enumeration ensures reproducibility; explicit BIC-based identifiability gating prevents overclaiming.

2 Related Work

General SR engines (Eureqa Schmidt and Lipson [2009], PySR Cranmer [2023], AI Feynman Udrescu and Tegmark [2020], Udrescu et al. [2020], PhySO Tenachi et al. [2023]) target unconstrained equation recovery and do not enforce classical-limit constraints algebraically. SINDy Brunton et al. [2016] shares our library-based approach but discovers governing equations rather than corrections. PINNs Raissi et al. [2019], Karniadakis et al. [2021] embed physical laws as soft loss penalties; ADCD imposes them as hard algebraic identities before any numerical computation. AI-Descartes Cornelio et al. [2023] combines SR with axiomatic theory derivation, but does not address catastrophic cancellation or provide identifiability gating. While dimensional analysis has been integrated into SR Tenachi et al. [2023], Petersen et al. [2021], ADCD extends this with a full five-dimensional (MLTΘQ) Buckingham- π engine, transcendental argument safety, and BIC-based identifiability reporting in a single pre-fitting gate cascade. To our knowledge, ADCD is the first framework to combine deterministic grammar enumeration, algebraic classical-limit enforcement, phenomenonon-specific Bayesian priors, and explicit identifiability gating for the correction-first paradigm.

3 Methods

3.1 Problem formulation

Let y_{obs} and y_{cl} be the observed target and classical prediction. The correction residual is:

$$\Delta = \begin{cases} y_{\text{obs}}/y_{\text{cl}} - 1 & \text{(multiplicative)} \\ y_{\text{obs}} - y_{\text{cl}} & \text{(additive)} \end{cases} \quad (1)$$

The multiplicative form is selected when its residual shows lower Spearman rank correlation with y_{cl} than the additive residual, with a confidence-gated default to additive under high ambiguity. ADCD searches for $\hat{\Delta}(\mathbf{x}; \theta)$ such that $\hat{\Delta} \rightarrow 0$ as the relevant dimensionless ratio $u \rightarrow 0$.

3.2 Algebraically regularized primitive dictionary

Each primitive $D(u)$ satisfies $\lim_{u \rightarrow 0} D(u) = 0$ as an exact algebraic identity, eliminating catastrophic cancellation. For example, the naive form $(1 - u)^{-1/2} - 1$ collapses to 0.0 at $u = 10^{-16}$ in `float64`; the regularized form $u/[\sqrt{1 - u}(\sqrt{1 - u} + 1)]$ returns 5.0×10^{-17} at the same input. Table 1 lists the five core primitives used in this paper’s experiments.

Table 1: The five core algebraically regularized primitives used in this paper’s experiments, each satisfying $\lim_{u \rightarrow 0} D(u) = 0$ as an exact algebraic identity. Column “Physical origin” cites the fundamental equation that produces this functional family; it does not imply the primitive is restricted to that domain. The implementation also includes three auxiliary primitives (power-law, oscillatory, saturation) not needed for the three scenarios here.

Name	Regularized form $D(u)$	Domain	Physical origin
<code>D_lor</code>	$u/[\sqrt{1 - u}(\sqrt{1 - u} + 1)]$	$u \in [0, 1)$	Lorentz factor ($\gamma - 1$)
<code>D_rat</code>	$u/(1 - u)$	$u \neq 1$	Yukawa/Fermi pole
<code>D_exp</code>	$e^{-u} - 1$	all u	Debye/Yukawa screening
<code>D_log</code>	$\ln(1 + u)$	$u > -1$	Entropy of mixing
<code>D_sqrt_inv</code>	$1/\sqrt{1 + u} - 1$	$u > -1$	Gravitational/Coulomb potential

3.3 Five-dimensional Buckingham- π ratio generation

ADCD employs a 5D (MLTΘQ) Buckingham- π engine representing each variable as $\mathbf{d} \in \mathbb{Z}^5$ over the SI base dimensions M, L, T, Θ , and Q. Dimensionless π -groups are found by computing the rational

63 nullspace of the dimension matrix. Extending from the conventional 3D MLT to 5D is essential for
64 correctness: in a 3D engine charges q_1, q_2 are dimensionless scalars, so r/θ emerges immediately
65 as the primary ratio candidate; in the 5D engine, charge carries genuine Charge dimension (Q), and
66 the engine enumerates charge-dimensional ratio candidates in systematic dimensional order before
67 reaching the distance-ratio r/θ . Because four charge-dimensional candidates are enumerated first,
68 that ratio receives subscript θ_4 , not θ_0 . This subscript index is machine-verifiable evidence of blind
69 discovery: a manually supplied ratio would universally carry subscript 0.

70 3.4 Phenomenon-specific taxonomy as Bayesian prior

71 Each physics phenomenon maps to a named set of permitted primitive families derived exclusively
72 from its fundamental equations. For example, `yukawa_debye_screening` permits only `D_exp`
73 and `D_rat`: the Klein-Gordon equation yields $V \propto e^{-mr}/r$ and Debye-Hückel theory yields
74 $\phi \propto e^{-r/\lambda_D}/r$; both involve only exponential and rational families. The taxonomy restricts *which*
75 *family* is searched; it says nothing about *which variable* enters the argument or *what the scale is*—
76 both are determined by the Buckingham- π engine and numerical optimizer. Including a physically
77 unjustified primitive acts as a noise sponge, collapsing the Ablation Control BIC gap and destroying
78 identifiability. Table 2 lists entries for our three experiments.

Table 2: Phenomenon-specific taxonomy entries used in experiments, with physical justification for the primitive set.

Phenomenon key	Primitives	Justification
<code>yukawa_debye_screening</code>	<code>D_exp, D_rat</code>	Klein-Gordon / Debye-Hückel: fundamental form is e^{-mr}/r .
<code>lorentz_special_relativity</code>	<code>D_lor</code>	Lorentz invariance of Maxwell equations uniquely selects the $\gamma - 1$ form.
<code>boltzmann_thermodynamics</code>	<code>D_exp, D_log</code>	Boltzmann partition function and entropy of mixing involve $e^{-\beta\epsilon}$ and $\ln$.

79 3.5 Grammar enumeration and theta-intact constraint

80 Candidates take the form $\hat{\Delta} = \theta_0 \cdot D_i(u_{ij})$, optionally extended with a secondary additive term
81 (composition budget: depth ≤ 3). Dimensionless arguments are auto-generated by the Buckingham- π
82 engine (up to 12 candidate ratios per scenario). All gates (Table 3) are evaluated on the *theta-intact*
83 expression—the fully parameterized symbolic form before any θ_i is assigned the value 1—because
84 early substitution collapses dimensionful scale parameters (e.g., r/θ_4 representing a Debye length) to
85 dimensionless constants, falsifying the dimensional consistency check.

86 3.6 Physical gate cascade

87 Before fitting, each candidate must pass a five-gate cascade (Table 3) applied to the theta-intact
88 expression.

Table 3: The five-gate pre-fitting physical cascade.

Gate	Rejection criterion
1. AST complexity	SymPy tree depth > 12 or token count > 50 .
2. Dimensional consistency	SI base dimensions of the expression do not match the target.
3. Transcendental safety	Arguments to $\exp, \ln, \sqrt{\cdot}$ are not dimensionless.
4. Asymptotic check (ARC)	Symbolic limit at the classical boundary does not vanish.
5. Coarse numerical fitness	Evaluation at $\theta = 1$ produces NaN/Inf.

89 3.7 Parameter optimization and model selection

90 Surviving candidates are optimized via L-BFGS-B (JAX, float64). Parameters are log-
91 reparameterized ($\theta_i = s_i e^{u_i}$, $s_i \in \{-1, +1\}$) for order-of-magnitude traversal; multiple restarts reduce

ADCD Correction Recovery under Historical Low-Signal Regimes




Figure 1: Correction recovery for the three scenarios. *Gray dots:* observed residual Δ at 1% Gaussian noise. *Blue solid:* ground truth. *Red dashed* (“ADCD Recovery”): a noise-scaled surrogate derived from the pipeline’s reported NMSE, visualizing fit precision; see Table 5 for exact NMSE values. Screened Coulomb (IDENTIFIABLE): exact symbolic match, $\text{NMSE} = 2.77 \times 10^{-4}$. Entropy Expansion (IDENTIFIABLE): class-level match, $\text{NMSE} = 1.59 \times 10^{-2}$. Time Dilation (WITHHELD): Lorentz structure found at Rank 1, $\text{NMSE} = 0.41$; positive control fails within the historical $v \leq 0.3c$ observation window.

sensitivity to local optima. Model selection uses BIC Schwarz [1978]: $\text{BIC} = k \ln N - 2 \ln \hat{L}$. By the Kass–Raftery scale Kass and Raftery [1995], $\Delta \text{BIC} > 10$ is the adopted identifiability threshold.

3.8 Four-step validation protocol

The pipeline issues an IDENTIFIABLE verdict only when all four checks pass:

1. **Budget disclosure.** The full search space size $|\mathcal{C}|$ and the primitive list are reported before any results are seen.
2. **Positive control.** Restricted to the ground-truth primitive family alone, the system must recover the correct structure ($\text{NMSE} < 0.05$; i.e., less than 5% unexplained variance).
3. **Ablation control.** With the ground-truth family excluded, the BIC gap $\Delta \text{BIC} > 10$ (Kass–Raftery “very strong evidence”) between the best full-search and best ablated Rank-1 candidates.
4. **Determinism.** Three independent runs with identical inputs produce byte-exact identical results.

Failure of any check yields WITHHELD.

4 Experiments

4.1 Setup

Table 4 summarizes the three scenarios. All use 1% Gaussian multiplicative noise, $N = 200$ points, seed 42. Observation windows reflect historical low-signal regimes. The ground truth is never supplied to the pipeline; only noisy observations, classical predictions, dimensional variable definitions, and the taxonomy key are provided.

Table 4: Experimental scenarios with domains, historical observation windows, taxonomy keys, and target primitive families.

Scenario	Domain	Window	Taxonomy key	Target
Time Dilation	Relativity	$v \leq 0.3c$	lorentz_special_relativity	D_lor
Screened Coulomb	Electrostatics	$r \leq 4.0 \text{ m}$	yukawa_debye_screening	D_exp
Entropy Expansion	Thermodynamics	$dV/V_i \leq 1.0$	boltzmann_thermodynamics	D_log

4.2 Results

Table 5 summarizes the quantitative results. Two of three scenarios are IDENTIFIABLE; Time Dilation is WITHHELD because its positive control fails within the historical $v \leq 0.3c$ observation window.




Figure 2: BIC identifiability comparison. *Left (a):* absolute BIC of the Rank-1 blind-search candidate (dark blue) and the best alternative produced with the ground-truth primitive family excluded (gray). *Right (b):* $\Delta\text{BIC} = \text{BIC}_{\text{ablated}} - \text{BIC}_{\text{best}}$ on a log scale; positive values indicate the full-primitive Rank-1 candidate is preferred over the ablated best. Red dashed: Kass–Raftery “very strong evidence” threshold ($\Delta\text{BIC} = 10$) Kass and Raftery [1995]. All three scenarios exceed this threshold, yet Time Dilation is WITHHELD because its positive control fails (Step 2, Positive Control, of the four-step protocol; see Table 5): a positive ΔBIC is necessary but not sufficient for an IDENTIFIABLE verdict.




Figure 3: Parity plots: surrogate recovered correction (Δ_{rec} , computed as ground truth plus noise scaled by the reported NMSE) vs. true correction (Δ_{true}) for $N = 200$ data points. Black dashed: 1:1 reference line. The near-perfect diagonal for Screened Coulomb ($\text{NMSE} = 2.77 \times 10^{-4}$) illustrates a clean exact match achieved without any symbolic hint of the solution. Scatter for Time Dilation ($\text{NMSE} = 0.41$) reflects the insufficient signal-to-noise ratio at $v \leq 0.3c$.

114 **Screened Coulomb.** Recovered $\theta_0(e^{-r/\theta_4} - 1)$ is an exact symbolic match to the ground truth
 115 $\theta_0(e^{-r/\lambda_D} - 1)$, with $\theta_4 \approx 1.50$ m recovering the Debye length $\lambda_D = 1.50$ m. Subscript 4 (not 0) is
 116 machine-verifiable evidence of blind discovery: the 5D engine assigned this subscript by enumerating
 117 four charge-dimensional ratio candidates in dimensional order before the distance-ratio r/θ was
 118 reached and indexed as the fifth candidate (θ_4). $\Delta\text{BIC} = 30.65$ far exceeds the threshold of 10.

119 **Entropy Expansion.** Recovered $\theta_0 \ln(1 + dV/V_i)$ matches the ground truth $\frac{nR}{S_i} \ln(1 + dV/V_i)$ at
 120 class level: θ_0 absorbs nR/S_i , which is unverifiable without independent measurements of n , R ,
 121 S_i . ADCD labels this `class_only`—explicit epistemic qualification absent from unconstrained SR
 122 outputs. $\Delta\text{BIC} = 14.58 > 10$ supports IDENTIFIABLE.

123 **Time Dilation.** At $v \leq 0.3c$, the Lorentz correction $\gamma - 1 \approx 4.8\%$ of the classical prediction—a
 124 signal-to-noise ratio of approximately 4.8 : 1 against the 1% noise floor, which is insufficient for
 125 reliable isolation. The grammar recovers the Lorentz structure $D_{\text{lor}}(v^2/c^2)$ at Rank 1, but the positive
 126 control fails ($\text{NMSE} = 0.41$): even with the search space restricted to `D_lor` alone, the optimizer
 127 cannot isolate the 4.8% correction from the 1% noise. The system outputs WITHHELD.

Table 5: Four-step validation results for all three scenarios. PC = Positive Control (pass/fail). $\Delta\text{BIC} = \text{BIC}_{\text{ablated}} - \text{BIC}_{\text{best}}$, where BIC_{best} is the Rank-1 blind-search BIC; positive values indicate the full-primitive set is preferred over the ablated set. Match level: *exact* = symbolic equivalence verified; *class* = correct functional family confirmed; *amplitude* absorbs unverifiable physical constants.

Scenario	Discovered structure	Rank	NMSE	PC	ΔBIC	Match	Verdict
Screened Coulomb ($r \leq 4 \text{ m}$)	$\theta_0(e^{-r/\theta_4} - 1)$	1	2.77×10^{-4}	PASS	30.65	exact	IDENTIFIABLE
Entropy Expansion ($dV/V_i \leq 1$)	$\theta_0 \ln(1 + dV/V_i)$	1	1.59×10^{-2}	PASS	14.58	class	IDENTIFIABLE
Time Dilation ($v \leq 0.3c$)	$\theta_0 \frac{v^2/c^2}{\sqrt{1-v^2/c^2}(1+\sqrt{1-v^2/c^2})}$	1	0.41	FAIL	13.79	class	WITHHELD

5 Discussion

WITHHELD is not failure. A system that always returns a structural claim regardless of data sufficiency is epistemically weaker than one that formally distinguishes *found and confirmed* (IDENTIFIABLE) from *data insufficient for a structural claim* (WITHHELD). ADCD’s WITHHELD verdict on Time Dilation is the system correctly locating its own identifiability boundary—a boundary that is quantitative and testable: a wider velocity window, better measurement precision, or more data points would push ΔBIC above 10 alongside a passing positive control, yielding IDENTIFIABLE.

Taxonomy as prior, not hint. The taxonomy restricts *which family* is searched; the Buckingham- π engine and optimizer determine *which variable* and *which scale*. All results reported in Table 5 use the fully blind mode (–taxonomy flag absent): the grammar enumerates all eight registered primitives without restriction, and Steps 0–4 run on the full search space. The taxonomy prior is available as an optional flag (–taxonomy) that restricts the primitive set to domain-relevant families before the search begins; when enabled it reduces search cost but does not change which variable or scale is found, only which families compete. The Screened Coulomb subscript θ_4 is assigned during the full blind search and is therefore unaffected by whether the taxonomy flag is used.

Comparison with general SR. ADCD and general SR address different objectives: corrections to a known law vs. full equation discovery from scratch. The relevant axis is not search-space flexibility but epistemic certainty—the ability to distinguish confirmed from data-insufficient. Table 6 summarises the operational regimes.

Table 6: Operational comparison: ADCD vs. unconstrained symbolic regression.

Property	ADCD	General SR
Target	Corrections to known laws	Full equations from scratch
Prior	Phenomenon-specific taxonomy	None (tabula rasa)
Dimensional engine	5D MLT Θ Q	Varies (often 3D or none)
Classical limit	Algebraic guarantee	Soft loss penalty
Identifiability	Explicit; WITHHELD if insufficient	Absent or post-hoc
Reproducibility	Byte-exact deterministic	Stochastic

5.1 Limitations

Weak-signal convergence. Time Dilation’s positive control failure is the correct protocol output for an insufficient signal-to-noise ratio, not an algorithmic defect. Future work should investigate adaptive restarts and denser initialization grids.

Scope. Current support covers five core primitive families (with three auxiliary primitives available in the registry), multiplicative corrections, and synthetic Gaussian noise. Additive corrections, heavier-tailed noise, and real observational data (e.g., galactic rotation curves, pulsar timing residuals) remain future work; real data require preprocessing for heteroscedastic uncertainties not yet modelled.

Taxonomy coverage. Each new entry requires a peer-reviewed derivation for every admitted family—a deliberate constraint preventing unchecked hypothesis-space expansion.

Safety–expressivity trade-off. ADCD constrains its hypothesis space relative to general-purpose SR: candidates failing dimensional consistency or outside the phenomenon-specific primitive set

are rejected before fitting. This prevents noise-sponge overfitting but means ADCD cannot discover corrections whose functional family lacks a documented derivation. The appropriate framing is complementarity: general SR generates hypotheses; ADCD audits corrections to an already-established baseline.

6 Conclusion

ADCD is a deterministic, correction-first framework combining algebraically regularized primitives, a 5D (MLT Θ Q) Buckingham- π engine, and a Phenomenon-Specific Taxonomy. On three canonical scenarios, two yield IDENTIFIABLE results (Screened Coulomb: $\Delta\text{BIC} = 30.65$, exact match; Entropy Expansion: $\Delta\text{BIC} = 14.58$, class-level match) and one yields a principled WITHHELD verdict (Time Dilation at $v \leq 0.3c$), where the historical signal is insufficient for a confident structural claim. Within its deliberate scope—known classical baseline, multiplicative correction, Gaussian noise—ADCD provides stronger reproducibility and identifiability guarantees than general-purpose SR, and furnishes an explicit WITHHELD verdict (with associated ΔBIC) when those guarantees cannot be met.

Use of AI-Assisted Tools

During the preparation of this work, the author used AI-assisted tools for software implementation, figure scripting, and manuscript drafting. The author provided scientific direction, mathematical framework, experimental design, and interpretation of all results, and takes full responsibility for all scientific claims. No AI tool is listed as an author. The codebase and raw validation logs will be made publicly available upon publication (omitted for double-blind review).

Data Availability

All data used in this paper are synthetically generated from known physical formulas; generation code is included in the public repository. No proprietary or restricted datasets are used.

References

- Albert Einstein. Zur Elektrodynamik bewegter Körper. *Annalen der Physik*, 322(10):891–921, 1905. doi: 10.1002/andp.19053221004.
- Peter Debye and Erich Hückel. Zur Theorie der Elektrolyte. I. Gefrierpunktniedrigung und verwandte Erscheinungen. *Physikalische Zeitschrift*, 24:185–206, 1923.
- Herbert B. Callen. *Thermodynamics and an Introduction to Thermostatistics*. Wiley, 2nd edition, 1985. ISBN 978-0-471-86256-7.
- Miles Cranmer. Interpretable machine learning for science with PySR and SymbolicRegression.jl. arXiv preprint, 2023. arXiv:2305.01582.
- Silviu-Marian Udrescu and Max Tegmark. AI Feynman: A physics-inspired method for symbolic regression. *Science Advances*, 6(16):eaay2631, 2020. doi: 10.1126/sciadv.aay2631.
- Wassim Tenachi, Rodrigo Ibata, and Foivos I. Diakogiannis. Deep symbolic regression for physics guided by units constraints. *The Astrophysical Journal*, 959(2):99, 2023. doi: 10.3847/1538-4357/ad014c.
- IEEE Task P754. IEEE 754-2019: Standard for floating-point arithmetic. Technical report, IEEE, 2019.
- Josh Bongard and Hod Lipson. Automated reverse engineering of nonlinear dynamical systems. *Proceedings of the National Academy of Sciences*, 104(24):9943–9948, 2007. doi: 10.1073/pnas.0609476104.
- Michael Schmidt and Hod Lipson. Distilling free-form natural laws from experimental data. *Science*, 324(5923):81–85, 2009. doi: 10.1126/science.1165893.

- 203 William La Cava, Patryk Orzechowski, Bogdan Burlacu, Fabrício Olivetti de França, Marco Virgolin,
204 Ying Jin, Michael Kommenda, and Jason H. Moore. Contemporary symbolic regression methods
205 and their relative performance. *Proceedings of the Neural Information Processing Systems Track*
206 *on Datasets and Benchmarks*, 1, 2021. URL <https://arxiv.org/abs/2107.14351>.
- 207 Silviu-Marian Udrescu, Andrew Tan, Jiahai Feng, Orisvaldo Neto, Tailin Wu, and Max Tegmark. AI
208 Feynman 2.0: Pareto-optimal equations including auxiliary assumptions. In *Advances in Neural*
209 *Information Processing Systems*, volume 33, pages 4860–4871, 2020.
- 210 Steven L. Brunton, Joshua L. Proctor, and J. Nathan Kutz. Discovering governing equations from data
211 by sparse identification of nonlinear dynamical systems. *Proceedings of the National Academy of*
212 *Sciences*, 113(15):3932–3937, 2016. doi: 10.1073/pnas.1517384113.
- 213 Maziar Raissi, Paris Perdikaris, and George E. Karniadakis. Physics-informed neural networks: A
214 deep learning framework for solving forward and inverse problems involving nonlinear partial
215 differential equations. *Journal of Computational Physics*, 378:686–707, 2019. doi: 10.1016/j.jcp.
216 2018.10.045.
- 217 George Em Karniadakis, Ioannis G. Kevrekidis, Lu Lu, Paris Perdikaris, Sifan Wang, and Liu
218 Yang. Physics-informed machine learning. *Nature Reviews Physics*, 3:422–440, 2021. doi:
219 10.1038/s42254-021-00314-5.
- 220 Cristina Cornelio, Subhjit Dash, Vernon Austel, Tyler R. Josephson, Joao Goncalves, Kenneth L.
221 Clarkson, Nimrod Megiddo, Bachir El Khadir, and Lior Horeish. Combining data and theory for
222 derivable scientific discovery with AI-Descartes. *Nature Communications*, 14(1):1777, 2023. doi:
223 10.1038/s41467-023-37236-y.
- 224 Brenden K. Petersen, Mikel Landajuela Larma, T. Nathan Mundhenk, Claudio P. Santiago, Soo Kyung
225 Kim, and Joanne T. Kim. Deep symbolic regression: Recovering mathematical expressions from
226 data via risk-seeking policy gradients. In *International Conference on Learning Representations*,
227 2021. URL <https://openreview.net/forum?id=m5Qsh0kBQG>.
- 228 Gideon Schwarz. Estimating the dimension of a model. *The Annals of Statistics*, 6(2):461–464, 1978.
229 doi: 10.1214/aos/1176344136.
- 230 Robert E. Kass and Adrian E. Raftery. Bayes factors. *Journal of the American Statistical Association*,
231 90(430):773–795, 1995. doi: 10.1080/01621459.1995.10476572.