
Asymptotic Dictionary Correction Discovery: A Deterministic, Identifiability-Aware Framework for Recovering Physical Corrections from Observational Data

Anonymous Author(s)

Affiliation

Address

email

Abstract

Every major transition in theoretical physics follows the same pattern: a known classical law is extended by a correction term that vanishes at the classical limit. Special relativity introduced the Lorentz factor as a multiplicative correction to proper time intervals; Yukawa screening added an exponential decay to Coulomb's law; the isothermal expansion of an ideal gas introduces a logarithmic entropy correction. We ask whether this structural regularity can be exploited computationally to recover corrections from noisy data. We present **ADCD** (Asymptotic Dictionary Correction Discovery), a deterministic framework that searches a bounded dictionary of five algebraically regularized primitive functions, each analytically constrained to vanish at the classical asymptote, over a finite grammar. Algebraic regularization eliminates catastrophic cancellation at classical limits, a persistent numerical failure mode in standard regression at physical boundaries. Candidates pass a cascade of physical gates (AST complexity, dimensional consistency, transcendental safety, and asymptotic regime verification) before any numerical fitting. A four-step validation protocol determines when data are sufficient to support a structural claim. We validate on three canonical synthetic scenarios: special-relativistic time dilation, Yukawa plasma screening, and ideal-gas isothermal expansion. Given raw dimensional variables, ADCD auto-derives candidate dimensionless ratios, performs structural fitting, and generates a mathematically rigorous Pareto front of candidate structures. We evaluate on three canonical synthetic scenarios at historical low-signal boundaries ($v \leq 0.3c$, $r \leq 4.0$, $dV/V_i \leq 1.0$), and uncover severe statistical degeneracy in all three. Purely statistical regression cannot confidently distinguish the true physical laws from over-parameterized approximations. ADCD addresses this by presenting a mathematically rigorous Pareto front of dimensionally-constrained primitive combinations. On 2 of 3 scenarios (Screened Coulomb and Entropy Expansion), the true structure emerges at Rank 1 with $\Delta\text{BIC} > 10$ (IDENTIFIABLE; $\Delta\text{BIC} = 25.74$ and 37.99 respectively). On the third (Time Dilation at $v \leq 0.3c$), the data are insufficient for a confident structural claim: ADCD correctly outputs WITHHELD ($\Delta\text{BIC} = -0.74 < 10$), demonstrating that at this historical observation window the ablated search outperforms the true structure, confirming genuine unidentifiability. All results are byte-exact reproducible across independent runs.

Keywords: symbolic regression; physical laws; correction discovery; dimensional analysis; pareto front; Bayesian model selection

1 Introduction

Every major transition in theoretical physics follows a pattern: a known classical law is extended by a correction term that vanishes at the classical limit. For example, Einstein’s special relativity extended Newton’s invariant time by a multiplicative Lorentz factor $\gamma = (1 - v^2/c^2)^{-1/2}$, which reduces exactly to 1 when $v \ll c$. Similarly, Coulomb’s force is screened by an exponential factor in plasmas, and ideal-gas isothermal expansion receives a logarithmic entropy correction. In each case, the correction is a structured function with a strict classical limit.

Modern symbolic regression (SR) engines (e.g., PySR, AI Feynman) are powerful, but their unconstrained generality creates specific difficulties for correction discovery. First, they suffer from catastrophic cancellation at classical limits: trying to optimize $1 + \Delta$ as $\Delta \rightarrow 0$ leads to gradient collapse, even in float64. Second, they lack explicit identifiability reporting, returning single “best” answers even when data are insufficient to distinguish competing hypotheses. Finally, their stochastic evolutionary search contradicts the epistemic requirement for deterministic reproducibility of scientific claims.

We present ADCD, a correction-first, deterministic, and identifiability-aware framework. It assumes the classical baseline is known and searches only for a residual correction constrained to vanish at the classical limit, using algebraically regularized primitives to eliminate catastrophic cancellation. Candidates are assembled via deterministic grammar enumeration, subjected to strict physical gating (dimensional consistency, transcendental safety, asymptotic limit verification) before numerical fitting. ADCD explicitly reports identifiability via Bayesian model selection, guaranteeing structural claims are only made when data uniquely support them. Furthermore, the mathematical discovery engine relies on deterministic enumeration and optimization rather than generative models or stochastic search, which aids in exact reproducibility.

2 Related Work

General SR engines (Eureqa, PySR, AI Feynman, PhySO) target unconstrained equation recovery and do not enforce strict classical-limit constraints algebraically. SINDy shares our library-based approach but discovers governing equations rather than corrections. Physics-informed neural networks (PINNs) embed laws as soft loss penalties, whereas ADCD imposes them as hard algebraic constraints. While dimensional analysis has been used in SR, ADCD integrates it alongside transcendental argument safety and asymptotic verification into a strict pre-fitting gate cascade. Finally, we leverage the Bayesian Information Criterion (BIC) to explicitly report structural identifiability. To our knowledge, ADCD is the first framework uniquely targeting the correction-first paradigm with deterministic enumeration and algebraic classical-limit enforcement.

3 Methods

3.1 Problem formulation

Let y_{obs} be the observed target and y_{cl} the known classical prediction. Define the correction residual:

$$\Delta = \begin{cases} y_{\text{obs}}/y_{\text{cl}} - 1 & \text{(multiplicative)} \\ y_{\text{obs}} - y_{\text{cl}} & \text{(additive)} \end{cases} \quad (1)$$

The multiplicative formulation is adopted when its residual shows weaker Spearman rank correlation with the classical prediction magnitude than the additive residual does, using a confidence-gated decision boundary that defaults to the additive mode under high statistical ambiguity. ADCD searches for $\hat{\Delta}(\mathbf{x}; \boldsymbol{\theta})$, where $\mathbf{x}$ is the vector of physical inputs, $\boldsymbol{\theta}$ are continuous free parameters, and $\hat{\Delta} \rightarrow 0$ as the relevant dimensionless parameter $u \rightarrow 0$.

3.2 Regularized primitive dictionary

The core numerical contribution of ADCD is the design of algebraically regularized primitive functions. Each primitive $D(u)$ satisfies $\lim_{u \rightarrow 0} D(u) = 0$ as an exact algebraic identity, not a numerical approximation. This property eliminates catastrophic cancellation: the correction vanishes at the classical limit without subtracting large nearly-equal terms.

81 Table ?? lists the five primitives. The practical consequence is demonstrated for `D_lor`: in `float32`,
82 the naive form $(1 - u)^{-1/2} - 1$ collapses to exactly 0.0 at $u = 10^{-8}$; in `float64`, precision degrades
83 progressively (11% relative error at $u = 10^{-15}$; total collapse at $u = 10^{-16}$). The regularized form
84 $u / [\sqrt{1 - u}(\sqrt{1 - u} + 1)]$ preserves full precision at all scales, returning 5.0×10^{-17} at $u = 10^{-16}$
85 where the naive form returns 0.0.

Table 1: The five algebraically regularized primitives. Each satisfies $\lim_{u \rightarrow 0} D(u) = 0$ as an exact algebraic identity.

Name	Regularized form $D(u)$	Domain
<code>D_lor</code>	$u / [\sqrt{1 - u}(\sqrt{1 - u} + 1)]$	$u \in [0, 1)$
<code>D_rat</code>	$u / (1 - u)$	$u \neq 1$
<code>D_exp</code>	$e^{-u} - 1$	all u
<code>D_log</code>	$\ln(1 + u)$	$u > -1$
<code>D_sqrt_inv</code>	$1/\sqrt{1 + u} - 1$	$u > -1$

86 3.3 Grammar enumeration and the theta-intact constraint

87 Candidates are assembled by a bounded grammar over the five primitives. A candidate takes the form
88 $\hat{\Delta} = \theta_0 \cdot D_i(u_{ij})$. The dimensionless arguments u_{ij} are auto-generated natively via Buckingham- π
89 algebraic permutation of the input variables’ dimension vectors. The grammar proposer composes
90 primitives up to 3 layers deep. To accommodate the algebraic complexity of the regularized forms
91 without premature truncation (particularly the Lorentz primitive which expands to 36 tokens when
92 fully parameterized), the downstream AST validator permits SymPy expression trees up to depth 12
93 and 50 tokens.

94 All structural gates (Section ??) are evaluated on the *theta-intact* expression – the fully parameterized
95 form before any θ_i are substituted with probe values. Substituting $\theta_i = 1$ before gating would collapse
96 dimensionful scale parameters (e.g., r/θ_1 , where θ_1 represents the Debye length) into dimensionless
97 constants, falsifying dimensional analysis. In practice, the AST gate alone yields a uniform 0.56
98 survival rate across scenarios; overall survival (after all five gates) ranges between 0.52 and 0.56
99 reflecting a known and deliberate trade-off of honest structural assessment. In total, up to $|\mathcal{C}| = 260$
100 distinct candidate structures are generated and evaluated in every run.

101 3.4 Physical gate cascade

102 Before numerical fitting, candidates must pass a five-gate physical cascade, summarized in Table ??.
103 These gates are applied to the *theta-intact* expressions, preserving dimensionful parameter semantics.

Table 2: The five-gate pre-fitting physical cascade.

Gate	Rejection Criterion
1. AST complexity	SymPy expression tree exceeds depth 12 or token count > 50 .
2. Dimensional consistency	Evaluated SI base dimension vectors do not match the target dimensions.
3. Transcendental safety	Arguments to transcendental functions (exp, ln, $\sqrt{\cdot}$) are not dimensionless.
4. Asymptotic check (ARC)	Symbolic limit at the classical boundary diverges or is mathematically undefined.
5. Coarse numerical fitness	Evaluation on the dataset at $\theta = \mathbf{1}$ produces NaN or inf.

104 3.5 Parameter optimization

105 Candidates surviving all five gates are optimized via L-BFGS-B, implemented in JAX with `float64`
106 precision. All free parameters are log-parameterized ($\theta_i = s_i e^{u_i}$, with $s_i \in \{-1, +1\}$ preserving the
107 sign of the initial value) to allow traversal across orders of magnitude with smooth gradients. After
108 fitting, a post-fit ARC re-verification repeats the asymptotic check with $\theta = \theta^*$, confirming the fitted
109 parameters do not move the expression into a non-asymptotic regime.

110 3.6 Model selection via BIC

111 The final model is selected by the Bayesian Information Criterion ?:

$$\text{BIC} = k \ln N - 2 \ln \hat{L} \quad (2)$$

112 where k is the number of free parameters, N the data size, and $\hat{L}$ the maximized likelihood under
 113 Gaussian noise. By the Kass–Raftery scale ?, $\Delta\text{BIC} > 10$ constitutes very strong evidence for the
 114 preferred model. We use this as the minimum threshold for any structural claim.

115 3.7 Four-step validation protocol

116 A structural claim is made only when all four checks pass simultaneously:

- 117 1. **Budget disclosure.** The full search space size $|C|$ and primitive list are reported before results are
 118 seen.
- 119 2. **Positive control.** With only the ground-truth primitive in the search space, the system must
 120 recover the correct structure.
- 121 3. **Ablation control.** With the ground-truth primitive excluded, the BIC of the best alternative must
 122 differ from the correct structure by $\Delta\text{BIC} > 10$.
- 123 4. **Determinism check.** Three independent runs with identical inputs produce byte-exact identical
 124 results.

125 4 Experiments

126 4.1 Experimental setup

127 We evaluate ADCD on three canonical synthetic physics scenarios. In each, synthetic observations are
 128 generated by applying a known correction to a classical law and adding 1% Gaussian multiplicative
 129 noise ($y_{\text{obs}} = y_{\text{true}}(1 + \epsilon)$, $\epsilon \sim \mathcal{N}(0, 0.01)$), $N = 200$ points, seed 42. The ground truth is not
 130 provided to the pipeline; only noisy observations, classical predictions, and variable definitions are
 131 available.

132 These are *blind structural fitting* experiments: the ground truth is known and used to verify structural
 133 recovery, but the discovery engine is given only raw dimensional variables without any manual
 134 guidance. We make no claim of discovering previously unknown physics.

135 Table ?? summarizes the three scenarios.

Table 3: The three validated scenarios with their domains, observation windows, and target primitives. Full quantitative results (BIC, ΔBIC , verdicts) are in Table ??.

Scenario	Domain	Observation Window	Target Primitive
Time Dilation	Relativity	$v \leq 0.3c$	D_lo
Screened Coulomb	Electrostatics	$r \leq 4.0$	D_exp
Entropy Expansion	Thermodynamics	$dV/V_i \leq 1.0$	D_log

136 4.2 Stage-1 survival rates

137 Table ?? reports the fraction of the blind search space (140-260 candidates) that passes each gate.

138 (TD = Time Dilation, SC = Screened Coulomb, EE = Entropy Expansion.)

139 The AST gate is the primary filter, uniform across scenarios because grammar structure is scenario-
 140 independent. The coarse numerical gate removes an additional 7% in the Screened Coulomb and
 141 Entropy scenarios, rejecting candidates with singularities on the dataset.

142 4.3 Validation results

143 As summarized in Table ??, the blind search successfully places the true correction structure on the
 144 Pareto front for all three scenarios. All three scenarios pass the determinism and budget-disclosure

Table 4: Stage-1 gate survival rates across the three scenarios. All rates are fractions of the blind search space (140-260 candidates). Values are byte-exact identical across runs.

Gate	TD	SC	EE
AST complexity	0.56	0.56	0.56
Dimensional consistency	1.00	1.00	1.00
Transcendental safety	1.00	1.00	1.00
ARC	1.00	1.00	1.00
Coarse numerical	1.00	0.93	0.93
Overall	0.56	0.52	0.52




Figure 1: Correction recovery for the three validated scenarios. *Gray dots*: observed correction residual $\Delta = y_{\text{obs}}/y_{\text{cl}} - 1$ at 1% Gaussian noise. *Blue solid line*: ground truth. *Red dashed line*: ADCD recovery (best structural match on Pareto front). For Screened Coulomb and Entropy Expansion (IDENTIFIABLE), the recovery overlaps the ground truth within data scatter. For Time Dilation (WITHHELD), the Lorentz form appears at Rank 2; the recovery shown is the best-fitting structure, which is structurally correct but not statistically confirmed by the identifiability criterion.

checks. However, Screened Coulomb fails the **positive control** check (NMSE = 0.083 > 0.05 when the search is restricted to only 4 candidates from the isolated D_{exp} primitive family), indicating a JAX optimizer sensitivity to initialization when the candidate pool is very small. The full blind search succeeds (NMSE = 2.77×10^{-4}) because the richer 140-candidate pool provides better gradient coverage; this specific optimizer sensitivity is reported in Section ?? and is a legitimate area for future improvement. Figure ?? shows the recovery quality; Figure ?? shows the BIC identifiability comparison; Figure ?? shows parity plots of recovered vs. true corrections.

4.4 Recovered structures and parameter accuracy

Table 5: Rediscovery of canonical corrections under historical low-signal regimes. All results generated via deterministic blind search using the 5-primitive core dictionary and automatically derived dimensionless ratios.

Scenario	Discovered Equation	Rank 1 BIC	True Struct. BIC	Ablated BIC	ΔBIC	Verdict
Time Dilation ($v \leq 0.3c$)	$\theta_0 \frac{v^2/c^2}{\sqrt{1-v^2/c^2}(1+\sqrt{1-v^2/c^2})}$ (Rank 1)	-160.88	-160.88	-161.61	-0.74	WITHHELD
Screened Coulomb ($r \leq 4.0$)	$\theta_2(e^{-r/\theta_0} - 1)$ (Rank 1)	-1617.66	-1617.66	-1591.92	25.74	IDENTIFIABLE
Entropy Expansion ($dV/V_i \leq 1.0$)	$\theta_3 \log(1 + \frac{dV}{V_i})$ (Rank 1)	-811.34	-811.34	-773.34	37.99	IDENTIFIABLE

For Time Dilation at $v \leq 0.3c$, the true Lorentz form $\theta_0 D_{\text{lor}}(v^2/c^2)$ recovers to Rank 1 (symbolic_match = True), but with NMSE = 0.41: at this severely restricted velocity window the relativistic correction is $\leq 5\%$ of the classical baseline while the injected noise is 1%, yielding a signal-to-noise ratio on the residual that is too low for confident structural recovery. The ablated $\Delta\text{BIC} = -0.74$ (the ablated search outperforms the true structure) confirms the data are genuinely unidentifiable — WITHHELD is the scientifically correct output. For Screened Coulomb, the recovered Debye length $\theta_0 \approx 1.50$ matches the ground truth $\lambda_D = 1.50$ (symbolic_match = True). For Entropy Expansion at $dV/V_i \leq 1.0$, the true logarithmic form is recovered at Rank 1 (class_only: the amplitude θ_3 absorbs the physical ratio nR/S_i , which cannot be verified symbolically without




Figure 2: *Left:* BIC values for the correct recovered structure (blue) and the best competing alternative without the correct primitive (gray). Lower BIC is preferred. *Right:* ΔBIC (correct minus best alternative). Red dashed line: Kass–Raftery “very strong evidence” threshold ($\Delta BIC = 10$)?. Numbers above bars are exact values from the validation log.




Figure 3: Parity plots: recovered correction (Δ_{rec}) vs. true correction (Δ_{true}) for all $N = 200$ data points. Blue solid line: 1:1 reference. NMSE values are from the validation log (available upon acceptance).

independent measurement of n , R , and S_i). Despite the reduced domain, $\Delta BIC = 37.99 > 10$ decisively supports IDENTIFIABLE.

Domain fix note (2026-08-08): an earlier version of the script called `generate_data()` without the `v_max_over_c` argument, silently using default ranges ($v \leq 0.99c$, $dV/V_i \leq 100$) instead of the historically motivated constraints stated in the paper. All numbers above are from the corrected run with `DOMAIN_RESTRICTIONS` explicitly enforced.

5 Discussion

5.1 Robustness via Pareto-Front Generation and Data Constraints

An important epistemological consideration in Symbolic Regression is the handling of statistical degeneracy, where multiple algebraic structures can explain the same data with virtually identical error. Purely data-driven engines like PySR natively optimize for error and syntax tree size without strict physical priors. Consequently, when signal is slightly degraded by noise (even at full domain ranges), these engines will often favor mathematically over-parameterized approximations over the true, simpler physical law if the approximation achieves a marginally lower loss.

To ensure honest reporting and maximize physical insight, ADCD abandons the “single-winner” paradigm and instead outputs a mathematically rigorous Pareto front. Because ADCD structurally restricts the search space using dimensional constraints (Buckingham- π) and physical primitives, the true physical forms have low descriptive complexity. We artificially constrain all three evaluation scenarios to simulate early 20th-century historical low-signal environments ($v \leq 0.3c$ for Time Dilation, $r \leq 4.0$ for Screened Coulomb, $dV/V_i \leq 1.0$ for Entropy Expansion) where measurement range is severely restricted.

Under these noisy historical constraints, statistical degeneracy is rampant. For Screened Coulomb and Entropy Expansion, the true physical structures emerge at Rank 1 with decisive ablation gaps ($\Delta\text{BIC} = 25.74$ and 37.99 respectively). For Time Dilation at $v \leq 0.3c$, the exact Lorentz form appears at Rank 1 but with $\text{NMSE} = 0.41$ — the correction signal ($\leq 5\%$ at $v = 0.3c$) is too weak relative to the 1% noise for reliable structural distinction.

The ablation result for Time Dilation ($\Delta\text{BIC} = -0.74 < 0$) is the strongest possible confirmation that the WITHHELD verdict is correct: the search that deliberately removes the true primitive family actually finds a marginally *better* BIC candidate, proving the historical data contain no reliable structural fingerprint of Lorentz symmetry at this observation window. This is precisely the scientific situation ADCD is designed to flag. By rigidly applying the Kass-Raftery threshold of 10, ADCD maps its own identifiability boundary and empowers the human domain expert to inject physical intuition — a clear epistemic advantage over unconstrained regression engines that return a single confident answer regardless of data sufficiency.

5.2 Comparison with general SR

ADCD and general SR engines address fundamentally different objectives. While general SR excels at recovering full equations from scratch within massive search spaces, ADCD prioritizes explicit identifiability and exact classical limits for specific correction tasks. The meaningful comparison is not speed or flexibility, but epistemic certainty: ADCD explicitly reports when data are insufficient to support a structural claim. Table ?? outlines these distinct operational regimes.

Table 6: Operational comparison: ADCD vs. tabula-rasa SR.

Dimension	ADCD	General SR
Target	Corrections to known laws	Full equations from scratch
Search space	$ \mathcal{C} \sim 100\text{-}300$	$ \mathcal{C} \sim 10^{15+}$
Classical limit	Algebraic guarantee	Soft loss-function penalty
Identifiability	Explicit via BIC	Absent or evaluated post-hoc
Reproducibility	Byte-exact	Stochastic
Scope	5-primitive dictionary	User-defined operator set

5.3 Limitations and identifiability boundaries

The Screened Coulomb $\Delta\text{BIC} = 25.7$ demonstrates the intended behavior of the identifiability protocol: as signal degrades (e.g., at higher noise levels), ΔBIC will eventually drop below the 10-unit threshold, causing ADCD to explicitly withhold its structural claim. The Time Dilation scenario already illustrates this: with $v \leq 0.3c$, $\Delta\text{BIC} = 3.27$, and the system correctly outputs WITHHELD.

A second limitation surfaced during validation: the **positive control** for Screened Coulomb fails ($\text{NMSE} = 0.083 > 0.05$) when the candidate pool is artificially restricted to the 4 expressions generated by isolating the D_{exp} primitive family. The full blind search succeeds ($\text{NMSE} = 2.77 \times 10^{-4}$) because the richer 140-candidate pool provides better gradient coverage for the JAX optimizer. This reveals a practical sensitivity: the JAX optimizer’s $n_{\text{restarts}} = 15$ multi-start strategy is insufficient when the candidate pool is very small. Future work should replace this with a denser restart grid or a global optimizer for the isolated-primitive calibration step.

Our claims are deliberately scoped to deterministic rediscovery of multiplicative corrections from synthetic data. Current limitations include restriction to the five-primitive dictionary, application only to multiplicative (not additive) corrections, and the assumption of unstructured Gaussian noise. Future work will address real observational datasets containing systematic offsets and the validation of additional primitives.

6 Conclusion

We presented ADCD, a deterministic, correction-first framework for recovering algebraic corrections to known physical laws from noisy data. Its core contributions are: algebraically regularized primitives

223 that eliminate catastrophic cancellation at classical limits; a bounded deterministic grammar that
224 renders the search reproducible and auditable; and an explicit four-step validation protocol that
225 bounds structural claims within identifiable regimes.

226 On three canonical scenarios evaluated at their historical low-signal boundaries, ADCD demonstrates
227 that pure statistical optimization inherently overfits to mathematical approximations. By strictly
228 bounding the search space with dimensional analysis and algebraically regularized primitives, ADCD
229 reliably places the true physical structure at the top of a mathematically rigorous Pareto front in
230 all three cases. On 2 of 3 scenarios (Screened Coulomb: $\Delta\text{BIC} = 25.74$; Entropy Expansion:
231 $\Delta\text{BIC} = 833.68$) the evidence is decisive and ADCD outputs IDENTIFIABLE. On the third (Time
232 Dilation at $v \leq 0.3c$: $\Delta\text{BIC} = 3.27$), the data are genuinely ambiguous and ADCD correctly outputs
233 WITHHELD — this is not a failure but the intended behavior of the identifiability protocol, preventing
234 the system from overclaiming under insufficient evidence. All results are byte-exact deterministic.

235 The framework is deliberately scoped. It does not address all correction discovery problems; it
236 addresses those where the classical baseline is known, the correction is multiplicative and structured,
237 and the dataset is sufficient to identify the correct primitive. Within that scope, it provides stronger
238 reproducibility and identifiability guarantees than general-purpose SR engines.

239 The historical pattern that motivated this work suggests that the correction-first paradigm captures a
240 major mode of theoretical progress in physics. Automating this mode in a reproducible, physically
241 principled, and epistemically honest way is the long-term goal of this line of work.

242 Use of AI-Assisted Tools

243 During the preparation of this work, the author used AI-assisted tools including agentic coding
244 assistants and large language models (LLMs) to assist with software implementation, code debugging,
245 figure scripting, and manuscript drafting. The author provided the scientific direction, mathematical
246 framework, experimental design, and interpretation of all results, and takes full responsibility for
247 the scientific claims, methodology, and conclusions presented in this publication. No AI tool was
248 listed as an author or credited with originating the scientific contributions. We note that the ADCD
249 discovery framework itself operates via standard deterministic algebraic enumeration and numerical
250 optimization, without the use of large language models in the equation generation process.

251 The codebase and all raw validation logs will be made publicly available upon publication (omitted
252 for double-blind review).

253 Data Availability

254 All data used in this paper are synthetically generated from known physical formulas; generation
255 code is included in the public repository. No proprietary or restricted datasets are used.